

Aspectually Enhanced Modelling of Sociotechnical Systems

Tom Wallis

Submitted in fulfilment of the requirements for the Degree of Doctor of Philosophy

School of Computing Science

College of Science and Engineering

University of Glasgow

© Tom Wallis

Table of Contents

1	Introduction	11
1.1	A primer on aspect orientation	11
1.2	Prior Work	11
1.3	Terms & definitions	12
2	Relevant Literature	13
2.1	Aspect Orientation	13
2.1.1	Alternative methods in Aspect Orientation	17
2.1.2	Criticisms of aspect-oriented programming	19
2.2	Dynamic Aspect Weaving	21
2.3	Aspect Orientation in Simulation & Modelling	25
2.3.1	Aspect Orientation & Business Process Modelling	29
2.4	Process Variance in Simulation & Data Generation	32
2.4.1	Discussion of Variation & Motivations for Variations in Process Models	33
2.4.2	Representing Variations in Process Models and their Outputs	35
2.5	Discussion	37

2.5.1	Research Questions	40
3	Prior Work	41
3.1	Motivations in originally implementing PyDySoFu	41
3.2	Original PyDySoFu Implementation	43
3.2.1	Weaving mechanism	43
3.2.2	Applying Process Mutations	44
3.2.3	Limitations	45
3.3	Additional Simulation Machinery	48
3.3.1	Fuzzi-Moss	48
3.3.2	Theatre_Ag	50
3.4	Example Studies using PyDySoFu for Behavioural Simulation	51
3.5	Discussion	52
4	Rewriting PyDySoFu	55
4.1	Requirements for Change	55
4.2	Python3-compatible Weaving Techniques	57
4.2.1	Abandoned techniques	57
4.2.2	A viable technique: import hooks	58
4.3	Import Hooks	60
4.3.1	Implementing import hooks	60
4.3.2	Strengths and weaknesses of import hooks	61
4.3.3	Weaving process	63

TABLE OF CONTENTS	5
4.4 Optimisations	67
4.4.1 Deep Hook Weaving	68
4.4.2 Non-Dynamic Weaving	68
4.4.3 Priority Ordering of Aspect Application	69
4.5 Discussion	69
5 RPGLite: A Mobile Game for Collecting Data	71
5.1 An Overview of RPGLite	73
5.1.1 RPGLite’s Design Implications	74
5.2 Implementation of RPGLite	77
5.2.1 Consent to participation	77
5.2.2 Mobile app	78
5.2.3 API & Server-side Logic	80
5.3 Empirical Play and Data Collected	81
6 Simulation Optimisation with Aspect Orientation	83
6.1 Aims	83
6.1.1 PyDySoFu Suitability	85
6.1.2 Proposed Experiment	86
6.2 Naive Model	88
6.3 Experimental Design	91
6.3.1 Models of Learning	93
6.3.2 Modelling Learning in RPGLite	95

6.3.3	Modelling Player Confidence in RPGLite	98
6.4	Aspects Applied	99
6.4.1	Aspects for model improvement	100
6.4.2	Aspects for Instrumentation	103
6.4.3	Aspects Implementing Models of Learning	106
6.5	Experimental Design using Aspectual Learning Models	109
6.5.1	Quantifying Similarity of Character Pair Selection	109
6.5.2	Annealing toward Player-Specific Learning Parameters	111
6.6	Experimental Results	111
6.7	Discussion	111
7	A chapter title here for the experiment moving aspects to new systems, or systems with some changes	113
7.1	(Reword) Experimental motivations / motivation of research question	113
7.1.1	Coupling of Model and Behaviour	114
7.1.2	(Reword) Gaps/opportunities left by previous experiment	114
7.1.3	Research Question	114
7.2	Experimental Design	114
7.2.1	Changes to RPGLite	114
7.2.2	Applying Behavioural Variations to New System	115
7.3	Applying Aspects from a Control System to a New System	115
7.3.1	Implementation	115

7.3.2	Results	115
7.4	Discussion	115
8	Future Work	117
8.1	Aspect-Oriented Metaprogramming in real-world Software Engineering	117
8.2	Augmentation of pre-existing models	118
8.3	Aspect Orientation’s utility for Research Software Engineers	119
8.4	Hypothesising Possible System Dynamics via Aspects	120
8.5	Standards for Aspect Orientation in Research Codebases	122
8.6	Standard Aspect-Oriented Model Features	124
8.7	Testing Frameworks to detect Unrealistic Behavioural Variances	125
8.8	Discovery of System Features by Fitting Aspect Parameters	128
8.9	Optimisation of Multiple Aspects	128
8.10	Future Work pertaining to RPGLite	128
8.10.1	Causes of Game Abandonment	129
8.10.2	Patterns of Play within Cliques of Players	129
8.10.3	Additional Game Features Driving Engagement	130
8.10.4	Larger-scale Data Collection	131
8.11	Discussion	131
9	Closing Discussion	133

Todo list

1 | Introduction

1.1 A primer on aspect orientation

Explain aspect orientation briefly. Some useful notes for the explanation:

- AO originated in xerox parc, first described in [Kic+97]. There are lots of weaving mechanisms for regular, static aspect orientation, and there's a good early survey of them all (and implementation in a custom OO language specifically for this!) in [MK03].
- AO has some forebears: metaobject protocols, subject-oriented programming, adaptive programming, composition filters. The latter three are described by [CBB19] as being alternative kinds of AO — I disagree, but they're certainly attempting similar things.
- The original and still most widely used AO implementation is AspectJ, which comes with its own aspect language. It's grown over the years and is used sometimes in industry [citation needed...]. A smaller alternative would be Spring AOP [find a citation for spring AOP](#)

1.2 Prior Work

Write a section for the introduction describing the work done on pdsf before this, to delineate where we're starting from and avoid any claims of plagiarism. This can be short, the first sec of the lit review is a proper discussion, but the tool should be mentioned here. See ?? for what already exists.

1.3 Terms & definitions

Complete the glossary in section 1.3.

Decide whether terms like BPMN, simulation & modelling, etc also belong in the glossarysection 1.3.

Aspect

Advice

Joinpoint

Pointcut

Weaving

AspectJ The original aspect orientation framework, with language extensions to describe pointcuts and aspects.

Target The procedure an aspect is applied to via a join point, to affect advice.

PyDySoFu

2 | Relevant Literature

This thesis presents an aspect-oriented approach to simulation and experimental design, tooling to support these endeavours, some empirical assessment of both the paradigm and its application in the domain of simulation and modelling. In particular, the presented research makes use of aspects to model hypotheses and the complexities of variable and unpredictable behaviour in the simulation of socio-technical systems.

Relevant literature in this topic typically comes from a variety of fields which do not overlap significantly, meaning that this review must cover segments of several partly related fields. The presented material is eclectic in nature as a result. We therefore present context for unfamiliar readers as well as a motivating case for the work to follow in three distinct areas, which are, in order :

Section 2.1, which introduces aspect orientation and gives a background on the field;

Section 2.2, which details some existing approaches to the dynamic weaving of aspects;

Section 2.3, which discusses the use of aspect orientation in simulation and modelling; and

Section 2.4, which outlines literature on the modelling of process variance, particularly in simulation and modelling.

don't use a
description
list here, this
should be
a regular
paragraph

2.1 Aspect Orientation

Often, software engineers have to repetitively handle an issue in some codebase, where the issue is pervasive across many parts of the codebase and is necessarily interwoven through its core functionality.

Common examples of this are guarding against unsafe concurrency usage, memory management, and logging. Modularity is considered a key trait of maintainable, flexible and legible programs [Par72]. Modern design techniques often centre around the structure of a program to increase its modularity, with object-oriented programming being the standard approach to designing with modularity in mind in many industrially-relevant languages today. Citation needed?

Approaches to modularity typically section codebases into units of functionality. Concerns such as logging and memory management happen in effectively all areas of a codebase, as a result of engineering needs and the properties of a project’s language and environment. As a result, common devices employed with the aim of increasing modularity are unable to strip these “cross-cutting concerns” into some separate, modular unit. Programmers separating these concerns into additional modules are expected to see two key benefits:

1. A reduction of “tangling”, where program logic essential for a program’s intended purpose is intermixed with ancillary code addressing cross-cutting concerns, thereby making essential logic more difficult to maintain;
2. A reduction of “scattering”, where program logic for cross-cutting concerns is strewn throughout a codebase, making maintenance of this code more difficult.

The existence of cross-cutting concerns therefore expected to make maintenance of both ancillary program logic and a program’s core logic more difficult. Addressing this, Kiczales et al. introduced the notion of aspect-oriented programming [Kic+97]. Aspects are best described through their component parts:

- A “join point” defines some point in a program’s execution (usually the moment of invocation or return of some function or method).
- “Advice” defines some behaviour, such as emitting a logline, which can conceptually happen anywhere in program execution (i.e. what’s defined would typically represent behaviour which cuts across many parts of a codebase).
- An “aspect” is constructed by composing this advice with “point cuts”: sets of join points that define all moments in program execution where associated advice is intended to be invoked.

- An “aspect weaver” then adds the functionality defined by each aspect by adding the functionality defined by its advice at each join point defined by its point cut.

The definition of join points and advice, or how weaving occurs, is a matter left for aspect orientation frameworks and languages to define. In employing the technique, aspect oriented programming aims to separate cross-cutting concerns into aspects, removing the aforementioned repetitive code from the logic implementing a program’s functional behaviour so that additional pieces of functionality — logging, authentication, and so on — can be maintained in only one place in a codebase (thereby simplifying their maintenance and comprehension), and remaining program logic can be understood and maintained without the overhead imposed by the previously tangled cross-cutting concerns.

In [Kic+97], Kiczales et al. see these engineering concepts as universal throughout business logic, motivating the aspect-oriented approach for the first time. The authors present an implementation of AOP in Lisp, and compare implementations by way of e.g. SLOC count in an emitted C program to a comparable, non-AOP implementation, with two examples (its use in image processing and document processing). They find the idea — which they note is “young” and note many areas where research might help it to grow — can successfully separate systemic implementation concerns such as memory management in a way that reduces program bloat and simplifies implementation. It is noted that measuring the benefits of their approach quantitatively is challenging.

Tooling followed the theoretical work presented by Kiczales et al. [Kic+97] with a demonstration and subsequent technical description of AspectJ, a Java extension for aspect oriented programming [Hil+00; Kic+01]. AspectJ was introduced to satisfy the research community’s need for a tool with which to demonstrate the aspect-oriented paradigm in case studies. The tool is intended to serve as “the basis of an empirical assessment of aspect-oriented programming” [Kic+01]. The library makes use of standard aspect-oriented concepts: Pointcuts, Join Points, and Advice, bundled together in Aspects. They define “dynamic” and “static” cross-cutting, by which they refer to join points at specific points in the execution of a program, and join points describing specific types whose functionality is to be altered in some way. Their paper describes only “dynamic” cross-cutting, but presents tooling support, architectural detail of its implementation, and the representation & definition of pointcuts in AspectJ. AspectJ is compared to other tools for aspect orientation and related decompositional paradigms, and

the authors are explicit about their approach being distinct from metaprogramming in, for example, Smalltalk or Clojure.

Filman, Friedman, and Norvig isolate properties of aspect orientation which they assert are definitive of the paradigm [FFN00]. Specifically, they claim that aspect oriented programming should be considered “quantification” and “obliviousness”:

“A”OP can be understood as the desire to make quantified statements about the behaviour of programs, and to have these quantifications hold over programs written by oblivious programmers. [...] We want to be able to say, “This code realises his concern. Execute it whenever these circumstances hold.

These concepts, alongside “tangling” and “scattering”, became core concepts in aspect orientation literature. This is in spite of Filman, Friedman, and Norvig giving no concrete definition of the terms in their original paper (nor citing a source for definition). For the purposes of this thesis, we therefore provide the following definitions of the terms:

Quantification is the property of specifying specific points in a program in which that program should change;

Obliviousness is the property of a codebase that it contains no lexical or conceptual reference to advice which might be applied to it, and of the programmer of a target program that their code may be amended by an aspect programmer.

Filman, Friedman, and Norvig write about aspect orientation “qua programming language”, and so theorise around aspect orientation as a paradigm independent of a particular instantiation. They are therefore able to arrive at conclusions about the paradigm in the abstract, and can identify concerns for future investigation for researchers in the field and design goals for developers of aspect orientation tooling. They note:

“B”etter AOP systems are more oblivious. They minimize the degree to which programmers (particularly the programmers of the primary functionality) have to change their behaviour to realise the benefits of AOP. It’s a really nice bumper sticker to be able to say, “Just program like you always do, and we’ll be able to add the aspects later.” (And change your mind downstream about your policies, and we’ll painlessly transform your code for that, too.)

Whether obviously designed aspect-oriented systems achieve their intended goals empirically is outside the scope of their work, and the lack of empirical evidence for this is discussed in ???. Claims made such as changing one’s mind while developing or maintaining a program and having that “painlessly transformed” — an effect of the aforementioned programmer’s obliviousness — is incompatible with earlier writing on modularity. Yourdon and Constantine assert [YC79]:

“The more that we must know of module B in order to understand module A, the more closely connected A is to B. The fact that we must know something about another module is a priori evidence of some degree of interconnection even if the form of the interconnection is not known.

Aspect orientation’s critics describe similar incompatibilities with existing best-practices [Prz10; CSS04], as well as the lack of empirical evidence for the benefits of obliviousness [Ste06]. Claims about “better” aspect-oriented systems being more oblivious should therefore be regarded as suggestions from the literature, and while obliviousness and quantification are useful concepts in discussing research in the field. They also give context for the research community’s perspective that obliviousness and quantification are design goals for aspect oriented programmers [AspectCplusplusDesignImp; aspect oriented workflow; Kel08] (though other researchers suggest they may be best applied in moderation [LC07]).

2.1.1 Alternative methods in Aspect Orientation

Aspect-oriented programming’s goal of modularising cross-cutting concerns is shared by other paradigms. Seminal work in aspect orientation makes note of similarities to use of reflection, metaprogramming & program transformation, and subject-oriented programming [Kic+97; Kic+01]. They also observe that other disciplines have introduced “aspectual decomposition” independently.

The example of pre-existing aspectual decomposition by way of diagramming given by Kiczales et al. [Kic+97] is in physical engineering. To give a concrete example from their description, differing types of diagrams when engineering a system such as thermal and electric diagrams of a heater are described as “aspectual” because of the modular nature of the diagrams; though there might be many diagrams of different kinds, they compose together to give an overview of the system being designed.

Similar diagramming techniques have independently arisen in other domains since. The Obashi dataflow modelling methodology[WC10] by Wallis and Cloughley models all possible paths of dataflow through “B&IT” (business and IT) diagrams, where business-specific concerns (people, locations, groups, and business processes such as payroll, stock-check or budgeting) are modelled alongside IT concerns such as applications supporting business processes and the software and hardware infrastructure supporting them. Modelling dataflows in this way allows for a comprehensive understanding of assets and business processes. However, in order to understand how data flows between specific assets within a B&IT, sub-graphs (“DAVs”, or Dataflow Analysis Views) denote specific pathways through which data flows between source and sink assets. Alternatively, a B&IT can be viewed as a composition of all possible DAVs within an organisation. Dataflows are therefore broken into different diagramming techniques and specific business concerns can be described independently of others, even if these concerns interact in their dataflow pathways (and, therefore, cutting across each other). Obashi therefore allows for the aspectual decomposition of business processes, through the description of an organisation by individual dataflow analysis views, which compose into an overall model of a system in a B&IT diagram. Obashi models are an instance of aspect orientation which were designed for simplicity and comprehension[WC10; Seo11], but trade this for domain-specificity.

Research conducted by Keller and Hölzle [KH98] investigates solutions to the difficulties involved in the integration of software components and their evolution over time, where those components are re-used with differing requirements. By modifying binaries directly, incompatibilities in a program and one of that program’s dependencies can be resolved by way of mutating either after compilation. Their implementation defines a representation for the modification of pre-generated Java class binaries, the output of which can be verified as also being valid Java class binaries. Keller and Hölzle claim that BCA allows for dynamic modification of programs with little overhead. They believe BCA is unique in its combination of features, which include engineering concerns such as typechecking code which is subject to adaptation and its obliviousness to source implementation, as well as guarantees that modifications are valid even for later iterations of the program subject to adaptation.

Other engineering techniques can be used to modularise cross-cutting concerns. For example, metaobject protocols describe the properties of an object’s class (including, for example, its position within a class hierarchy) in an adaptable manner [KDB91]. These can be used to implement aspect

orientation [Esp04], therefore providing at a minimum the same functionality, though they achieve this through their reflective qualities and are designed with metaprogramming as a primary goal as opposed to modularisation of cross-cutting concerns [KDB91; Sul01]. Multiple-dispatch, where methods on objects are chosen to be run based on the properties of the parameters passed at point of invocation, allows for oblivious decomposition without the need for a weaver [DGM08], although this does not support the goals of aspect orientation in totality. For example, a programmer might want their program to exhibit differing behaviour when methods are called with differently-typed arguments, which is supported by multiple dispatch. However, they might instead want their program to exhibit some additional behaviour whenever a method is invoked, such as logging, but might not want to implement logging alongside the rest of their method implementation for clarity or length reasons. Multiple dispatch therefore offers comparable but different functionality to a software engineer. Engineering patterns such as decorators provide similar functionality to aspects [FBS17], in that cross-cutting concerns can be separated into their own module, but they differ in their approach to obliviousness: decorators annotate areas of a codebase they are applied to, and therefore do not offer obliviousness as aspects do. Add a paragraph after this on subject-oriented programming.

More notes Add notes on subject-oriented programming, bigraphs, maybe holons?

2.1.2 Criticisms of aspect-oriented programming

subsec:aop-criticisms

The growing aspect-oriented programming research community collected both proponents and detractors of the paradigm. The developments in aspect orientation pertinent to the research presented in this thesis are discussed in later sections of this literature review. However, common criticisms of aspect-oriented programming are important to present in two regards:

1. Discussions of advancements in the field pertinent to the work presented in this thesis should be understood within the context of some perceived weaknesses in the field, which helps to frame an understanding of literature reviewed in this chapter,
2. The presented work addresses some criticisms of aspect-oriented programming, meaning that

the criticisms of the paradigm writ large and properties of work published in awareness of those weaknesses will motivate some research presented in later chapters.

An early piece of scepticism in the aspect orientation community is [CSS04], in which Constantinides, Skotiniotis, and Stoerzer see AOP’s core concepts as having significant similarities to GO-TO statements, which have historically been the subject of some derision in the literature. [CSS04] is, in spirit, a child of Dijkstra’s “Letters to the editor: go to statement considered harmful”. The authors note that the notion of unstructured control flow makes reasoning about a program complicated — disorientating a programmer by way of “destroying their coordinate system”, leaving them unsure about both a program’s flow of execution and the states at different points of that flow — and discuss whether aspect orientated programs can have a consistent “coordinate system” for developers. They note that, while Go To statements are at least visible in disrupted code, the AOP concept of obliviousness makes such reasoning even more difficult than Go To statements, as even the understanding of where and how flow is interrupted is not represented structurally within an aspect-oriented program. They compare aspects to a Come From statement, noting that the concept is a literal April Fools’ joke for programming language enthusiasts who claim they’ve found an improvement over Go To statements. The authors conclude that existing techniques, specifically Dynamic Dispatch in OOP, provide similar benefits without the trade-off in legibility of a program’s intended execution.

A similar and more thorough critique of the aspect oriented paradigm can be found in Steimann’s “The paradoxical success of aspect-oriented programming” [Ste06]. The concern in this paper is that the popularity of aspect oriented programming — which was nearly 10 years old at time of publication — was founded on a perception that it assisted in engineering more than it was proof that such assistance viable in practice. The author notes that most papers are theoretical in their discussion on tooling, that examples were typically repetitive, and that the community’s discussion concerned more what aspect orientation is good for than what it actually is in practice. AOP is compared against OOP, AOP’s claimed properties and principles are examined in detail, and the impact on software engineering is reasoned about from a sceptical perspective, comparing claims such as improved modularity against classic papers on the subjects (such as Parnas’ work on the same). The paper presents a philosophical examination of aspect orientation, assessing the paradigm against its purported merits and discussing whether we should expect, rationally, that the claims made by the AOP research community would

hold true. The paper ends noting some benefits of AOP that do hold true under rational scrutiny, and notes that the true utility of AOP may be very different to those purported by the community. Overall, the paper is a philosophical and critical reflection on the state of AOP research and the community’s zeitgeist at the time, claims around which are not well-evidenced in literature. In particular, the author sees AOP’s promise of unprecedented modularisation as unfulfillable.

Similar sentiments are shared by Przybylek [Prz10], who looks to examine aspect oriented programming within the context of language designers’ quest to achieve maintainable modularity in system design. They frame the design goals of aspect orientation as being to represent issues that “cannot be represented as first-class entities in the adopted language”. The paper discusses whether the modularity offered by aspect orientation actually makes code more modular. In particular, they distinguish between lexical separation of concerns and the separation of concerns originally discussed by Dijkstra and Dijkstra in “On the role of scientific thought”. They assess the principles of modularity — modular reasoning, interface design, and a decrease in coupling, for example — and find that from a theoretical perspective, there are many reasons to believe that the aspect-oriented paradigm can detrimentally impact the expected benefits of proper modularisation in a program. They conclude that the benefits touted by AOP are a myth repeated often enough to be believed, but point to many papers which suggest improvements to the standard AOP approach which might reduce its negative impact or make it more practically viable. Przybylek presents a critical review of aspect orientation literature, but often hints at others’ solutions to the problems identified too.

2.2 Dynamic Aspect Weaving

One implementation of dynamic weaving is PROSE [PGA02; PAG03], a library which achieves dynamic weaving by use of a Just-In-Time compiler for Java. The authors saw aspect orientation as a solution to software’s increasing need for adaptivity: mobile devices, for example, could enable a required feature by applying an aspect as a kind of “hotfix”, thereby adapting over time to a user’s needs. Other uses of dynamic aspect orientation they identify are in the process of software development: as aspects are applied to a compiled, live product, the join points being used can be inspected by a developer to see whether the pointcut used is correct. If not, a developer could use dynamic weaving to

remove a misapplied aspect, rewrite the pointcut, and weave again without recompiling and relaunching their project.

Indeed, the conclusion Popovici, Alonso, and Gross provide in [PAG03] indicates that some performance issues may prevent dynamic aspect orientation from being useful in production software, but that it presented opportunities in a prototyping or debugging context.

PROSE explores dynamic weaving as it could apply in a development context, but the authors do not appear to have investigated dynamic weaving as it could apply to simulation contexts, or others where software making use of aspects does not constitute a product.

The performance issues noted by Popovici, Alonso, and Gross are explored in more detail by Chitchyan and Sommerville in [CS04]. Chitchyan and Sommerville present a review of early dynamic aspect orientation techniques. The paper reviews AspectWerkz, JBoss, Prose, and Nanning Aspects through the lens of the authors' prior work on dynamic reconfiguration of software systems and their generalisation of dynamic aspect orientation approaches:

1. "Total hook" weaving, where aspect hooks are woven at all possible points;
2. "Actual hook" weaving, where aspect hooks are woven where required;
3. "Collective" weaving, where aspects are woven directly into the executed code, "collecting the aspects and base in one unit".

Because of the paper's focus on software reconfiguration (rather than the mechanics and design of dynamic aspect weaving specifically), the analysis of the tools presented in the paper is of less relevance to the work presented in this thesis than their generalisation of dynamic weaving. The trade-offs of the three generalised philosophies are discussed. Chitchyan and Sommerville propose that total hook weaving allows flexibility in the evolution of a software product, at the expense of the performance of that product; this contrasts collected weaving, which shifts overhead out of the codebase and into the maintenance effort. Actual hook weaving is positioned as a compromise between the two, offering the best approach for none of their criteria but never compromising so much as to offer the worst, either. This suggests merit in a tool designed to flexibly offer any weaving approach appropriate for the task at hand. It's explicitly noted that, in practice, one could use many of the systems they describe. Though

the paper is an early publication in the field, no tool the authors review offers all three, and none offers collective weaving alongside either kind of hook weaving.

In contrast, Gilani and Spinczyk note that, while there are different approaches to dynamically weaving aspects, no approach is suitable for an embedded environment. This is due to these systems' low power and available memory. Gilani and Spinczyk therefore propose a framework through which weavers can be assessed for suitability in a given environment, or generated from a set of possible features (where, presumably, features would be enabled and disabled as per an environment's needs). Their families of weavers are defined by the similarities of the requirements in domains they are applied to, and specifically defined by their trade-off between dynamism and resource use (asserted to be broadly proportional). It is unclear that a carefully crafted "actual hook weaver", or JIT-compiled "collective weaving", in the parlance of Chitchyan and Sommerville (see [CS04]), would be meaningfully less efficient than static weaving in all but the extreme application areas outlined in the paper (embedded systems with resources in the range of 30kb memory).

Aspect oriented programming's criticism can often be that it doesn't know what it "aims to be good for", and its application in such extreme environments is arguably mistaken from the off. The families outlined in Gilani and Spinczyk's publication are unnecessary if dynamic aspects are not required in their target environments. Steimann's critique of aspect-oriented programming, contrasted against these families, presents an interesting question. If the goals of dynamism and resource efficiency are at odds, and Steimann's stance that aspect-oriented programs do not earn its proponents' plaudits in practice **don't be fancy here**, what can dynamic aspect weaving be appropriately applied to? In what environment does the trade-off presented by dynamic weaving not necessitate a theory like Gilani and Spinczyk's in the first place? Arguably, that environment is not found in low-resource systems, and a take-away of [GS04] could be that researchers should seek other contexts in which to apply aspect oriented programming. ¹ **Note from discussion with Tim to work into this: "the anticipated benefits of aspect-oriented programs are not observed in practice" Both paragraphs on Gilani2004 need to be heavily edited.**

et al. propose a new aspect-oriented invocation mechanism, which they call "Bind" [+06]. Bind's design is motivated by perceived opportunities to improve modularity from a design perspective.

¹As discussed alongside [GK99] in section 2.3, simulation & modelling might present a more appropriate field.

The impact of “scattering” and “tangling” in a codebase after weaving in some aspect orientation implementations leads to a more complicated post-weave codebase, which in turn leads to increased difficulty including the compilation of aspect-oriented code and the development and execution of unit tests on said code. In order to demonstrate Bind’s approach to simplifying post-weave codebases, the design of “Nu”, an aspect orientation framework in .NET supporting Bind, is explained and an implementation presented. They present Bind as an alternative to the weaving of aspect hooks (for load-time registration) into target code, in the style of PROSE (see [PGA02; PAG03]), or the weaving of calls directly, in the style of AspectJ (see [Kic+01]). Bind’s model for aspect orientation is to apply or remove implementations of cross-cutting concerns to arbitrary sets of join points at a time of a developer’s choosing. Nu’s model for this is designed with the aim of granularity of join point specification. What results is a flexible model for aspect orientation which is demonstrated to satisfactorily emulate many other models for aspect oriented programming, such as the models of AspectJ, HyperJ, and Adaptive Programming. It is noted that it is “very common in aspect-oriented programming research literature to provide language extensions to support new properties of aspect-like constructs”, and note that their work is similar to (yet distinct from) weaving approaches in run-time & load-time weaving, support for aspect orientation directly in a language’s virtual machine, and work towards general models of aspect oriented programming (models which can represent a variety of existing approaches). Their approach is flexible and considered enough to warrant impact in the introduction of aspect orientation within virtual machines (because they require no direct support), and in their ability to represent different weaving approaches, arguably because their approach is general enough in design to approach the general model worked towards, which qualifies their satisfaction of their motivation to provide a model distinct to the approaches initially discussed. In line with the complaints of AOP’s critics, this does not seem to qualify the satisfaction with which they achieve their practical engineering goals.

Relatedly, Dyer and Rajan explain in more depth than in the design and implementation of the Bind mechanism and the implementation of the Nu framework [DR10]. A more technical discussion is presented, in particular on implementation details including optimisation and benchmarking, largely against AspectJ. Notably, the implementation discussed is a Java implementation, rather than the .Net implementation presented in [rajan2006nu]. Many aspect orientation frameworks are language-specific;

the existence of Nu’s implementation on multiple platforms highlights the work’s most interesting facet being the design of the Bind primitive, rather than the framework itself. In a research area where tooling papers are common but the lack of design philosophy & analysis of case studies is frequent fodder for critics, the novelty of the Bind mechanism distinguishes this series of papers.

2.3 Aspect Orientation in Simulation & Modelling

Aspect orientation as applied to simulating systems, and building models of systems, has been researched from several approaches.

A very early example of aspect orientation in simulation & modelling is presented by Gulyás and Kozsik [GK99]. Gulyás and Kozsik observes that, in the study of complex systems through software models, the software developed typically serves two purposes: the experimental subject, and the observational apparatus used to conduct the experiment itself. Arguing that the separation of these roles ought to make both the implementation of an experimental system and its later analysis simpler, Gulyás and Kozsik proposes the use of aspect orientation as a means of separating what they perceive as cross-cutting concerns of systems modelling. They present their Multi-Agent Modelling Language, a language implemented in Objective-C via the Swarm simulation package and designed for aspect-oriented simulation of agent-based models. Their aspect orientation effectively makes use of Observer patterns to measure a pre-constructed system under simulation, without the observations being an intrinsic component of the simulated system. They find that AOP provides an intuitive and straightforward method by which simulated experimental systems can be composed, and that MAML’s simplicity and its philosophy on modelling are more “satisfactory” than Swarm’s standard approach, though the paper betrays that its implementation was more complex than initially conceived: the patch unix tool was intended for use as their weaver, though the team eventually developed a transpiler from MAML to Swarm instead (which they name xmc.). The deciding factors for the development of a custom transpiler are not discussed.

Gulyás and Kozsik’s work presents not only tooling for aspect oriented simulation, but some reasoning & philosophy on the potential benefit of using aspect orientation in these endeavours that extends further than the conclusions of modularity through separation of concerns and a reduction

of tangling & scattering. In particular, their work discusses specific scenarios in which the type of separation of concerns offered by aspect orientation is desirable, and the engineering approach to achieving the aim reasonable. This distinguishes the work in comparison to many aspect orientation papers reviewed in this chapter. Many papers describe the expected benefits by simply drawing from existing literature such as [Kic+97]. The fact that a rare example of detailed reasoning about the appropriateness of aspect oriented programming in a particular domain is highlighted because the domain in question is simulation & modelling; the subject of this thesis. That aspect orientation might be well suited to separating observer and experiment motivates, in part, this thesis’ work showing the plausible realism of a simulation in which behaviour is modified by aspects. Put another way: this thesis draws on the idea that, in response to Steimann asking what aspect orientation is good for, Gulyás and Kozsik would seem to answer, “simulation & modelling”, a premise this thesis shares.

Chibani, Belattar, and Bourouis discuss two issues in object-oriented programming: “tangling”, where separate design elements of a program are woven within each other in program source, and “scattering”, where a single design element is strewn throughout the source, rather than being contained within a single area of the codebase. **Remove the description of tangling and scattering from here and include in a review of more foundational AOP literature early in the lit review. They don’t belong buried so far down in this chapter.** They propose that aspect orientation solves these problems, and identify that there are potential benefits in discrete event simulation code in both regards, making DES frameworks with aspect oriented primitives a potentially fruitful contribution to the research community. Chibani, Belattar, and Bourouis identify cross-cutting concerns in DES codebases, including event handling, resource sharing, and the restoration of a simulation run. The contribution of the paper is the discussion of AOP’s potential application to DES codebases, and detail of the avenues available for research in the field. Japrosim is presented as a motivating example of an existing DES framework which they see as ripe for the aspect oriented enhancements they identify.

In a later publication [CBB19], Chibani, Belattar, and Bourouis identify opportunities for the use of aspect orientation in simulation tooling, aiming to increase “modularity, understandability, maintainability, reusability, and testability” by applying the paradigm [CBB19]. They present a case study of an application of aspect orientation to simulation tooling by identifying cross-cutting concerns in Japrosim, a discrete event simulation framework, and propose an aspect-oriented redesign

of the tool using AspectJ. Chibani, Belattar, and Bourouis describe Japrosim’s existing object-oriented design, followed by aspect oriented variations of some design elements, including concurrent process management and in Japrosim’s graphical animation features. A similar survey of areas in which Japrosim’s source might benefit from the application of aspect orientation is presented by Chibani, Belattar, and Bourouis in an earlier work [CBB14]. In both cases, the main contribution noted is the design itself. Counting the main improvements between the presented aspect-oriented design and the existing object-oriented one is left to future work in the authors’ later publication [CBB19], although a concrete implementation is linked to and some quantitative evaluation of that implementation presented in their earlier publication [CBB14]. The quantitative evaluation provides measurements based on Martin’s object-oriented design metrics and demonstrates a greater independence of packages in their aspect oriented version of Japrosim than in the original. However, the intended aim of aspect orientation is not to decouple existing packages, but to isolate those packages’ cross-cutting concerns into new ones. It is therefore unclear that their quantitative evaluation achieves its improvements as a result of aspect orientation. No further discussion of their results is provided, and it is possible that the improvement is due to the rewriting necessary in their maintenance of the Japrosim source, rather than due to their use of aspect orientation specifically.

In a manner similar to Chibani, Belattar, and Bourouis’s [CBB19; CBB13; CBB14], Aksu, Belet, and Zdemir observe that there are opportunities to be found in a simulation framework able to take advantage of aspect orientation [ABZ08]. Examining the DEVS framework Simkit, their proposal for aspect-oriented programming adoption is two-fold: refactoring of the framework itself and aspect-oriented tooling for use by modellers, who represent cross-cutting concerns within their models. Opportunities for improvements in production and development are discussed, and some implementation notes are detailed, although no concrete implementation or evaluation is provided; the work instead proposes design alterations, and the authors “leave it as a future work [sic] to explore the usability and efficiency” of aspect orientation used idiomatically alongside Java’s existing reflection offerings. The existence of multiple attempts to refactor differing simulation packages with aspect orientation indicates potential for modellers in the use of aspect-oriented patterns, but the real-world utility of the techniques are omitted. Chibani, Belattar, and Bourouis and Aksu, Belet, and Zdemir both seem to defer to the zeitgeist wisdom of the aspect orientation community in their unproven claim that it improves

modularity and maintainability of a codebase. ²

Neither Gulyás and Kozsik nor Aksu, Belet, and Zdemir detail a case study of their techniques with real-world examples. However, Ionescu et al. do in their work identifying an increased demand for computational power in simulation execution on supercomputers [Ion+09]. Existing known-good models might be unsuitable for the extreme requirements of code efficiency modellers contend with, but running the code on different environments requires modifications for suitability in different environments, around which there are regulations and risks of a reduction in quality during maintenance. The authors propose an aspect-oriented solution to the problem, where aspects modify the simulation codebase with minimal overhead. An implementation of a real-world model for disaster prevention is presented, and assessed both by comparison against an equivalent non-aspect-oriented codebase and by assessment of the aspect-oriented variant’s scalability and reliability in both cluster and multi-cluster environments. They find that a comparative analysis of generated code and of their simulations in various configurations both indicate that their simulation’s aspect-oriented implementation is suitable for use in disaster prevention, implying that aspect orientation could be suitable in scenarios with comparable requirements.

That the authors can conclude that aspect orientation is suitable in a real-world use case constrained by the requirements of supercomputer use seems promising for the aspect-oriented paradigm, which is criticised for its lack of practical evaluation [CSS04; Ste06; Prz10]. As it is therefore a rare example of aspect-oriented case studies, their evaluation methodology is important to highlight. Their code analysis makes use of significant lines of code as a core metric, which doesn’t reliably reflect code quality; quoting Rosenberg’s “Some misconceptions about lines of code” [Ros97]:

“(...) the best use of SLOC is not as a predictor of quality itself (for such a prediction would simply reduce to a claim about size, not quality), but rather as a covariate adjusting for size in using another metric.”

This is important because a common claim in the aspect-oriented literature for which there is little empirical evidence is that aspect orientation improves the “quality” of a codebase, but the related claims made by Ionescu et al. [Ion+09] are unreliable.

²Chibani, Belattar, and Bourouis do present some quantitative evaluation, but this is flawed as previously described.

It is important to note that improvements in code quality specifically are those which have come under scrutiny by the critical papers reviewed in ???. The results presented by Ionescu et al. do not satisfactorily address the requests for empirical evidence of improved code quality in these reviewed criticisms. This does not impact their aspect-oriented models' suitability given the motivations of Ionescu et al., which are that models should be amendable for new supercomputing settings without lack of performance. The models described in this work satisfy that aim. These models are also evaluated by way of performance, an important factor in supercomputer use where execution time is financially expensive and power-intensive. Quantitative evaluation of their simulation's execution time shows less than 5% slowdown compared to a non-aspect-oriented implementation. Ionescu et al. deem this a reasonable trade-off for the engineering improvements they observe. Their application to the supercomputing & disaster prevention simulation domains therefore seem satisfactory by way of performance, and meet Ionescu et al.'s aim of demonstrating a modelling technique which permits adapting existing models for use in new environments without directly maintaining the original model's source.

This result is notable with regards the findings presented in this thesis, which similarly aim to alter a pre-existing model without directly altering it, although for purposes such as model reuse and simplification of experimental design rather than for avoiding the regulatory overhead and financial cost of maintaining the models described by Ionescu et al. [Ion+09].

2.3.1 Aspect Orientation & Business Process Modelling

Several projects within the business process modelling research community make use of aspect orientation to design modelling languages which produce less monolithic business process models [Cap+09; Sil+20] and simplify the composition of models [CM07]. As business processes are inherently sociotechnical and this thesis presents tooling for and results in the modelling and simulation of socio-technical systems using aspect-oriented techniques, it is important to review this community's literature. **Editing tip from Tim: don't call things "important", this implies other things aren't. Rework areas where I do this.** This field is particularly relevant as work on this project prior to this thesis' research models software engineering processes that are conceptually similar to business process modelling (see chapter 3), and there also exists interest in modelling behavioural variance within the business

process modelling community (see section 2.4), which is highly relevant to this thesis’ concern with the representation of alteration to process and modelled behaviour as aspects.

Charfi and Mezini see opportunities in integrating BPEL, an executable business-process modelling language, with aspect-oriented concepts [CM07]. This is because when BPEL systems are composed together the static nature of the logic being composed is not always appropriate for BPEL’s use cases. The specific use-case examined is web service definitions, where changes affecting composition of multiple component parts can affect many areas of a final result, making modification error-prone. They specifically seek to support dynamic workflow definitions — “adaptive workflows” — which BPEL’s existing extension mechanisms do not sufficiently support, but the aspect-oriented literature discusses at length (an overview of which is presented in section 2.2). Therefore, they look to construct an aspect-oriented BPEL extension. Using the case study of modelling a travel agency’s web services, they create an aspect-oriented extension by first defining how such an extension would be represented graphically in BPEL’s workflow diagrams. Further detail is added to arrive at a technical definition with XML representations, weaving mechanics, and eventually the construction of a BPEL dialect, AO4BPEL. The authors find that their pointcut system (which describes join points on both processes and BPEL messages), support for adaptive workflows, and aspect-oriented approach to workflow process modelling make AO4BPEL unique at the time of publication, though related AOP implementations exist in each individual area of their contributions. The work is weakened by brittle semantics around pointcuts, join points, and the temporal nature of workflow modelling. For example, they note that defining contingent behaviour — only applying an aspect conditionally, based on a trace through a simulation of a modelled system — would allow the application of advice only when model state deems this appropriate.³ They also call for more generally theoretical AOP research, which mirrors the requests some critics of aspect orientation research make (as noted in ??).

In a PhD thesis describing AO4BPEL in detail [CM06] Charfi and Mezini presented a generalisation of the notation developed for AO4BPEL, which applies to any graphical workflow modelling language. Accompanying this are some examples of its use building a framework for enforcing certain requirements of BPEL models, and use of that framework to develop aspect-oriented frameworks for enforcing security and reliability within AO4BPEL models.

³The contingent application of model adaptation is a motivating case for some work presented in this thesis; see chapter 3 for a discussion.

In later work, Charfi, Müller, and Mezini define a similar aspect-oriented dialect of BPMN they name AO4BPMN [CMM10], after asserting that the concerns addressed by AO4BPEL [CM06; CM07] in the field of executable process languages also apply to business-process modelling languages, and can be solved similarly. The generalised notation of aspectual workflow models presented in Charfi and Mezini’s thesis [CM06] are applied to BPMN to produce an aspect-oriented language specifically for process modelling, as opposed to executable business process modelling.

Cappelli et al. also note that cross-cutting concerns exist in business process models, and are specifically motivated by monolithic design approaches common in business process modelling languages. Like Kiczales et al., they claim that a lack of modularity in business process models leads to cross-cutting concerns scattered throughout a model [Cap+09]. To alleviate the issue, they propose a meta-language, AOPML, which incorporates aspect orientation in a metamodel of business process modelling languages, and instantiate it within their own dialect of BPMN. Using a model of a steering committee as a case study, and separating cross-cutting concerns such as logging, the paper proposes reducing complexity and repetition graphically, thereby in a manner more in keeping with the language design philosophies of popular business process modelling languages, the design and use of which are typically graphical [Obj11; DOR95; Obj15]. They note that this is in contrast to other applications of aspect orientation in business process modelling — specifically AO4BPMN — where aspect definitions are written in XML concern not only the advice to be applied but also their relevant join points, as in general programming aspect orientation implementations such as AspectJ. In this way, the AOPML exhibits the spirit of business process modelling more stringently than does Charfi and Mezini’s notation for aspect-oriented workflow modelling.

The difference between Charfi, Müller, and Mezini’s approach in designing AO4BPMN [CMM10] and Cappelli et al.’s approach in designing AOPML [Cap+09] highlights design decisions taken when introducing aspect orientation in a new domain. There is an opportunity for a domain-specific aspect orientation framework to align its design with the traditions and idioms already present in models within that domain, but doing so may break the traditions and idioms which already exist in aspect-oriented approaches in other domains. Comparing the approaches of Charfi, Müller, and Mezini and Cappelli et al. does surface that there may be no clear “best” design approach when blending pre-existing modelling paradigms, such as the graphical modelling languages used in business-process modelling

and the abstract concepts of aspect orientation. The discussion around whether it is more desirable to adapt existing design elements of aspect-oriented frameworks to a given domain or adapt that domain's existing modelling traditions and idioms to incorporate aspect orientation as it is used elsewhere is outside the scope of this thesis.

New concepts within the design of aspect orientation frameworks are addressed in the business process modelling community. Jalali, Wohed, and Ouyang note that aspect oriented modelling frameworks often do not explicitly model the precedence of aspect application [JWO12]. They address this limitation by defining a mechanism to be used in capturing multiple concerns as aspects, where the invocation of advice must follow a certain precedence. The aim of the work is not to propose tooling around the precedence of aspect application so much as to contribute to aspect oriented design theory, providing a notation for precedence which is broadly applicable. The precedence model is, put simply, that a mapping exists for each application of advice to join point such that the mapping defines an ordering on advice for that join point. The definition defines “AOBPMN”, a formalised dialect of BPMN supporting aspect orientation with precedence. A case study is provided where AOBPMN is instantiated within a coloured Petri net. Their study expands on existing work by research teams led by Capelli [Cap+09; Sil+20] and by Charfi [CM07], in that it develops a mature formalism for and model of aspect orientation as applied to business process modelling. However, Jalali, Wohed, and Ouyang note that their case study is limited in scale. No tooling or evaluation of the practical benefit of their approach is provided.

2.4 Process Variance in Simulation & Data Generation

Brief intro of the section here [Brief intro of the process variance section here](#)

Typically, a simulation concerns a single process. This means that all expected behaviour must be included within that process; complex or contingent behaviour must be represented within it. The techniques reviewed here offer separately including some modification of a process (or represent the modifications of varied process within the simulated output [SA13; SA14]). The benefit of this approach is that possible changes to a process can be described once and applied to that process where appropriate. Process changes might describe attempts to circumvent security protocols, laziness or

confusion in a human actor within the model, or random “noise” so as to produce synthetic log traces containing aberrations which mimic those found in noisy empirical datasets. In all cases, behavioural variations can be described as some alteration to a process and applied to either a model or the product of that model (datasets or log traces) to represent the same alteration introduced at an arbitrary point of the simulation.

This decouples the expected behaviour in the original model from simulated behaviour, which is obtained by composing the model and behavioural variation using a given technique’s method for doing so. This approach to modelling behavioural variation allows the same altered behaviour, which would otherwise be described in many disparate points in a model, to instead be written once and introduced wherever required. The observation that the same variation might appear in many areas of a model, and that the variation can be separated from the model and introduced where necessary, frames the modelling of these variations in the same way as aspect orientation frames cross-cutting concerns. The work presented in this thesis explicitly applies changes to processes and simulated behaviour as aspects in the same manner. Therefore, although this aspectual connection is not made explicit in much of the literature to date, it is important to review literature on simulation and modelling which modularises these variations; this section reviews that literature. The work reviewed is highly relevant to the contributions in this thesis, in particular because the core motivations of this field are shared by this thesis; the section therefore leads with a subsection discussing those motivations, and their relationship to aspect orientation, in detail.

2.4.1 Discussion of Variation & Motivations for Variations in Process Models

Mitsyuk et al. are motivated by the field of process mining’s requirement for datasets of process logs made from well-understood process models, defined in a high-level manner [Mit+17]. They demonstrate a technique for generating event logs from BPMN models by introducing algorithms for the direct simulation of BPMN models and the collection of traces from those simulations. While their approach does not support the simulation of all BPMN concepts, notably message passing, they provide a tool which produces log traces for a BPMN model through PROM, a standard tool within the process mining community [Van+05]. This results in their technique providing high-level model simulation through already-standard tooling, meaning adopters of the technique need not rely on dedicated tooling

which may not be compatible with other researchers’ process mining techniques.

The algorithms presented by Mitsyuk et al. simulate processes described by BPMN models, but don’t include any provisions for representing variance. However, the technique could plausibly be combined with aspect orientation techniques for BPMN as discussed in ?? [CMM10; Cap+09] to represent alternate behaviour applied contingently. Demonstrating the viability of this approach is an avenue of research beyond the scope of this thesis. However, the motivation of the work mirrors that of other research projects reviewed in this section: a need for synthetic datasets of traces through a process, for use in scenarios where empirical datasets are difficult to obtain.

Difficulties arise when obtaining real-world datasets for many reasons. For example, large empirical datasets are typically produced by organisations which would prefer some level of secrecy around their operations, making publishing those operations for the investigation of research teams unlikely. Researchers collecting these datasets describe a “lengthy process” [Don20] and explain that traces of real-world processes are hard to obtain because “higher management [can be] worried about the risks” of publishing such datasets [Don20]. Another factor contributing to the difficulty of collecting empirical datasets is that they often cannot be collected, either because there is a need to study the process before implementing it (making synthetic datasets the only option available to researchers **Find citation for empirical datasets being the only ones available to researchers, maybe for disaster prevention or similar?**), because the process is not yet fully understood (making simulation of many variants of that process useful in aiding understanding **find a citation for simulating different systems for finding an optimal one. Arguably Genetic programming & hill climbing do similar things?**), or because the dataset itself is of use to researchers, not the real-world system that produced it (such as in the case of evaluating process mining techniques [VW04; AGL98]). **Find some nice way to round this subsection out, or refactor out subsections** Important to acknowledge somewhere in this subsection that synthetic data generation is a well-researched field, but that generating logs from simulations with variance is the specific area relevant to this thesis

2.4.2 Representing Variations in Process Models and their Outputs

Research undertaken by Stocker and Accorsi [SA13] aims to synthesise process logs which are representative of attackers' efforts to compromise the security of a modelled system. Their research project, named "SecSY", is an attempt to address issues arising from the difficulty of retrieving representative log traces for security-critical systems in which attacker activity is present. Logs are developed by process simulation through "well-structured" models, a mathematical property on which transformations were previously defined by Vanhatalo, Völzer, and Koehler [VVK09]. The authors develop a tool for the simulation of a process using well-structured process models, and apply transformations to both the model before execution and the log it produces through the trace of a simulation. They conclude that their tool is performant, and verify it can produce logs representing security violations by way of analysis through PROM, a popular framework for process mining, and pre-defined security constraints on their models. They note that log traces cannot be interleaved (due to a lack of parallel simulation of processes), may be incomplete (missing violations), and that mutated models and traces are not guaranteed to be sound by construction. However, they see their proposal as a necessary step in realistic data generation for business processes. A weakness of the work is that model and trace modifications are relatively rudimentary: processes can be added or removed, but complex graph transformations are presumably only permissible when representable through the composition of the mutation primitives they provide, on which there are only three for processes: swapping AND and XOR definitions of process gateways, and swapping process order. Mutations cannot be applied contingent on the state of a simulation run, for example, representing a decision taken by an attacker based on what had already happened. In later work, Stocker and Accorsi detail the technical aspects of SecSY, their tool for implementing the generation of synthetic logs which use their technique [SA13] to represent security violations security-critical business processes. A Java implementation of SecSY is described, which simulates well-structured models and applies mathematically-defined transformations on the model being simulated (before simulation occurs) and the logs obtained through simulation traces. An improvement on earlier work is that custom transformers can be written. However, a limitation of the original work remains, which is that users cannot easily dictate the degree to which variations are applied.

Pourmasoumi et al. also address the need for access to variations on business processes, though for the development of a research field, “cross-organisational process mining” [Pou+15]. Process mining can require many process logs, as does the benchmarking and evaluation of process mining techniques. Traces from business processes which are similar but not identical can produce log traces which reflect that similarity, but also reflect the variations in different instances of those processes. These log traces exhibiting variation can be used in the training and analysis of process mining tooling and techniques, which must contend with natural variation present in the execution of real-world traces. To support the field, log trace generation from a variety of process models is therefore required. Such logs are not in adequate supply, as explained in section 2.4.1. The authors’ approach to the problem is to present an algorithm for the mutation of business processes, such that simulation against variations of the business process can produce process logs reflecting those variations. Their algorithm makes use of structure tree representations of processes, which models processes as trees and permits conversion to BPMN models and Petri nets [Bui14]. Pourmasoumi et al. make use of this constraint to demonstrate that their models are block-structured, a mathematical constraint on model structure which 95% of models are shown to comply with [Li10]. Their contribution is a set of transformations on structure trees and block-structured models, and an algorithm applying these transformations to process models, and a tool which implements it built on PLG, a process log generation tool. They conclude that tools such theirs can be used to generate log traces representing process variation, in such a manner as to satisfy the requirements of the process mining research community.

Pourmasoumi et al. describe a list of transformations they explain is “not intended to be comprehensive” [Pou+15], which makes its full potential unclear. Additionally, processes their tool applies to must be block-structured. The importance of this requirement is that it limits their technique to business process models. It is not demonstrated that models of processes in other domains satisfy the condition, such as the flow of data [WC10] or behaviour of human or technical actors in socio-technical systems [WS18a]. Finally, the tool is limited by its lack of capacity to represent variations which are applied contingent on system state. A use case for modelling behavioural variance is to model changes which are impossible to anticipate from the vantage of a modeller, such as a socio-technical system’s human actors’ mistakes, security exceptions and violations, or corrective actions to mitigate undesirable system state. Pourmasoumi et al.’s tool produces variations on a process model, but

modelling behaviour which is expected to vary within iterations of the same model is outwith the scope of their project. **Shugurov & Mitsyuk's work should follow this. REVIEW IT.**

Machado et al. note that there are operational costs to the inefficient modelling of business processes. Specifically, they note that processes can be replicated across automated processes, and failure to identify such scenarios give rise to these operational costs. This work's core motivation is that the representation of variation in process models would allow for the capture of a replicated process once, with instances of similar processes described as deviations from that captured blueprint. On this basis, Machado et al. extend BPMN to support the notion of individual processes as transformations of an underlying process, i.e. that a given process can sometimes be expressed as a deviation from some pattern, and is therefore define-able as the composition of that pattern and a variation upon it. Their approach is illustrated through two small, broadly similar business processes initially modelled in BPMN and then represented in Haskell, allowing the authors to demonstrate their representation of variability as process deviation with realistic examples. While the work presented makes no empirical evaluation of their technique, Machado et al. note that their industrial partner responded positively to the research presented in this publication and that further technical improvements are to be made (support for around advice, and for quantification). They also express an intent to conduct a real-world evaluation in the HR domain. While we are unaware of any real-world evaluation undertaken by this research time to date, some formal proofs that the transformations their tool supports are always well-formed [MAG12].

2.5 Discussion

Aspect orientated programming is designed to permit highly modular software engineering in scenarios where cross-cutting concerns are identified, by isolating them as separate modules [Kic+97]. The aim in employing aspect oriented programming is to reduce tangling, where a cross-cutting concern is intertwined within a program's main concern, and scattering, where the same cross-cutting concern is re-written at many points in a program's source. Aspects which modularise that concern can be written once, separately from a program's main concern, and re-introduced to each point in a program to which the aspect relates by way of weaving. An aspect-orientation framework must therefore be

able to quantify the points at which the aspect applies [FFN00]. Aspects specify both advice (the implementation of its associated cross-cutting concern) and the join point defining where in a target program that advice should be woven. Aspect orientation therefore implies that the source aspects are applied to are oblivious to their application [FFN00].

In theory, the design of the paradigm is such that it should be expected to increase modularity in the software applying it [Kic+97; FFN00] **more citations** , and its proponents often claim this modularity as a benefit of the aspect-oriented approaches of their research [GS04; CM07; Cap+09; JWO12; CBB13]. However, its critics question the reasoning around these benefits, and note that there is little empirical study into whether aspect oriented programs truly benefit as a result of this modularity [CSS04; Ste06; Prz10].

One appropriate application area may be in the representation of behavioural variation in simulation and modelling. The application of aspects to models is already well-studied [ABZ08; CBB13], particularly within aspect-oriented business process modelling [CM07; Cap+09; JWO12], where modelling behavioural variation has also seen some prior research [Mac+11; SA13; Pou+15; Mit+17]. Outside of business process modelling, aspect orientation would reasonably be expected to support the development and observation of models themselves [GK99].

Research opportunities at the intersection of aspect orientation and the modelling behavioural variation occur because behavioural variation is an example of a cross-cutting concern. Changes to expected behaviour such as laziness, boredom, confusion or misunderstanding can impact many parts of a process in a socio-technical system, but modelling the variation caused by any one of these requires similar alterations to behaviour in many disparate parts of the model they occur within. Behavioural variations are therefore both scattered and tangled, and constitute a cross-cutting concern which might be well suited to modularisation in aspects.

Existing aspect orientation techniques and behavioural variation modelling techniques are ill-equipped to take advantage of this alignment. That behaviour changes when it varies is tautological; however, changes supported by existing aspect orientation techniques weave advice before, after, or around their join points, and therefore do not alter the definition itself. In some systems, some variations may be representable as additions inserted before and/or after some other behaviours,

but such techniques are unsuitable for representing modifications of the target behaviour itself, or behaviours which should be omitted instead of added. Additionally, join points available to an aspectual programmer may not be granular enough to permit representing the changes they require in such as system, and aspect orientation’s principle of obliviousness opposes the modification of target code to make new join points available. Techniques which would directly rewrite target source are typically extremely low-level, and therefore ill-suited to most modelling applications [KH98]. Other techniques which permit defining changes to processes at a high level may allow a modeller to describe intended changes (such as the high-level annotations supported in AOPML [Cap+09]), but such techniques are intended for human interpretation, not machine execution for simulation purposes. These techniques permit describing behavioural variation within another process, but only by virtue of the flexibility of natural-language annotation, and are therefore unsuitable for simulation and modelling purposes.

Such techniques also lack executable notion of “state”. Real-world behavioural variance can often be contingent on the environment an actor exists within. While variations such as security breaches might be predictable (by identifying weaknesses in existing processes, for example), variance in socio-technical systems often occurs in the behaviour of human actors. This might be in response to a degraded mode [JS07], where behaviour naturally drifts to a functional — but undesirable — state, or due to an individual making a mistake, forgetting procedure, or being in a state which alters their behaviour, such as tiredness or drunkenness [OLe20]. A framework for modelling behavioural variation using aspects should therefore apply that aspect to a simulated system contingent on the state of that system at a given point in time. High-level modelling technologies such as BPMN and OPM are executable [Mit+17; DOR95], but it is not trivially evident that executable versions of these technologies are compatible with aspect-oriented modifications of their modelling language [CMM10; Cap+09]. As noted, low-level program transformation technologies are also unsuitable. Techniques for applying variations to models exist, but are unsuitable for simulation (and therefore cannot represent application based on temporal state) [SA13], produce individual models representing each possible instance of a variation [Pou+15], do not support the dynamic weaving of aspects for contingent behaviour [Mac+11], or attempt to represent the changes one would expect in the output of a simulation by executing an unmodified simulation and amending its output directly [SM14]. None of these techniques are suitable for representing a behavioural change which is applied contingent on state. Incidentally, these techniques for modelling

behavioural variation also lack support for the alteration of a process definition, or changes “inside” a process definition, as discussed earlier, which also makes them unsuitable for simulated behavioural variation.

Should a tool exist which dynamically weaves definitions of behavioural variations for contingent application and which is capable of expressing changes within a join point rather than before or after it, a challenge would arise as to demonstrating the benefits the tool achieves. Contingent application of behavioural variation, and the ability to define changes to processes specifically, would fulfil the opportunity in marrying aspect-orientation and modelling behavioural variation, but the benefit offered by such a tool is unclear. The introduction of oblivious modification to a model may break its representation of its real-world analogue, making such a model difficult to reason about. Added to this is the complexity of the tool’s capacity to rewrite any join point’s definition. Though aspect oriented literature often lacks case studies demonstrating the benefits of the approach, it is particularly important to investigate whether such a tool could produce realistic models, and whether the expected benefits of aspect orientation as applied to the model hold in practice. In particular, would engineers benefit from the modularity afforded by isolating behavioural variation into an aspect, and would the resulting aspectual modules be re-usable when modelling other systems?

2.5.1 Research Questions

Not sure exactly how to write this. Eeep!

3 | Prior Work

An implementation of the main tool developed and used in this thesis, “PyDySoFu”, predates this thesis. As context for the contributions in this thesis, this chapter will describe the state of the project before the presented research was undertaken. Motivations for the tool’s original development are described in section 3.1, which are followed by its design and implementation in section 3.2, and that of related tooling for experiments and simulations in section 3.3. The chapter concludes with a description of the research undertaken using these tools in section 3.4, as some results in the representation of behavioural variance using aspect orientation were found using these tools which predate this thesis and offer important background for the research undertaken in it.

3.1 Motivations in originally implementing PyDySoFu

The original incarnation of PyDySoFu was developed with different motivations than those outlined in chapter 2. It is therefore important to elucidate the context in which it was designed and developed.

PyDySoFu was originally developed for the representation of behavioural variance in sociotechnical systems, and was first produced as a proof-of-concept; the core contribution was one of tooling. It was developed for use in Python codebases, by virtue of the language’s widespread use and its flexibility in its modelling of data and process.

The original version of the tool was to be applied to models of behaviour in socio-technical systems, where individual actions were represented as functions. Actions which could be decomposed further into more granular actions were to be defined as functions with calls to invoke their more granular counterpart functions. Invocations of low-level behaviours would implement some change to an

environment in the model which its modelled behaviour would be expected to incur. Invocations of high-level behaviours, containing the invocations of lower-level behaviours they compose in the model, would therefore apply the combined effect of the collected behaviours they represent. A benefit of this approach to modelling behaviour was that high-level behaviours could implement the “flow” of a behaviour. For example, a behaviour which would be modelled in a flowchart as having some loop could be modelled analogously in the method described through use of primitive control flow operators in Python, such as `for` and `while` loops.

Another benefit of this approach is that the behaviours modelled have a predictable structure which is amenable to metaprogramming. A low-level behaviour’s affect could be changed by changing the function definition; more structural changes could be made by altering the flow of less granular behaviours. A simple high-level behaviour containing a series of function invocations (modelling an ordered list of steps in the socio-technical system) can be represented as a literal list of function calls. The contents of such a list is trivially modifiable. Removing an item from a list or truncating it at a certain length, for example, are both achievable in a trivial manner using high-level languages such as Python. Notably, many behaviours can be conceived of which could be represented as high-level behaviours but would not be amenable to a simple list of more granular behaviours, such as a behaviour with a looping quality.

With a mechanism to rewrite either an implementation of a behaviour or a collection of behaviours (in the less granular functions mentioned), modelling in such a fashion could therefore lend itself to semantically simple metaprogramming that could represent real-world variations in behaviour. However, largely for reasons discussed in ??, metaprogramming for representing realistic behavioural variations in socio-technical simulations should be able to take advantage of system state. Many real-world behaviours are contingent on environmental state. Real-world actors in socio-technical systems might become tired after lots of work, or proportionally to time of day within a simulation. Therefore, the metaprogramming as described should be performed during runtime, for which no suitable candidate was available. PyDySoFu was developed to fulfil this requirement, so that behavioural variance in socio-technical simulation could be modelled as described and subsequently studied.

3.2 Original PyDySoFu Implementation

To disambiguate the improvements made to the original PyDySoFu implementation in the tooling contributions of this thesis — and to explain the fundamental concepts involved in both implementations’ approaches to aspect orientation and the application of behavioural variance — the implementation of the original PyDySoFu tool, which predates the work presented in this thesis, is described here.

3.2.1 Weaving mechanism

The original PyDySoFu implementation made use of a weaving mechanism which could be categorised as “total weaving” in the parlance of Chitchyan and Sommerville [CS04]: hooks to apply advice are woven into every possible join point. The library was implemented in Python, which offers a flexible object model PyDySoFu is able to take advantage of in order to weave its hooks.¹

Python’s object model has three key properties which the original implementation of PyDySoFu takes advantage of:

1. Everything in Python is an object, including types, functions, and classes
2. Objects are, in essence, implemented as a dictionary² with string keys. All attributes of an object — such as a method on an instance of a class — are values in this dictionary, and their identifiers are the string keys of the dictionary. In essence, `someObject.val` is notionally equivalent to `someObject.__dict__['val']`, though the subtleties of this mechanism will be explained later.
3. Operations on objects are handled by “magic methods”, which are specifically named methods on objects which Python calls for fundamental behaviour in the language. For example, `objA == objB` is interpreted by Python as `objA.__eq__(objB)`. Other built-in behaviours in Python have similar reserved method names which represent the implementation of some behaviour.

These methods can be overridden, or specified by a programmer if they wish.

¹The weaving mechanism of PyDySoFu was eventually factored out into another library, Asp [WS23]; as PyDySoFu was mostly used before this refactoring, particularly in case studies (see section 3.4), the weaving mechanism will here be described within the context of PyDySoFu specifically.

²Python’s name for what other languages might call a `map` or `hashmap`.

PyDySoFu weaves aspect hooks into classes by taking advantage of these three properties of Python. At a high level, PyDySoFu operates by replacing the `__getattribute__()` method of a class object with a custom one. `__getattribute__()` is the magic method responsible for performing lookups in an object’s underlying dictionary. The replacement `__getattribute__()` also looks up attributes in the relevant object’s underlying dictionary, but it also searches for advice to be applied when performing these lookups, and applies any advice it finds. The replacement `__getattribute__()` method represents the aspect hooks woven by PyDySoFu.

Hooks are applied to a class by way of an invocation to a function, `fuzz_clazz`, which takes a class as a parameter and weaves aspect hooks into that class [WS18c; WS23]. `fuzz_clazz` replaces the `__getattribute__()` method of the class with a new function object which it constructs. The replacement function object injects aspect hooks to callable attributes, thereby discovering aspects, applying them, and invoking an attribute which is the target of advice within a wrapper of the attribute itself.

The replacement `__getattribute__()` method makes a call to the class’ original `__getattribute__()` method to retrieve an attribute when required. If this attribute is not a function or method, it is returned by the woven `__getattribute__()` function and the program affected class behaves as if it was never altered. However: if an attribute to be retrieved is a method or function, a new function is constructed and returned. This function looks up and applies advice to the attribute originally retrieved by the program, and contains a reference to the original aspect to enable its execution. This wrapping function is therefore the core of the weaving process: a replaced `__getattribute__()` method injects the aspect hooks which constitute the aspect-orientation framework. Advice to be run before, around, and after a join point were implemented as invocations of advice functions before a target was executed for “before” advice, after a target was executed for “after” advice, and with an around function being run with a reference to the advice’s to invoke at its own discretion.

3.2.2 Applying Process Mutations

The development of PyDySoFu was intended to support simulation & modelling research by providing a way of applying program modifications at runtime. It sought to fulfil the needs of a

program which determined that a (potentially non-deterministic) change to the behaviour it modelled was required, and enable the program to apply that change mid-process in an aspect-oriented fashion. Aspect orientation libraries typically support advice woven before, around, or after a join point; modifying the join point itself essentially allows changes inside its definition, introducing a fourth type of weaving. PyDySoFu achieves this through a special type of “before”-style aspect, which it calls a “fuzzer”.

Fuzzers implement transformations on abstract syntax trees. They are implemented as functions which receive a list of AST objects representing the definition of a function which satisfies a fuzzing advice’s join point, and return another list of AST objects, which replace the target function’s definition. Any transformation resulting in a valid AST is permitted. A code snippet demonstrating the implementation of this process is shown in fig. 3.1.

As can be seen in fig. 3.1, fuzzing aspects are implemented by way of “prelude” advice (PyDySoFu nomenclature for advice run before a target invocation, inspired by early work on TheatreAG, which is discussed in section 3.3.2). Prior to the target being run, the fuzzer replaces its underlying code object with one produced by compiling a fuzzing aspect’s returned AST objects. Replacing the underlying code object with a modified one allows the fuzzer to define arbitrary modifications to the definition of the function, thereby achieving runtime metaprogramming.

3.2.3 Limitations

The original implementation of PyDySoFu was intended to demonstrate the feasibility of runtime metaprogramming and its potential in simulation and modelling. However, its design presents limitations.

There is an overhead involved in running the wrapped function for every invocation of `__getattrute__()`, and also in running aspect hooks for all possible join points, even when those hooks are not targets of advice at a given moment. When an attribute is the target of advice, aspects are discovered and applied. However, aspects are discovered by lookup within the scope of the function creating a replacement `__getattrute__()` method. This design requires multiple instances of advice weaving to create multiple replacement `__getattrute__()` calls, all of which are invoked on any attribute lookup. A

```

class FuzzingAspect(IdentityAspect):

    def __init__(self, fuzzing_advice):
        self.fuzzing_advice = fuzzing_advice

    def prelude(self, attribute, context, *args, **kwargs):
        self.apply_fuzzing(attribute, context)

    def apply_fuzzing(self, attribute, context):
        # Ensure that advice key is unbound method for instance methods.
        if inspect.ismethod(attribute):
            reference_function = attribute.im_func
            advice_key = getattr(attribute.im_class, attribute.func_name)
        else:
            reference_function = attribute
            advice_key = reference_function

        fuzzer = self.fuzzing_advice.get(advice_key, identity)
        fuzz_function(reference_function, fuzzer, context)

def fuzz_function(reference_function, fuzzer=identity, context=None):
    reference_syntax_tree = get_reference_syntax_tree(reference_function)

    fuzzed_syntax_tree = copy.deepcopy(reference_syntax_tree)
    workflow_transformer = WorkflowTransformer(fuzzer=fuzzer, context=
        context)
    workflow_transformer.visit(fuzzed_syntax_tree)

    # Compile the newly mutated function into a module, extract the
    # mutated function code object and replace the
    # reference function's code object for this call.
    compiled_module = compile(fuzzed_syntax_tree, inspect.getsourcefile(
        reference_function), 'exec')

    reference_function.func_code = compiled_module.co_consts[0]

```

Figure 3.1: A code snippet from the original PyDySoFu implementation, implementing Fuzzing aspects and applying fuzzing to a function definition.

single target of advice which has multiple pieces of advice woven therefore incurs a performance penalty for every piece of advice applied, which must be incurred when any attribute is looked up on the target's class. If the join point defining the target may apply to many classes, each class must incur the same penalty, even if none of their attributes are targets of advice in practice. The original library was developed to demonstrate the feasibility of the idea underlying PyDySoFu (runtime metaprogramming), but the weaving mechanism implemented left room for improvement. A robust implementation with attention paid to reducing this overhead is introduced in chapter 4.

Similarly, an additional overhead is incurred by a lack of caching of the modifications fuzzers make to function definitions. One can envisage a need for runtime metaprogramming which produces different function definitions at different times: an example could be modelling different degrees of degraded modes introduced to an actor's behaviour in safety-critical systems research [JS07]. One can also envisage no such need: an example could be minor temporary modifications otherwise permanently made in program maintenance, such as to constants within a function definition, to the format of a function's return value, or adding control flow which exits a function early on a termination condition. The requirement is a product of the tool's use case in different scenarios. In scenarios where the same modification is to be made every invocation, a fuzzer need only be run once; optimisations enabling the caching of fuzzing aspects' effects would provide better performance in use cases where such a feature is appropriate.

Other aspect orientation frameworks offer support for other types of advice. Handi-wrap and AspectJ both support features related to the processing of exceptions thrown by a program [BH02a; Kic+01] and these features have inspired work into improved exception handling in object-oriented systems [MD11]. However, this version of PyDySoFu offers no direct support for exception handling. Opportunities to support the feature were therefore available for future revisions of the library to capitalise on; such a revision is also presented in chapter 4.

A final limitation is that the weaving technique this early of PyDySoFu used is incompatible with Python3, as replacing `__getattr__()` is not possible in Python's newer version. It was determined that a tool which was of practical use to the simulation and modelling community should be produced which would remain useful to future researchers making modern models; as the existing weaving technique lacked performance, an opportunity presented itself for a complete redesign. The

resulting new design is presented in chapter 4 which makes use of a new weaving technique.

3.3 Additional Simulation Machinery

Other related projects developed tooling for sociotechnical simulation & modelling. Fuzzi-Moss was a project collecting a library of standardised behavioural modifications for use in sociotechnical simulation & modelling; Theatre_AG was a project offering a model of time against which actors within sociotechnical simulations & models could act. While these projects ultimately were not used in producing the contributions of this thesis, they are outlined here as relevant to the original PyDySoFu project as they were originally developed as a suite of simulation & modelling tools to be employed together.

3.3.1 Fuzzi-Moss

Fuzzi-Moss³ was a library of standard behavioural variations written as fuzzers to be applied by PyDySoFu [SW16]. It was primarily created for use in a model of the impact of inconsistency in teams’ executions of software engineering methodology, which is discussed further in section 3.4.

Fuzzi-Moss contained utilities and fuzzers such as:

- Several probability mass functions representing chance of unexpected behaviour given a length of time and an actor’s...:
 - conscientiousness, a lack of which would increase chance of behavioural adaptation due to lack of effort;
 - concentration, a lack of which would increase chance of behavioural adaptation due to distraction
- A missed_target fuzzer, which terminated a while loop early if activated via a probability mass function of an actor’s propensity for negligence. As an example of a Fuzzi-Moss fuzzer, this is shown in ??.

³A backronym for “Fuzzing Models of sociotechnical simulations”


```

def missed_target(random, pmf=default_distracted_pmf(2)):
    """
    Creates a fuzzer that causes a workflow containing a while loop to be
    prematurely terminated before the condition
    in the reference function is satisfied. The binary probability
    distribution for continuing work is a function of
    the duration of the workflow, as measured by the supplied turn based
    clock.
    :param random: a random value source.
    :param pmf: a function that accepts a duration and returns a
    probability threshold for an
    actor to be distracted from a target.
    :return: the insufficient effort fuzz function.
    """

    def _insufficient_effort(steps, context):

        break_insertion = \
            'if not self.is_distracted() : break'

        context.is_distracted = IsDistracted(context.actor.clock, random,
            pmf)

        fuzzer = \
            recurse_into_nested_steps(
                fuzzer=filter_steps(
                    fuzz_filter=include_control_structures(target={ast.
                        While}),
                    fuzzer=recurse_into_nested_steps(
                        target_structures={ast.While},
                        fuzzer=insert_steps(0, break_insertion),
                        min_depth=1
                    )
                ),
                min_depth=1
            )

        return fuzzer(steps, context)

    return _insufficient_effort

```

- An `incomplete_procedure` fuzzer, which truncated the steps⁴ taken by an actor if activated via a probability mass function representing an actor’s propensity for distraction

Plans were also made for fuzzers representing an actor “becoming muddled”⁵ and making mistakes in decision-making, but neither have been completed at time of writing. A discussion around the revival of this project in the context of the PyDySoFu rewrite presented in chapter 4 is given in ??.

3.3.2 Theatre_Ag

Theatre_Ag (“Theatre”) is a project defining a model of time against which actors in sociotechnical models & simulations can act [Sto23]. In the project’s overview, it describes itself as...:

“Theatre_Ag is a workflow oriented agent based simulation environment. Theatre_Ag is designed to enable experimenters to specify readable workflows directly as collections of related methods organised into Plain Old Python Classes that are executed by the agents in the simulation. All other simulation machinery (critically task duration and clock synchronization) is handled internally by the simulation environment.

The central metaphor underlying Theatre’s model of timing is theatrical: actors in a simulation or model are members of a “cast” (a collection of actors) who enact a “workflow” (simulation steps) in a “scene” (domain model within which the actors interact). Central to the library is its clock: tasks are given durations, and a clock which synchronises all agents’ position in time ticks to complete different tasks. The theatrical model theatre introduces is the context for PyDySoFu’s nomenclature for its types of advice: “prelude” advice happens before a task and “encore” advice is invoked afterward, as a prelude and encore would be in a literal theatre.

Theatre has been used as the environment in models of TCP/IP, algorithmic trading, the spread of disease [OLe20], and the impact of behavioural variation in software engineering methodologies as described in section 3.4.

⁴Represented by lines of code

⁵A behavioural variation caused by confusion.

3.4 Example Studies using PyDySoFu for Behavioural Simulation

The viability of encoding behavioural variations as aspects using PyDySoFu has been demonstrated in earlier studies [WS18a; OLe20]. The study modelled software engineers working to different methodologies of software engineering: waterfall, in which requirements are gathered, software is developed to meet requirements, quality assurance steps are undertaken, and the resulting software is delivered to customers; and TDD, where the development of tests for quality assurance precedes the development of features. The study sought to investigate whether, when software engineers were working suboptimally, there was a difference in the rate of bugs introduced to a program developed under each methodology.

Maybe I should include something on Aran's masters dissertation too, rather than just citing it?

The study began with a “naive” model of software engineers following each paradigm, developed in Python using PyDySoFu, Theatre, and Fuzzi-Moss. Engineers would produce “chunks” of code, which could contain bugs. In quality assurance, engineers were modelled as attempting to identify bugs in different areas of the codebase, fixing them if they were discovered. Developers could commit chunks of code toward features identified through requirement engineering, which were eventually completed, but could potentially contain undiscovered bugs within chunks of code.

This model was then augmented aspectually using PyDySoFu. Distraction was represented through the truncation of functions representing workflow steps. Developers were modelled with different levels of distraction, affecting a probability mass function (PMF) which would activate when a developer was modelled as being distracted in a given moment. If the PMF activated, the workflow step invoked at that moment was truncated using PyDySoFu. The model showed that developers following the TDD methodology could successfully complete a larger number of features on average than those following waterfall, concurring with the prevailing consensus on the two methodologies.

In replicating the community understanding of the model, the paper demonstrated the feasibility of aspectually augmenting modelled behaviour: the simulation took a naive model with no capacity for analysing errors, and introduced new features of the model supporting an avenue of investigation otherwise impossible with the methodologies represented by the naive model. That the resulting simulation matched the expectations existing within the community gave confidence that the tool

Add a citation for TDD vs agile! Citations are missing from the copy of the CAiSE paper I can find, but I'd like to use the same one

could be used to build realistic simulations where some features of a model could be separated from its core codebase.

3.5 Discussion

The existing case study employing PyDySoFu demonstrated that realistic model features could be separated from their core codebase, and gave credence to PyDySoFu’s use as a tool for aspectually augmenting models with behavioural variance. It also left many research questions unanswered and

Does this want to be a list? Consider reworking to a paragraph.

tooling flaws unsatisfied, however:

- Aspects were believed to be “realistic” as they represented the expected outcomes of the simulation. However, no real-world data was used to corroborate the claim, and it was unclear that aspectually augmented behaviour could capture the variations present in real-world human behaviour.
- These aspects also demonstrated variations which were identifiable in the emergent properties of a system (for example, mean time to failure of a software system under development, or successfully completed features). The variations applied to individual developers might have poorly modelled individual behaviour, but produced accurate emergent properties of the system individual developers acted within.
- As discussed in section 3.2.3, PyDySoFu’s implementation at the time was a proof-of-concept which, while successfully demonstrating the potential of aspect-oriented runtime metaprogramming, was also inefficient, feature-incomplete, and lacked compatibility with modern software engineering tooling.
- Models of distraction were adopted from the common library provided by Fuzzi-Moss. However, this model was not applied to other codebases. It remains unclear that Fuzzi-Moss’ model of distraction is broadly applicable in other projects: different models of distraction might be required by different researchers. Further, a model of distraction which realistically represents the behaviour of an individual (rather than the emergent properties of the system that individual acts within) might not apply to other systems the individual acts within. Briefly put, The portability of aspectually modelled behavioural simulation has not been investigated, and literature within

the aspect orientation community lacks evidence to support a belief in their portability [Prz10; CSS04; Ste06].

This left opportunities to improve both the tooling offered for aspect-oriented runtime metaprogramming, and the evidence supporting its use to encode behavioural variations in sociotechnical systems. Improvements to the tooling follow in chapter 4; later chapters propose — and discuss the implementation of — real-world systems which are suitable for modelling using PyDySoFu in chapter 5, and a study of those systems using aspectually augmented models in chapter 6 and chapter 7.

4 | Rewriting PyDySoFu

Check the most up-to-date pdsf implementation — is it actually on that backup drive?

The work undertaken in this thesis required in improved implementation of old tooling. When previously used, PyDySoFu was a proof of concept which could feasibly produce scientific simulations, but was implemented in a manner which was not optimised for speed (making it a burden for large simulations), lacked granularity in the application of its aspect hooks (hooks could only be applied to entire classes), and most importantly, did not work with Python3 (Python2 support officially ended during this PhD).

This chapter briefly outlines the new implementation of PyDySoFu, discusses improvements made to design and performance, and explains some contributions made to the design of aspect orientation frameworks which addresses some core issues raised with the paradigm . Consider adding references to the sections through this PDSF chapter, depending on how beefy it becomes...

Rewrite chapter outline after the structure of the chapter is known to be OK.

4.1 Requirements for Change

As time wore on with PyDySoFu's original implementation, it became increasingly clear that a rewrite was required . PyDySoFu grew out of an undergraduate project, and accrued technical debt as a result of being written under extreme time constraints with little experience. On revisiting, and on reflecting on other aspect orientation frameworks (as discussed in section 2.2 and [CS04]) and the use previously found for PyDySoFu (see [WS18a; WS18b]), it was clear that there were a series of improvements which could be made in the process of rewriting the tool:

Too informal

sloppily put

- Before this body of work, PyDySoFu made use of techniques for applying aspect hooks which did not translate to the changes Python 3 made to its object model. In particular, Python 3 changed its underlying object model, using a read-only wrapper class that made the replacement of `__getattr__` impossible via the previous route.
- PyDySoFu’s original implementation made no serious accommodations for efficiency. It could be seen as the “total weaving” described by Chitchyan and Sommerville in [CS04], and it was not possible to provide additional options to ensure that aspects could be as efficiently woven as possible at runtime given a particular use-case .

Could be less
sloppy

- The original PyDySoFu implementation wove onto a class, meaning that even properties of the class which were not considered join points were still affected by the weaving, even if in a minor way . Because `__getattr__` retrieves all attributes including special builtin attributes and non-callable attributes, these are also returned via the modified implementation of `__getattr__`, incurring an overhead, albeit small, for all attribute resolutions instead of a desired subset.

Could be less
sloppy

- The original PyDySoFu implementation made no accommodations for scenarios where fuzzing of source code was applied in a “static” manner. That is to say, where a deterministic modification to source is woven as advice, instead of dynamically modifying source code, the same modification would still be made every time the target attribute was executed, unless caching of results was specifically managed by the aspect applying the change. No optimisations were made pertaining to this, but compilation and abstract syntax tree editing have the potential to be PyDySoFu’s most expensive operations.
- Unlike other aspect orientation frameworks such as AspectJ [Kic+01], join points could not be specified by pattern. Instead, each individual join point must be supplied as a Python object. This means that, while the target attributes are still oblivious to the advice applied to them, the application of that advice could not be written obliviously.

As a large number of requirements were left unfulfilled by the original implementation of PyDySoFu, a new implementation satisfying them was deemed necessary.

4.2 Python3-compatible Weaving Techniques

Replacing `__getattritude__` on the class of a targeted method was no longer viable in Python 3. A replacement method therefore had to be found. For clarity: replacing `__getattritude__` allowed for hooks to be woven (at runtime) into likely future targets for advice. These hooks would then discover and manage the execution of advice around each target. Because advice can be run before and around a target, and dynamic weaving implies that advice could be supplied or removed at any time, we look to intercept the calling of any target, and manage advice immediately before execution. So, the task at hand is to find a method of attaching additional work to the calling of any potential target, before that target is executed. We refer to code woven around a target which manages applied advice as aspect hooks.

4.2.1 Abandoned techniques

Rather than “monkey-patching”¹ a new version of `__getattritude__` with hooks for weaving aspects, the rewritten method could be patched to the object itself at a deeper level than used in the original PyDySoFu implementation. This would make use of Python’s `ctypes` api to patch the underlying object. Similar work has been done in the python community in a project called `ForbiddenFruit` [con21]. Efforts were made to add the required functionality to `ForbiddenFruit` — patching `__getattritude__` directly on the object, or “cursing” it in `ForbiddenFruit` jargon — but this was abandoned as the underlying mechanism is particularly unsafe, Python API changes could render the work unusable in future versions easily, and the implementation would only work with particular implementations of Python (for `ctypes` to exist, the Python implementation must be written in C). Community patches existed for cursing `__getattritude__` which did not work, and attempts proved challenging, indicating that this would also be complicated to maintain over time. There are also efficiency concerns with this technique depending on its use: weaving advice around a function would mean monkey-patching the built-in class of functions, which would incur an overhead from

¹Monkey-patching is the practice of making on-the-fly changes to object behaviours / definitions by taking advantage of language properties such as flexible object structures. Common examples of these structures are objects literally being maps from string attribute / method names to the associated underlying value, as in Python and Javascript. Monkey-patching makes use of these simple structures by replacing values such as the function object mapped to by the original function’s name in the dictionary, effectively changing its behaviour. This is the method by which PyDySoFu originally replaced `__getattritude__` on a class object.

running aspect hooks on every invocation of every function.

Another approach involved making use of existing Python functionality for interrupting method calls. As PyDySoFu wraps method calls at execution time, what is required is to add functionality to the beginning and end of the execution of a method. Python has built-in functionality for implementing debuggers, profilers, and similar development tools, which provides exactly this functionality, as debuggers must be able to — at any point during execution marked as a breakpoint — pause a running program and inspect call stacks, the values of variables, and so on. As a result, the method `settrace()` allows a developer to specify a hook providing additional functionality to a program. Making use of this also has issues in our case. Most significantly, `settrace()` catches myriad events in the Python interpreter which PyDySoFu may not need to concern itself with, incurring significant overhead. In addition, use of the function overrides previous calls to it, meaning that any debuggers used by a user of PyDySoFu would be replaced with PyDySoFu’s functionality, which was deemed untenable. However, it is worth noting that the technique could work in theory, and if future versions of Python allow for multiple trace handlers being managed by `settrace()`, this could provide an interesting approach when implementing future dynamic aspect orientation frameworks.

4.2.2 A viable technique: import hooks

A final available technique was to continue to monkey-patch hooks to discover and weave aspects, via an alternative method which did not make use of `__getattr__`. This approach would change the use of PyDySoFu slightly to make a compromise between performance and obliviousness of aspect application: when importing a module targeted for aspect weaving, methods which are potential weaving targets are invisibly monkey-patched with a wrapper method with a reference to the original² and hooks to detect and run dynamically supplied advice.

An important note for discussing the implementation of PyDySoFu is that almost all Python functionality operates by use of its “magic methods”. “Magic methods” are methods beginning and ending with two `_` characters. The Python language documentation specifies sets of magic methods and their required function signatures which are used internally to implement functionality; for example, any object with the method `__eq__()` defined can be compared against using the `==` operator, and

²Necessary to run the originally targeted method.

the `__eq__()` magic method is run to determine the outcome of the operator. Magic methods support more than operator overloading. For example, anything which defines `__len__()` and `__getitem__()` is treated as an immutable container, and adding `__setitem__()` and `__delitem__()` makes that container mutable. Any class defining `__call__()` is treated as a callable object (not unlike a function). More can be found in Python’s documentation[VD09], although more focused guides exist in the Python community [con16]. For the purposes of implementing import-based aspect hook weaving, magic methods have the affect of making the language ideal, particularly as an environment to implement dynamic aspect orientation. Python’s functionality for importing modules is managed by builtins `__import__`, which receives module names as strings and handles package resolution. By monkey-patching the import system, modules can be modified during the process of importing. As this technique allows for control over where aspect hooks are applied, PyDySoFu can target only function and method objects to apply aspect hooks to, avoiding the overhead its previous iteration introduced when applying hooks to all attribute lookups including non-callables, such as variables or Class objects.

Monkey-patching builtins `__import__` is as simple as replacing the function object with a new one, which has the effect of changing the behaviour of Python’s `import` keyword: because all Python functionality relies on magic methods implicitly, its behaviour can be altered in this way. However, our intent is not necessarily to manipulate all modules, but a subset of imports specified by a modeller as suitable for manipulation. If all invocations of `import` wove hooks into modules, including those made in the process of importing packages, an unnecessary overhead would be introduced when invoking any callable in any module. Therefore, it is important to have a mechanism to enable and disable the weaving of aspect hooks for a given `import` statement (effectively, to enable and disable PyDySoFu’s modified import logic).

This can be achieved through another use of magic methods in a manner which also makes clear to a modeller exactly where aspect hooks are being applied: making use of Python’s `with` keyword. A more technical discussion on the process of weaving, and the nature of the hooks applied, follows in section 4.3.

```
with AspectHooks():
    import mymodule
```

Figure 4.1: Example of importing a module, mymodule, using PyDySoFu’s import hook design.

4.3 Import Hooks

4.3.1 Implementing import hooks

We are interested in manipulating builtins `.__import__` only when imports are made which introduce modules containing prospective join points; a developer might import many modules, which do not all require aspect hooks to be introduced. We enable this new import behaviour with a syntax of the form shown in section 4.3.1, which weaves aspect hooks into all functions and (non-builtin) class methods within the mymodule module object added to the local namespace of the importing stack³. Less formally: importing with `AspectHooks()` applies aspect hooks to all potential targets of advice in the mymodule package. The behaviour of Python’s `with` keyword is defined by more magic methods: any object with `__enter__()` and `__exit__()` defined can be used here, where `__enter__()` is run at the beginning of the enclosed block, and `__exit__()` when leaving the block.

PyDySoFu caches the original builtins `.__import__` object in an instance of the class, and replaces it with `AspectHooks.__import__`, in its `__enter__()` method. This is reversed by replacing builtins `.__import__()` with the cached object in its `__exit__()` function. `__enter__()` and `__exit__()` are the magic methods corresponding to entering and exiting Python’s `with` blocks, meaning that this technique modifies Python’s importing only when imports are made within a block such as that in section 4.3.1. The resulting implementation for weaving aspect hooks is uncomplicated, as can be seen in fig. 4.2.

³Python’s use of the stack namespace in its importing system means that careless re-importing a module can lead to multiple copies of it in different function stacks, meaning that the same name resolution (such as resolving a class by its name in a module) might, after applying aspect hooks in PyDySoFu, change the behaviour of procedures depending on where they are called. Scenarios where this might arise are deemed unlikely enough that the risk of this design decision becoming troublesome are considered negligible. Still, it would be remiss not to make note of the fact.

```

class AspectHooks:
    def __enter__(self, *args, **kwargs):
        self.old_import = __import__
        import builtins
        builtins.__import__ = self.__import__

    def __import__(self, *args, **kwargs):
        # ...replacement import logic for performing weaving...

    def __exit__(self, *args, **kwargs):
        builtins.__import__ = self.old_import

```

Figure 4.2: Magic methods used to enable the with keyword usage for PyDySoFu

4.3.2 Strengths and weaknesses of import hooks

As a technique for weaving aspect hooks, this new method provides multiple benefits. Application of aspect hooks is straightforward from the perspective of a modeller using PyDySoFu, whose code clearly applies aspect hooks and does so in a legible way for future maintainers. Import hooks’ explicit application of hooks to modules means that places where aspect hooks might be applied equally explicit, and thereby implements an interpretation of aspect-oriented programming’s principle of “obliviousness” moderated by clarity. While join points are oblivious to potential aspect application, callers of join points can expect to be more aware, and signs for codebase maintainers are left directly within source code. Aspect hooks can be applied to specific modules or every module depending on the use of the supplied with statement, allowing for total weaving or actual hook weaving [CS04] depending on their preferences. Further, performance is optimised in comparison to the previous implementation of PyDySoFu, as hooks are weave-able at a more granular level (on the level of procedures such as functions and methods, rather than all attributes of a class).

There are also caveats of this approach that are necessary to address. As aspect hooks are woven in the new implementation of PyDySoFu via Python’s import functionality, any procedure not imported from a module cannot have aspect hooks attached. *Consider adding local namespace weaving to pdsf3: should be easy to implement as a cheeky little monkey-patch...* However, as aspect orientation is primarily concerned with a separation-of-concerns approach to software architecture, targets are expected to exist in other modules, and we do not consider this to be a significant limitation.

A more significant limitation of the import hook approach is that the object with aspect hooks

woven exists in the namespace of the module importing the join point. In other words, this method makes it impossible for a module to make use of aspect hooks that are woven in an unrelated piece of code. *is this true??? verify, I think not, if imported normally. tested but didn't realise that pdsf won't fuzz functions with single-character IDs, so inconclusive.* We therefore have a “semi-oblivious” property to our aspect orientation approach: targets of advice are unaware of any adaptations made, but any code making use of those adaptations must be aware enough to at least apply aspect hooks.⁴

In a manner of speaking, this can be considered to alleviate some concerns with aspect orientation as a paradigm. Aspect Orientation is criticised for making reasoning about programs more difficult [Prz10; CSS04; Ste06]. One cause of this is that aspects separate logic from where it is run; Constantinides, Skotiniotis, and Stoerzer’s comparison with the jokingly proposed *come from* statement [Cla73; CSS04] is a reminder that it can be effectively impossible to understand how a program will execute if the path of execution is not at least linear or clearly decipherable from source code. Aspect orientation as a paradigm violates this linearity by design. However, import hooks as implemented in fig. 4.2 allow for aspect-oriented code which can be interpreted in one of only two ways:

1. Looking at the original implementation of a procedure, its intended execution is clear. A programmer can make use of this directly and it is guaranteed to behave as expected.
2. Any program making use of a procedure imported from a module will see, when the procedure is imported, whether it has had aspect hooks applied. In this case its behaviour is unknown — falling prey to the design flaws discussed in the aspect orientation literature ?????? — but this unpredictability is highlighted to the programmer.

As a result, while import hooks are somewhat limited in that they are applied specifically to imported code and break the traditional AOP concept of obliviousness in at least a weak manner, these two facts combine to arguably fix a latent issue in the design of the aspect oriented paradigm. The original PyDySoFu implementation was able to modify any procedure in a more traditional, oblivious manner. While this new implementation is clearly more limited as a result, we consider these limitations an overall benefit to the design of the tool, and a contribution to aspect orientation framework design.

⁴Note that once aspect hooks are applied, advice can still be supplied from anywhere in the codebase.

4.3.3 Weaving process

Weaving in PyDySoFu’s updated implementation takes place via monkey-patching of aspect hooks, as described in ???. Aspect hooks replace executable targets within a module at the moment the module is imported. When the target is invoked, the wrapping aspect hook is executed in lieu of the original target object. The wrapping aspect contains the target function within its closure, allowing it to execute the original target; however, it was also created by the AspectHooks class, and so has reference to aspects registered against it.

The process of weaving with this implementation of PyDySoFu is therefore as follows:

1. An aspect is imported using AspectHooks, as shown in section 4.3.1.
2. An aspect is registered against AspectHooks as shown in fig. 4.3, by providing a regular expression matching the ID of any join points to apply an aspect to and the aspect to apply. Specifically relevant to aspect registration:
 - Methods permitting aspect application are:
 - (a) AspectHooks.add__prelude(rule, aspect), which registers advice to be run before a target is invoked
 - (b) AspectHooks.add__encore(rule, aspect), which registers advice to be run after a target is invoked
 - (c) AspectHooks.add__around(rule, aspect), which registers advice to wrap an aspect invocation, effectively providing the functionality of prelude and encore advice with a single aspect
 - (d) AspectHooks.add__error_handler(rule, aspect), which registers advice to catch and process exceptions raised by a join point
 - (e) AspectHooks.add__fuzzer(rule, aspect), which registers advice to modify a target before it is invoked, effectively providing aspect application within a join point, and permits arbitrary ast-level modifications of the target
 - Each method described above also accepts an optional urgency integer parameter, enabling optional additional features of PyDySoFu which are discussed in section 4.4.3.

- Invoking the above aspect hook registration methods compiles and caches the regular expression provided for efficiency when identifying matching join points in later invocations of aspect hooks.
- Each method outlined above returns a callback which de-registers the aspect provided, to facilitate the ephemeral application of aspects. This could be useful in experimental codebases where aspects represent behavioural deviations, and a researcher looked to run simulations with different deviations applied to compare datasets, for example.

A bit clunky
— take the
last sentence
out in favour
of notes in
discussion or
rewrite.

3. Invocations of executable objects within the module imported in 1. are replaced with aspect hook wrappers of themselves; an invocation of an executable object within that module therefore triggers PyDySoFu's aspect hooks, which seek aspects registered against the ID of the target they wrap. Aspects registered against a join point identifying rule which matches the ID of the wrapped target can therefore be identified, and executed as appropriate.

All aspects applied when registering advice are callable objects; however, they have different function signatures. The arguments expected by each type of advice is:

prelude advice takes `target`, `*args`, `**kwargs`

encore advice takes `target`, `original_return_val`, `*args`, `**kwargs`

around advice takes `next_around`, `target`, `*args`, `**kwargs`

error_handler advice takes `target`, `handled_exception`, `*args`, `**kwargs`

fuzzer advice takes `ast_steps_from_target`, `*args`, `**kwargs`

...where `target` is the target an aspect is applied to; `original_return_value` is the value returned by a target after it was run; `next_around` is what an around-style advice runs when it has finished its pre-target component, and intends to run the target itself; `ast_steps_from_target` is a list of AST objects representing the definition of the target a fuzzing aspect is applied to; and `*args`, `**kwargs` is a Python idiom collecting arbitrary positional and keyword arguments passed to a function but not specified within its function signature, which are here employed to collect arguments passed to a


```
def log_invocations(target, *args, **kwargs):
    """
    An aspect to print invocations of a target.
    """
    print("Invoking " + target.__name__)
    target(*args, **kwargs)
    print("Invocation of " + target.__name__ + " finished.")

from pdsf import AspectHooks
with AspectHooks():
    from some_module import some_func

AspectHooks.add_around("some_func", log_invocations)
```

Figure 4.3: Registering an aspect against AspectHooks

target, and can be used by any aspect to inspect the expected behaviour of the target (or to pass into the target, if the aspect invokes it).

The execution of prelude, encore, and error_handler aspects work predictably: prelude aspects are executed before the wrapped target, encore aspects after the wrapped target, and error_handler aspects within a try block encompassing the execution of all advice. Notably, the return values of around aspects replace those of their target, and encore aspects replace that of the target if they return any value which is not None. By comparison, the implementations of around and fuzzer aspect weavers are non-trivial.

As with all aspects, around-style aspects are given the target being invoked alongside its supplied positional and keyword arguments. However, these aspects are also given the next around-style aspect woven against the target to run. Each aspect should be provided its next_around argument. However: with a series of many around-style aspects to apply, and each calling the next in the series, a given aspect would need to know the next_around parameter for the aspect it calls; the second could call a third, meaning it would need to know the next_around parameter for its own successor, which could only come from the first in the series. The third could call a fourth, requiring another next_around parameter, which the second would have to pass to the third, and the first to the second. The inelegance of this naive solution to the problem seems confusing and inelegant for modellers to interact with or understand.

Decide whether the “-style” prefix I sometimes use should stay or go when copyediting. Or: when do I use it? Be consistent.

However, two alternatives present themselves. One is to use Python’s generator pattern, which can be simply described as a callable which returns successive items from a series on each invocation.

```

from functools import partial, reduce
final_around = lambda target, *args, **kwargs: target(*args, **kwargs)
nest_around_call = lambda nested_around, next_around: partial(next_around
    , nested_around)
target_with_around = reduce(nest_around_call, around_advice, final_around
    )

```

Figure 4.4: Simplified example of weaving around aspects in PyDySoFu

An alternative approach is to supply each around-style aspect with its first argument (the next in the series) before the aspects are invoked, unburdening individual aspects with responsibility to retrieve and apply the next aspect successfully. This approach is employed by PyDySoFu. The first argument of every around-style aspect is provided by constructing a partial function. Partial functions are a feature of many languages which allow for some arguments of a function to be provided, but for the function not to be invoked; instead, a new function is returned, with any remaining arguments left to be provided on invocation. With the first argument provided, the resulting partial function has the signature `target, *args, **kwargs`, where all values are known and trivially supplied by any around-style aspect.

The problem to solve is therefore the construction of the partially applied aspects in such a way that each refers to the next correctly. To solve this, each function is provided to its precursor; iterating through all around-style aspects in this manner results in a single aspect which contains a reference to the second, the second a reference to the third, and so on until the final registered aspect. The final aspect has no successor it can reference however; to overcome this, a function which does nothing but execute the original target is provided, with the same signature as the partially-applied aspects. The final function terminates the chain, and is responsible for executing the target itself; every other aspect is simply responsible for the execution of the next around-style aspect in its chain.⁵ In this way, around-style aspects can be woven in a manner which simplifies a developer's interaction with the aspects. A simplified snippet of the source code of PyDySoFu which implements this process is shown in fig. 4.4 to illustrate.

The final kind of advice to apply are fuzzer-style aspects, which PyDySoFu contributes to aspect orientation framework design. These aspects make modifications to the definition of the target they are applied to, allowing advice to be woven within a target. Arbitrary transformations of the target also

⁵Note that this is technically optional; if it determined it was necessary, any around-style aspect could break the chain by calling the target with its arguments directly.

```

if fuzzers is not None and fuzzers != []:
    code = dedent(inspect.getsource(t))
    target_ast = ast.parse(code)
    funcbody_steps = target_ast.body[0].body
    for fuzzer in fuzzers:
        non_inline_changed_steps = fuzzer(funcbody_steps, *args, **kwargs)

        if non_inline_changed_steps:
            funcbody_steps = non_inline_changed_steps

    target_ast.body[0].body = funcbody_steps
    compiled_fuzzed_target = compile(target_ast, "<ast>", "exec")
    if not isinstance(t, FunctionType):
        t.__func__.__code__ = compiled_fuzzed_target.co_consts[0]
    else:
        t.__code__ = compiled_fuzzed_target.co_consts[0]

```

Figure 4.5: A code snippet representing the process of modifying a target of an aspect application using a fuzzer-style aspect. `t` refers to the target on which modifications are applied.

allow removal or change of any part of the target definition. This is achieved by acquiring the original source of the function to modify using Python’s built-in reflection library, `inspect`. The source can be parsed into an abstract syntax tree using Python’s built-in library, `ast`, which contains a list of AST objects representing a target’s original definition. This list of AST steps is then passed to a fuzzer aspect as an argument, which returns a new list of AST steps if any changes are to be made. The new steps are compiled and the resulting Python code object containing a modified function definition is monkey-patched into the original target, effectively changing its definition. The original code object is cached during aspect hook weaving, and is used to replace the modified code object after target execution to avoid unexpected behaviours in future invocations of the target. A simplified example of the process of acquiring, modifying, compiling, and monkey-patching a target’s underlying definition is found in fig. 4.5 .

4.4 Optimisations

Additional features were implemented in the presented incarnation of PyDySoFu which improve its usability in the research software engineering context, and optimise performance where necessary: deep hook weaving, non-dynamic weaving, and aspect priority support.

```
from pdsf import AspectHooks
AspectHooks.deep_apply = True
```

Figure 4.6: Code snippet enabling deep hook weaving

```
from pdsf import AspectHooks
AspectHooks.treat_rules_as_dynamic = True
```

Figure 4.7: Code snippet enabling dynamic weaving

4.4.1 Deep Hook Weaving

Under ordinary circumstances, PyDySoFu weaves aspect hooks into an imported module, but not modules imported by that module. Deep hook weaving enables weaving of aspect hooks into all potential join points by applying hooks to functions and methods imported by the original module, but also any recursive imports. A code snippet enabling the feature is presented in [fig. 4.6](#).

Deep hook weaving introduces performance overhead due to additional checks for aspects which were woven dynamically. However, it is plausible that the modules within which desired join points reside are not imported directly by a developer, but are available to them indirectly through another package they make use of. It is also plausible that a developer may look to instrument the entirety of — and investigate details of — a call stack without attaching a debugger to a process. In both examples, PyDySoFu’s default behaviour is insufficient. Deep weaving is therefore provided to support these use cases.

4.4.2 Non-Dynamic Weaving

PyDySoFu’s runtime weaving of aspects allows for aspects to be applied at any time. Aspects can be removed and re-applied during program execution. However, programs may be written with the intention of applying aspects once — as program modifications to be introduced without direct manipulation of a codebase — so that, once defined, the set of applied aspects would remain static. In this scenario, the dynamic weaving of PyDySoFu introduces unnecessary overhead. Weaving aspects dynamically means that every invocation of a prospective join point requires searching for aspects which presently apply to it.

```
from pdsf import AspectHooks
AspectHooks.manage_ordering = True
```

Figure 4.8: Code snippet enabling priority ordering of aspect application

To avoid this overhead in scenarios where it is unrequired, PyDySoFu can cache the aspects applied when any callable wrapped by an aspect hook is invoked for the first time. On its first execution in this mode, the aspect hook stores the set of aspects it matched to the invoked target, and future invocations retrieve the set of aspects to apply from the cache. This avoids expensive regular expression matches, which fail in all cases but those where an invoked target is to be augmented by the application of an individual aspect the regular expression is paired with.

For performance reasons, the default mode of PyDySoFu is to run with non-dynamic weaving. Dynamic weaving functionality is enabled by toggling the `treat_rules_as_dynamic` flag on the `AspectHook` class, similarly to deep hook weaving (see section 4.4.1). A code snippet demonstrating this is shown in fig. 4.7.

4.4.3 Priority Ordering of Aspect Application

As dynamic weaving allows for non-deterministic application of aspects, it may be that the order in which aspects are intended to be applied is not the aspect in which they are woven. To support these use cases, aspects may be sorted by priority during the lookup of aspects to apply when a join point is invoked.

As mentioned in section 4.3.3, when aspects are registered against the `AspectHooks` class, an optional urgency parameter is available. The parameter takes an integer representation of priority of aspect application, with higher numbers representing more urgent application. Aspects with no urgency applied default to `urgency=0`. High-urgency aspects are applied before low-urgency aspects. This feature is disabled by default, but can be enabled using the code snippet shown in fig. 4.8.

4.5 Discussion

The new implementation of PyDySoFu makes a few contributions , particularly in comparison to

Add metrics here on the performance implications; they're pretty significant, I actually added them because the regex matches slowed my own code down to a crawl. Our experiments make use of non-dynamic weaving for this reason.

contributions to what?

Is this list
complete?

the previous version :

- Introduces a new technique of weaving aspect hooks on import, improving its design over a typical aspect orientation framework by making use of Python's with keyword when weaving hooks to trade a degree of obliviousness for clarity
- Aspect hooks can be applied with more precision than the previous implementation of PyDySoFu, meaning:
 - Users of the framework can better delineate between total and actual hook weaving
 - Unnecessary overheads from checking dynamically applied aspects at each join point are reduced.

PyDySoFu's current incarnation also provides opportunities for improvements and for future work. Our intended use case for aspect orientation for simulation & modelling is in scientific codebases specifically; direct integration with the scientific package ecosystem (which is vibrant in Python's community) should be made. A good initial project would be integration of aspect application in sciunit tests [Tha+17] . A discussion on potential use cases of PyDySoFu together with existing research software engineering technologies is provided in section 8.7 .

5 | RPGLite: A Mobile Game for Collecting Data

RPGLite is a game designed to mimic existing popular games, while being structured to permit exploring all game states via formal methods; Kavanagh and Miller have produced PRISM models which can be model-checked to identify ideal play strategies in all game states [KM20]. Some experiments were conducted around RPGLite to answer the question: “Over time, do players converge on an ideal strategy of play?” Fix formatting and correct wording of William’s hypothesis. Insert my own research question here

Real-world play of the game invites many research questions. For example, an alternative question to answer would be, “what strategy of play do players typically adopt?” or, relatedly, “do all players adopt the same strategies?”¹ Kavanagh and Miller’s work can identify the cost of an action [KM20; Kav21], allowing for richer datasets to be produced and more in-depth analysis of real-world play to be conducted. One possibility invited by these datasets is that it may permit modelling of real-world players’ interactions with the game, and therefore offers an opportunity to investigate the effectiveness of aspect-oriented simulation and modelling. Referring back to TODO INSERT RESEARCH QUESTION HERE, it would be feasible to build a model of RPGLite play which did not satisfactorily model real-world data, inviting aspectual augmentation of the model in an attempt to improve its performance, and thereby investigating the research question proposed.

However, it would not be possible to perform analyses of player behaviour without real-world player

¹These are not of research interest for the purposes of this thesis, but examples of interesting questions to ask of a game where “correct” and “incorrect” actions can be categorised.

Lift screen-shots in this chapter from game-on paper source where possible

data. Datasets collected empirically would be required to compare against synthetically generated datasets, to ascertain their similarity and thereby compare different methods of producing synthetic datasets. As RPGLite is a real-world sociotechnical system which already yields a formal analysis of play, and datasets from this system can be used to support multiple avenues of research in different disciplines,² producing a dataset of real-world RPGLite could also support future work in a variety of fields as a non-commercial dataset for research in fields including at least game design and gameplay research [Kav21], formal methods [KM20], sociotechnical simulation research, and research software engineering tooling demonstration.

To that end, a collaboration was undertaken with Kavanagh to develop and release a mobile implementation of RPGLite which would collect player data for research purposes. The dataset produced would enable Kavanagh to demonstrate the utility of their model checking in an empirical scenario [Kav21], and for the purposes of this thesis, it would also enable the analysis of models representing player behaviour, by supporting the comparison of these models against the collected data.

The development and release of RPGLite as a mobile game offers an opportunity to make use of aspect orientation in a new context: a model of naive RPGLite play can be produced which represents random play, and aspects could be written which augment the naive model representing hypothesised player behaviour. If the aspectually augmented models generate data which correlates with empirically sourced data more closely than that generated by the naive model (with random play, and no aspectual augmentation), we can dismiss naive play as “realistic”, and understand the aspectually augmented behaviour as “more realistic”. Many aspects can be written representing different styles of play, which might be adopted by different players, a concrete benefit of aspect orientation in modelling & simulation.

The modelling & simulation research opportunities RPGLite presents are realised in chapter 6 and chapter 7; this chapter focuses only on the design and development of RPGLite as a mobile game, and the dataset collected from real-world play and released [Wal+20; Wal+21]. It starts with an overview of RPGLite itself in section 5.1, and progresses to a discussion of the game’s implementation in section 5.2 before concluding with a discussion around the empirical data collected using the game

²Later contributions in this thesis are supported by datasets produced by RPGLite, as were the contributions of Kavanagh’s PhD thesis [Kav21].

in section 5.3.

5.1 An Overview of RPGLite

Needs copy-
editing

RPGLite is a simple two-player game played in turns. Each player selects characters independent of the other, with each character having a unique set of abilities and properties, which are generally health, chance of success on attack, and damage dealt on a successful attack. The abilities of some characters necessitate additional properties. Each player selects an “alive” character (one with health greater than 0) to perform their action against a chosen “alive” target (or occasionally targets). A successful attack — randomly determined by chance of successful attack for the selected attacking character — results in that character’s unique ability being inflicted on their target[s]. A random player is chosen to take a first move, players may always skip their turn as a valid action, and players continue to take alternating turns until a victor is left with the only “alive” characters.

Eight characters are available for selection, with the following abilities:

Knight Deals damage to an opponent character on a successful hit.

Archer Deals damage to two opponent characters on a successful hit.

Wizard Deals damage to an opponent character on a successful hit, disabling (or “stunning”) them for the duration of the opponent’s next turn.

Healer Deals damage to an opponent character on a successful hit, and heals themselves or, optionally, the other player character instead (assuming that character is still alive).

Barbarian Deals damage to an opponent character on a successful hit, dealing additional damage if their health is low when attacking.

Rogue Deals damage to an opponent character on a successful hit, dealing additional damage if the target’s health is low when attacked.

Monk Deals damage to an opponent character on a successful hit, and immediately takes another turn, until their attack is unsuccessful.

Gunner Deals damage to an opponent regardless of success, dealing additional damage on a successful hit.

Specific details of each character — their health, chance to hit, and damage on hit as well as character-specific details (such as the threshold for additional Barbarian or Rogue damage, for instance) — are defined as a “configuration” of RPGLite. Different configurations change the game’s “balance”, a term referring to the relative strengths of different characters or character pairs. For example, if a configuration leaves many characters with initial health values close to a Barbarian’s threshold for additional damage, then they become a very powerful character due to their ability to inflict additional damage. If the Monk’s chance to hit is high, the repeated turns it offers can be very advantageous. Character skills can work in concert with each other: choosing a Barbarian and Healer such that the barbarian can be kept at low health for additional damage, but the healer can be used to keep them alive, may be an effective strategy depending on the game’s configuration. Kavanagh et al. found that model-checking a configuration of the game could discover the relative strengths of characters and character pairs when played optimally [Kav+19].

RPGLite’s design has two objectives it must meet. First, that it is interesting to players, which requires that it is approachable and complex enough not to be immediately solvable. This is necessary for real-world data collection, and to demonstrate a design representative of something that could conceivably be a real-world game with an active playerbase. Second, RPGLite’s design must be sufficiently simple for model-checking. Model checking is a necessary requirement of design because of our need to identify optimal moves: analysis of player behaviour rests on our understanding of how close to “ideal” players are, and whether players approach ideal strategies over time. This is the crux of the work found in formal methods research on RPGLite [Kav21; KM20], which relies on a reduced state space in order to calculate optimal moves and character pairings.

5.1.1 RPGLite’s Design Implications

The state space of RPGLite makes it unusually well-suited to analysis through formal methods. Because of this, datasets produced through simulation of RPGLite can be compared against two other datasets: one of real-world play, and another of what can be mathematically shown to be “correct”

player behaviour.

To demonstrate this state space, note that RPGLite games can have their states described by a set of values: the healths of characters on each team, plus a stunned character. With eight characters having a maximum health value of w, x, y, z^3 , two players, and an indicator of which character is stunned.⁴ The entropy of a game’s state is therefore $\log_2(w \times x \times y \times z \times 3 \times 2)$, where the multiplications by 3 and 2 represent the stunnedness indicator for the current player (either character, or neither), and the player whose turn it is to play, respectively. The maximum health for a character is 10 hit points, which makes the maximum entropy of a game state $\log_2(10^4 \times 3 \times 2) = \log_2(60000) \approx 16.87\text{bits}$. Each player picks 2 of 10 characters, and can choose to attack either opponent character with either of their own, or skip their turn, for a maximum of 6 possible actions. We can therefore see that the total entropy of the entire RPGLite game is no greater than $\log_2(10^4 \times 6 \times 2 \times \binom{10}{2}) \times 6 \approx 24.36\text{bits}$.⁵

Iterating through these states allows us to map the entire state space of RPGLite. As the state space the game defines is relatively small⁶ it is feasible to analyse every possible game state. Note that the figure of $\approx 24.36\text{bits}$ includes movements between states as well as the states themselves, meaning that it is feasible to map the potential progressions through all possible games of RPGLite using formal methods. Further, this allows us to understand the chances of a player winning given transitions between different states, for example by representing moves in the game as transitions in a decision diagram where nodes are the game’s state. We can calculate exactly how “good” a move is, by comparing chances of success making a given move in a given state against the chances of success when making the objectively optimal move.

In this way, RPGLite’s design allows it to be understood formally, yet it also draws on common game design elements and is sufficiently “interesting” to generate data from a real-world playerbase.⁷

³Maximum health values are dependent on RPGLite’s configuration.

⁴Note that stunnedness is valid for exactly one character for one turn, meaning that only one character may be stunned at any time, and the status effect immediately resets, meaning there are only three possible states for stunnedness: either character belonging to the player taking a turn, or neither.

⁵The actual figure is smaller:

- Players may choose two characters with special abilities that prevent them from attacking both opponent characters at once (this accounts for 9 out of 10 characters), giving 5 possible actions in a turn, rather than 6.
- The “metagame”, which refers to the perceived “best” strategies at any given point in time, would impact the chance of a player selecting making certain moves or choosing to play with certain characters. RPGLite is designed to be slightly “unbalanced” in the parameters of different characters such as health, attack damage, or potency of special abilities, meaning long-term players are expected to learn effective playstyles and adjust accordingly.

Players’ behaviour is therefore less uniform than this calculation would imply, but the calculation provides a higher bound on the entropy of the game.

⁶For comparison, mapping valid positions in chess takes about 136 bits [Nie77], and this figure does not account for valid moves within the game, which our calculation for RPGLite does

⁷Demonstrated through data collected from several thousand completed games, which can be found at [KWM20].

This yields some useful properties for the purposes of aspect-oriented simulation:

1. Simulated moves can be selected naively, i.e. at random, but can also be made perfectly according to the known-correct move in a given game state, or made with some calculated “cost” as to the chance of winning.
2. Real-world players’ behaviour can also be analysed according to the same metrics: for example, moves made by players can be analysed to understand bias, whether players learned to play “better” moves over time, or whether they selected known-strong characters more frequently than those who can be formally shown to have a relatively low chance of winning games.
3. As the actions taken when playing RPGLite are consistent — such as deciding the target character of an attack, or a character to use in an attack — random play can be simulated as a “naive” play style, which can be compared against real-world players. Where player behaviour does not correlate to naive play,⁸ the biases of players may be represented as aspects which are applied only to specific actors within the simulation.
4. Should aspects be suitable as a manner of accurately representing biased play, aspects offer a separation of concerns within the simulation: any nuance found within the playstyle of specific real-world players would be replicated and applied not to the model itself, but to specific simulated players. Playstyles might also be mixed with the application of multiple aspects.

Whether aspect orientation is suitable for the realistic simulation of RPGLite gameplay is the topic of the remainder of this thesis. However, the design of RPGLite allows for a controlled system where a clear notion of “good” and “naive” behaviours can be defined, the system is closed insofar as all interactions within the system are known and all game elements are precisely understood, and all interactions take place between experiment participants for data collection purposes, allowing for a large dataset to be collected without information being removed due to players not consenting to their data being collected and disseminated for science. In short, RPGLite’s design constitutes a system where all aspects are well-understood, no interference is anticipated from system components which are unknown or outside of experimental control, and lots of data can be collected for analysis. Therefore, the data generated by gameplay is suitable for comparison against datasets outputted through aspect-oriented

⁸The concept of play correlation is introduced in chapter 6.

simulation. This can be used to assess whether simulated players with aspect-affected behaviour accurately reflect the playerbase, and so can help to assess whether aspect orientation is suitable for realistic simulation.

5.2 Implementation of RPGLite

Knowing ideal play is useful. However, to understand how real-world players would interact with RPGLite, empirical data needed to be collected. To produce this, a mobile online multiplayer version of the game was developed for data collection purposes. Play constituted engagement with an experiment for data collection, and after several months, a database logging player behaviour offered a dataset which could be used to simulate real-world player behaviour.

This section describes details of RPGLite’s implementation as a data collection tool. Some lessons learned after reflection on the implementation process were documented for the benefit of others’ avoiding our errors[Wal+21] which provides some reflection on the process of RPGLite’s development .

5.2.1 Consent to participation

Players gave consent to participate in a scientific study as a part of creating an account to play RPGLite. Players were required to explicitly scroll through an information sheet and consent agreement and agree to both to make an account. Copies of both were made available to download for player reference in both the game and RPGLite’s website⁹ . Players were instructed to contact the involved researchers to withdraw participation, and were instructed that they could do so at any time. Email addresses to contact were noted in both the app and the website.

The study, information sheet, consent form, and data collected were reviewed by the University of Glasgow Science & Engineering ethics committee before the game was publicly released.

⁹RPGLite’s website is available at <https://rpglite.app/>, where copies of the consent form and information sheet are linked.

Should a summary of the paper be a section here, or maybe an appendix? Keen not to plagiarise, but it’s also maybe worth including. William and I both spent a lot of time getting that paper out. I presented it at GameOn2020. He’s included some in his thesis, so I

5.2.2 Mobile app

As a mobile game, RPGLite’s user-facing component was an application, distributed through the Google Play Store on Android and the Apple App Store on iOS. This was developed in Unity, a framework for developing games in C# which can be distributed to almost any platform¹⁰. Most assets were developed in GIMP, with character designs contributed by a commissioned artist online. Unity allowed for a “WYSIWYG” or what-you-see-is-what-you-get interface builder, with event handlers defined in C# code which would “hook” into events signalled by interface element interactions. User-facing components of the game were largely produced by Kavanagh, the collaborator on this project and original RPGLite designer. Therefore, in an attempt not to take credit for this component of the work involved in collecting the RPGLite dataset, curious readers are referred to Kavanagh’s notes on the development process for full details [Kav21].

Beta testing required user engagement. Apps were deployed to Android and iOS devices of colleagues, who played a series of games to check that game logic was sufficiently robust and graphic design sufficiently adequate for public distribution of the game. In a manner inspired by design science methodology [JP14], beta tests comprised of iterations on the artefact of the game and its server-side component (discussed in section 5.2.3) which were distributed to beta testers. Informal feedback and enhancement requests were sought on each iteration until the game behaved correctly in all edge cases (no major bugs were reported) and a final design was settled upon (no major complaints about interaction or aesthetics were reported), at which point the game satisfied criteria for public distribution.

An example of the game’s visual evolution is given in fig. 5.1, where colourful buttons replace a tabular, text-heavy interface. Another example is the evolution of the game’s main screen, “Active Games”, which loaded when players logged into the app and allowed users to see and interact with active games they were involved in. The visual identity and colour palette of this screen was refined iteratively, as can be seen in fig. 5.2. Other features which were developed following user feedback while beta testing include the implementation of a notification system and a leaderboard showing a

player’s experience relative to their peers; these are shown in fig. 5.3.

¹⁰Meaning that there are technically also versions of RPGLite playable on, say, a games console or web browser.

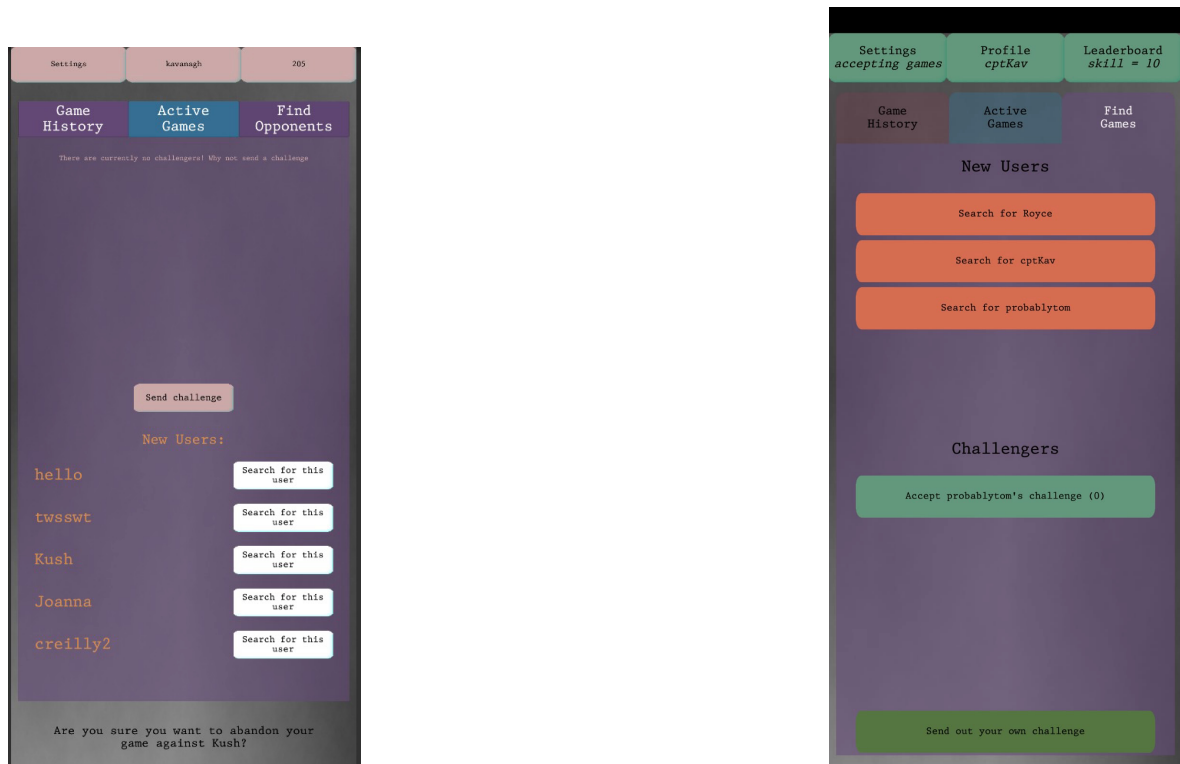


Figure 5.1: Evolution of the “Find users” matchmaking screen. Prototype left, final design right.

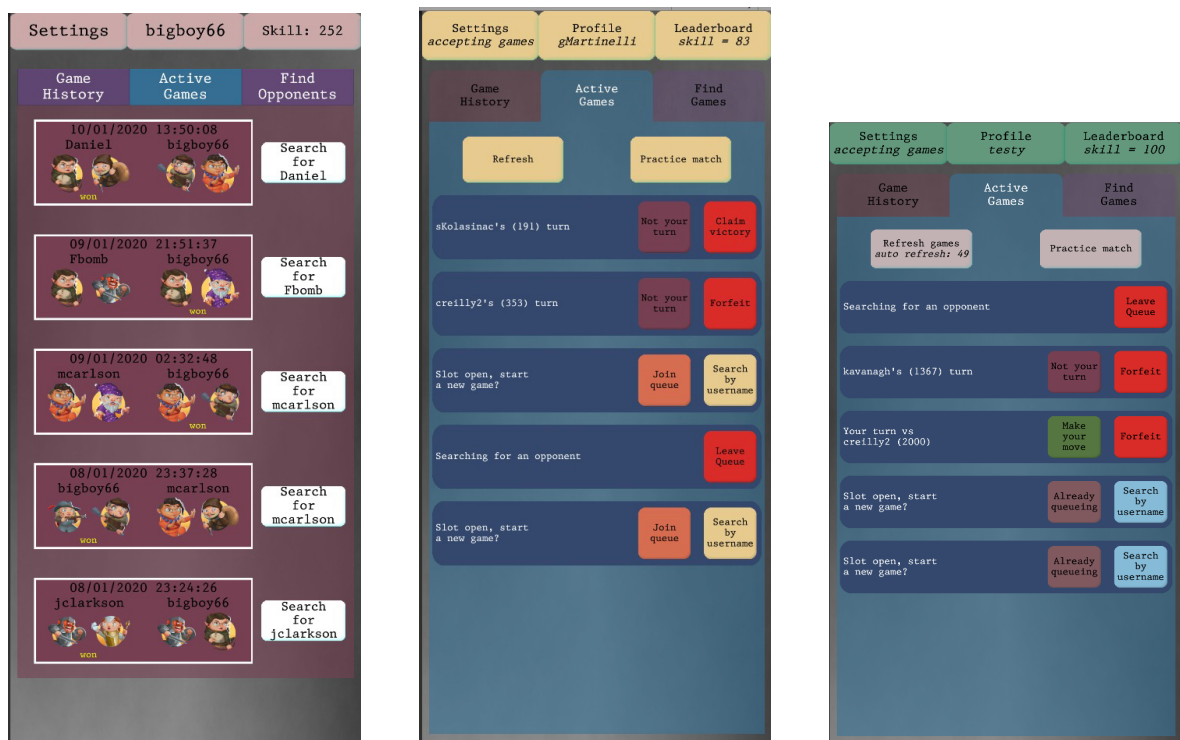


Figure 5.2: Evolution of the “Active Games” screen, which allowed RPGLite users to see and interact with games they were playing. An early prototype is shown to the left; a refinement through beta testing in the center; and the final version released to the public to the right, with an improved colour palette.




Figure 5.3: Features developed following requests from beta testers. A leaderboard of player engagement is shown to the left; examples of notifications when an opponent has acted are shown to the right.

5.2.3 API & Server-side Logic

As the data collected ought to be empirical, RPGLite was developed as an online game. This required a server and API for a client to communicate with.

A REST API was developed with Python's Flask framework. Endpoints were created to support all in-game actions requiring centralised state (or a recording of that state), allowing for player search, matchmaking, player profile design, game history and statistics analysis, ranking calculations, login and password reset, synchronisation of access to sensitive information, and other in-game activities. The API also allowed moves to be made, and rejected erroneous game states or unauthorised input from any malicious actor, to ensure that the data collected, analysed, and published was not manipulated. The API also sent push notifications to an opponent's device when moves were made. This was a feature requested during beta testing, which was anecdotally observed to increase engagement with the game.

On each of these actions, data was collected about the action performed, and logged in a database. In addition, in-game activities which required no server-side input but were considered to have potential

in later analysis would log datapoints through the API into the same database.

A MongoDB database instance was installed and managed on a University of Glasgow Computing Science virtual machine. The no-SQL nature of the database permitted flexible structuring of the data, and simplified analysis of the games' results after data collection was completed. The API was also hosted on the same virtual machine. A combination of port access rules applied via Nginx and hardening the security of the database itself through MongoDB's account features and controlled permissions within the server prevented unauthorised access to the database, ensuring that the data remained untampered.

Game state was interpreted client-side. The intention in this design choice was to un-burden a centralised service we were responsible for maintaining, moving functionality such as calculating game states after moves were created to clients on players' devices. This left the serverside logic mostly concerned with executing database transactions and ensuring the integrity of game states was uncompromised. Reflections on the suitability of this approach made after the game's launch have been published separately [GameOn2020].

5.3 Empirical Play and Data Collected

In total, players produced a dataset comprising around 9,693 games. It also includes 1,105,065 datapoints generated by gameplay or player interaction with the client, such as players checking their position on a leaderboard, searching their game history, or rolling to hit as part of an attack. The data is drawn directly from the MongoDB database used to run the game. With player consent on signup (see section 5.2.1) this data is published and available for all researchers to analyse in the format of a JSON object [KWM20].

Completed games drawn from the MongoDB instance contain many fields, and the in particular the breadth of datapoints logged by player interaction with the RPGLite app means there are an extremely broad range of datapoints available to curious researchers. A full list of available data is available in a text document within the dataset ??; in brief, some of particular interest include:

- The history of moves made, and the times those moves were made

- The players involved and the winning player of each game (by username as no personally identifying data was collected)
- The ELO scores of players in each game
- The characters chosen by each player
- The “score” of each player¹¹

RPGLite was played through two “seasons”: after an initial 3,000 games played in a first configuration of RPGLite, a second configuration was released. This meant that, having learned a strategy to play with the initial configuration (a preference of characters, for example), changes were made which may invalidate what players had learned. The dataset therefore contains two conceptual subsets: data collected from a system with the same mechanics but slightly tweaked

¹¹RPGLite’s mobile app presented users with a naive scoring mechanism used to rank users on a leaderboard which some users then used to identify other players of a similar notional skill.

6 | Simulation Optimisation with Aspect Orientation

With a game deployed to experiment participants and a dataset of empirical play collected, it was possible to determine optimal play in any game state. This entirely separate body of work is documented in another student's PhD thesis [Cite William's PhD thesis](#) . This dataset leads to further research. If we understand how players should play, and we have data to indicate how they do play, we can investigate how real-world players might be modelled.

6.1 Aims

Aspect orientation's use in previous simulation and modelling efforts have typically focused on the use of aspects to compose model or simulation details [There must be tons of good citations for aspects being used to compose together a simulation / model](#) . Critics of aspect orientation note that the act of process composition makes visually understanding codebases difficult, and so ensuring that a simulation properly models real-world behaviour is made trickier with the introduction of aspect orientation. However, aspect orientation might instead be used to augment an existing model, by rethinking what aspects are used to represent.

An alternative use of aspects would be to first build a non-aspect-oriented model of expected behaviour, and separately build aspects which describe deviations from this. For example, one might more realistically simulate safety procedures by first producing an idealised, “naive” model of what

We don't think this is the case — those citations just don't exist! Rework to discuss use of AOP in business process modelling, and note that nobody made

employees are expected to do, and separately model alterations to prescribed behaviour as an employee’s boredom, expectation that checks and balances are unnecessary wastes of their time, and so on — effectively, separating out models of degraded modes[JS07].

Previous research on the use of aspect orientation to model degraded modes adopted the traditionally claimed benefit of aspect orientation: separation of cross-cutting concerns, allowing for a greater reusability of codebases[WS18a]. A repository of cross-cutting concerns in socio-technical simulation such as boredom was developed as a library to be applied to any future models[SW16]. However, aspects used in simulation have no intrinsic need to represent concerns that are cross-cutting. Indeed, whether they can be accurately used to represent cross-cutting concerns in simulation is the topic addressed in [Add a cross-reference to the chapter on cross-cutting concern simulation accuracy when it exists](#). Aspects might instead be used to represent amendments to processes which deviate from an expected norm, in this case represented by the idealised model aspects are applied to.

To more concretely relate this to the experiment at hand: play of RPGLite can be modelled as players matchmaking, picking characters, and then mutually taking turns until one player’s characters are entirely expired. Once a player’s characters are dead, new matches can be made. This can continue indefinitely. Lacking a heuristic to select next moves or characters, players might be modelled as picking random moves. However, heuristics for move selection can be added to the naive model of play by way of augmenting the processes already defined through aspects. This approach can be of significant utility in both modelling player behaviour and accurately modelling different players:

1. Different players might use their own unique heuristics to model play. Each player’s behaviour is therefore well described by separating what play “looks like” to what makes a given player play differently to their peers.
2. Different players might lean more heavily on different heuristics, or mixes thereof. Play might be characterised by reliance on experience, on recent games, on knowledge of an opponent, and so on; these different variables can be expected to be weighted differently by each player, adding complexity to the code which models this individualised play.
3. A modeller might discover a new idea for a heuristic long after developing an original concept for a model. Due to the impact of 2., ideal architectures for an approach such as this should allow

these heuristics to be defined entirely separately to the base model to maximise flexibility when maintaining experiments.

Considering 1., 2., and 3., architectures and paradigms which enable separation of concerns are well-suited to defining alternative approaches to play. Some architectural approaches such as mixins or plugin design patterns might support this structure well, but they typically rely on language features (in the case of mixins) or knowledge of software engineering (in the case of design patterns). Aspect orientation is typically provided to developers as a framework or runtime in a language (such as AspectJ[Kic+01] or PROSE[PGA02]) and can require minimal architectural understanding to use: concepts are simple, and the effort of composition is alleviated by the supporting framework or runtime.

The approach makes little use of aspect orientation’s significant contribution — cross-cutting concerns — as whether behaviour cross-cuts different parts of a codebase is not of interest in this use case. Instead, aspect orientation is treated as a composition mechanism with a reasonably low degree of technical knowledge required.

6.1.1 PyDySoFu Suitability

Some aspect orientation frameworks do not adequately achieve this requirement. For example, the most influential framework, AspectJ, requires the use of language extensions to define integrate aspect orientation[Hil+00], and similar additional complexity is added in seemingly every alternative framework, through the use of bespoke virtual machines, compilers, translators, or languages[rajan2006nu_towardsAO_invocation; PAG03; SL07; BH02b].

PyDySoFu, however, requires very little additional knowledge to use. Its design prioritises simplicity and a shallow learning curve that makes its adoption by researchers without a software engineering background feasible: **maybe cut this list of reasons PyDySoFu is fantastic...**

- PyDySoFu is implemented as a pure-python library, meaning that it can be installed through Python’s package manager (pip) and imported like any other Python library. No additional supporting infrastructure is required.
- Aspects in PyDySoFu are simple functions which take as arguments whichever pieces of informa-

tion are pertinent for the function’s use as an aspect¹.

- To weave a PyDySoFu aspect requires only a method call, which returns a `callable` which unweaves that aspect.
- Defining PyDySoFu pointcuts requires only a regular expression matching a method name. This can apply to a wide range of join points if required, but where method names are provided directly, the join point is made clear.
- Additional clarity over where aspects can be woven is introduced by PyDySoFu’s transparent weaving of aspect hooks, mitigating some of aspect orientation’s most prominent criticisms.

PyDySoFu therefore satisfies the requirements of this work well: it offers composition of procedures outside of the scope of an original codebase, makes what is being composed where clear to a programmer, and makes no significant changes to Python as a language (thereby requiring users to specialise in fewer tools).

6.1.2 Proposed Experiment

Aspect orientation’s use as a composition tool for model components makes sense in principle, but it is unclear whether the addition of behaviours to a naive model would make the model more “realistic”. Furthermore, changes to a model could alter its representation so as to weaken its mimicry of the system it simulates; adding behaviours could make it less realistic. The fundamental issue at play is that it is unclear whether the changes made would properly represent what might be empirically observed. While PyDySoFu’s design makes understanding what is being composed simpler than other aspect orientation frameworks, a composed model under this paradigm is still split across multiple areas of a codebase, making a visual assessment of whether a model accurately reflects the intended behaviour impractical. It is therefore important to demonstrate the efficacy of PyDySoFu and the modelling paradigm it introduces, by confirming the realism of a model to which behavioural variation is applied.

We can confirm whether aspects can realistically represent changes to a naive understanding of

¹For example, an “encore” aspect which is woven after a target procedure returns will be provided that target’s return value.

the real world by comparing their output against empirical data. For example, if a such a model of behaviour in a system outputs data which correlates poorly against empirically collected data, a change to that system would make it more realistic if it improved this correlation, and could be said to be realistic if the generated data appeared sufficiently “close” to the empirical dataset — which here means that the correlation between the two is of statistical significance. Such a change can be aspect-oriented. Therefore, we can see the application of aspects as the application of packages of potential improvements to a base model, which can be verified by way of comparison to known-good datasets.

This is the basis of the experiments in this thesis.

With datasets collected empirically on RPGLite’s play, we can build a naive model of play and aspects to apply that should realistically model data from players. This can be used to answer the question:

“ Can aspect-oriented models be said to exhibit realism?”

To answer this question, a naive model of play is produced. Aspects are developed which model learning within the system defined by the naive system. The synthetic datasets produced by models with naive or aspect-applied models can be compared to an empirical dataset sourced from real-world players of the game, and their similarity compared.

“Naive” is used here to describe a model which does not encode any understanding of the players of the game being modelled. Traits such as learning, distraction, or aptitude for similar games are irrelevant to the naive model. We need a naive model to demonstrate the effectiveness of an aspect-oriented alternative: we can measure how closely it reflects empirical data, and compare this against the same measurement drawn from another model with behavioural variations encoded using aspects. A closer match from our aspect-oriented model would demonstrate that the technique can enhance a model’s realism. Our naive model also provides a null hypothesis: if no improved similarity is observed, the technique brought no improvement to the model’s realism. No measurable difference indicates that weaving behavioural variations as aspects has no impact on a model’s realism.

Write a little here on the learning aspects.

Flesh this out as a brief wrap-up of our experimental technique. The end of our “naive” model explanation might be useful to move down here / rework into this para. Contrasting the similarity of the empirical dataset to both naive and aspect-applied datasets... This discussion is provided so as to provide context for the following sections; experimental design is discussed in more depth in section 6.5.

6.2 Naive Model

A naive model of play was developed by separating each stage of the actions taken by players in the client-side app, and separating them into individual procedures.

To facilitate the retrieval of information pertinent to a simulation in an applied aspect, the model was written so as to contain simulation state as mutable function arguments. The model was written as a workflow, and state of workflow execution was separated into three components: the actor that a function invocation (or “step”) represents activity from; the context of that step in the execution of a workflow; and the context of that workflow’s execution in a broader environment. Incidentally, we found this structure to allow a flexible and natural implementation of a procedural simulation, which should translate easily to existing simulation frameworks such as SimPy[Mat08]:

Actor — allows the function to identify the actor performing the activity defined by the function.

This argument is any object uniquely identifying an actor.

Context — allows the function to determine details of the current thread of work being undertaken by the actor. This is necessary because in some simulations, the same actor might pause and resume multiple occurrences of the same activity — for example, they might concurrently play three different matches in RPLite. As a result, it is necessary to understand the context of the action being performed by the actor in question. This argument can be any object uniquely identifying the context of a piece of work, but should be mutable (such as a class or dictionary-like object) to permit the communication of information across invocations of different action-representing functions.

Environment — an actor’s actions are often determined by the global environment they act within.

There may be ancillary details to the actor’s actions and the context of their particular thread of work which they are undertaken within which are used to determine behaviour, such as a landscape they traverse or other actors they might choose to interact with. Because all actors share access to a global environment, this also provides a message passing space, or a space where actors can set values and flags other actors might look to, should those details be more general than their specific thread of work at a given point in time. ²

Each simulation step receives these three arguments at a minimum. Because steps of the model are functions, and therefore valid join points, aspects applied to these have access to the entire state of the simulation.

The naive model of RPGLite follows a simple workflow mimicking player interaction with the client-side application used by real-world players. A graphical representation is provided in fig. 6.1. **Position this figure properly on the page. Maybe reword? Pretty verbose.** The model it describes produces synthetic data. It does so by simulating players interacting with each other in a broader RPGLite ecosystem. The large number of players is important: future behavioural variations introduce behaviours which guide future actions based on past experiences. Non-determinism early in the model can therefore cause different players to act in different ways. Simulating many players allows the system to represent a wide variety of early experiences, as would occur to real-world players, too. As the aspects being woven simulate learning, early experiences which cause players to learn maladaptive strategies — or simply strategies less optimal than those learned by their peers — should later contrast against winning strategies, so that players with less useful early inferences can learn from players who play more optimally.

Consider changing fig. 6.1 to only include details about the model of the game itself, leaving the steps about data analysis to the later aspect-applied version (which is the real experiment). Move fig. 6.1 to .svg for fidelity The stages of the naive model as laid out map to those encountered by real-world players. Two randomly-selected players repeatedly select characters to play (from the pool of 8 characters available in the real-world game), and a player is chosen to play first at random. That

²This is different to environments in some other simulation frameworks, such as SimPy[dev21], where the environment controls scheduling and execution: this structure imposes no constraints such as models of time, and anticipates that any such functionality should be implemented by the programmer. However, an environment such as SimPy’s might satisfy a programmer’s needs when using this particular pattern.

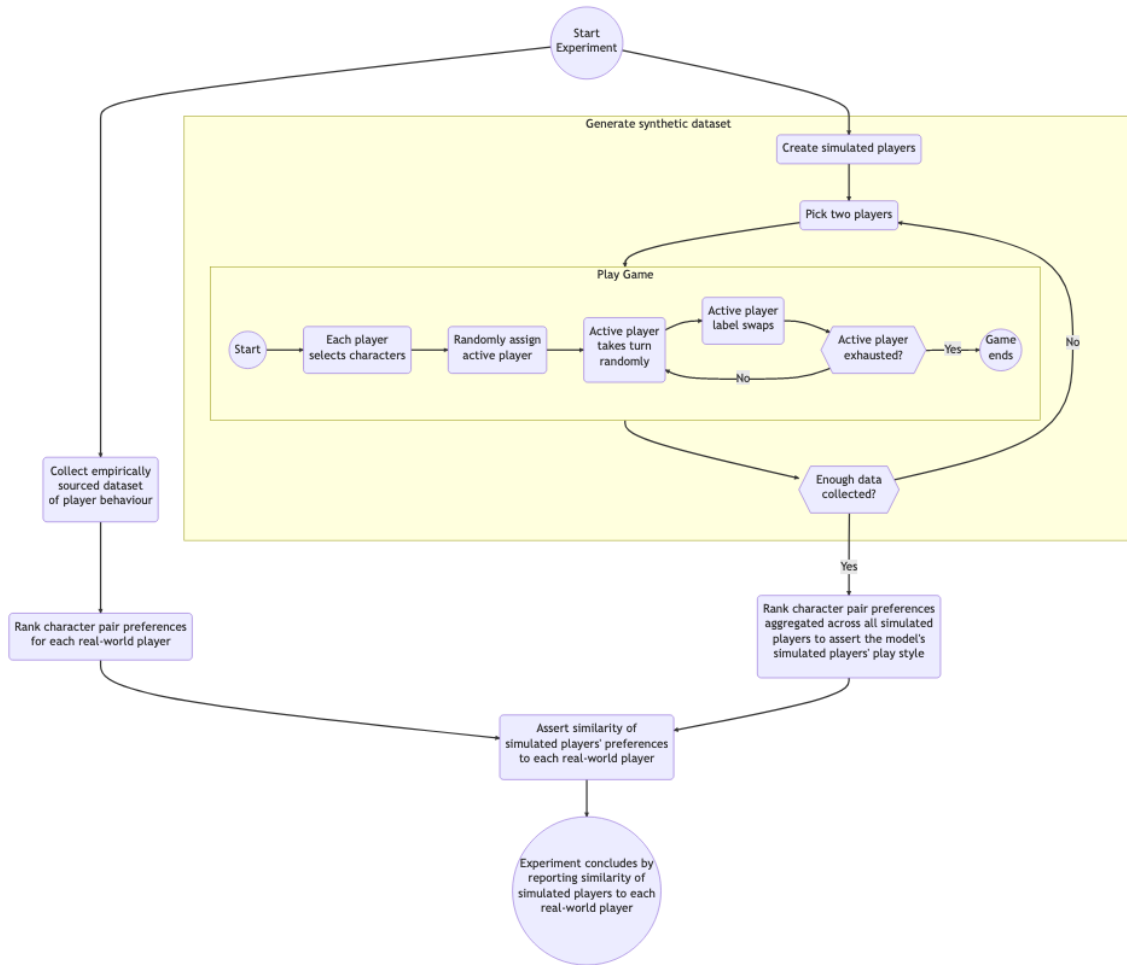


Figure 6.1: A flowchart diagramming a simple model of RPGLite play, without aspects applied.

player selects a random valid move to make.³ The active player alternates, and the process repeats, until such time as an active player starts their turn with both of their characters fully depleted of health. The player with remaining characters is the victor, and another game is started by picking random players and starting a game between them, until a predetermined number of games has been played.

After a sufficient number of games are played, analysis of the various datasets collected can begin, although this is discussed in more detail in section 6.5.

6.3 Experimental Design

The naive model above describes a simulation of RPGLite itself; however, the models exists to support investigation into research questions, and cannot do this alone. To answer our research questions, the model must be augmented so as to represent real-world play. This section describes the modifications made to the naive model (through aspects applied by PyDySoFu) and how data from those simulations can be analysed to arrive at answers to our research questions.

The questions these simulations seek to answer are:

1. Can we fit model details per-player to get realistic player behaviour?
2. Can we cluster based on accuracy of different models for each player?
3. Can we generate predictive data for unknown models from known ones?

These RQs are taken from an early outline of a thesis structure. Maybe they need rewording? To investigate these, we produce models of learning which can be applied to the naive model to augment player behaviour. Simulated players will have their behaviour augmented in relation to the character pairs they choose to play. Real-world players could reasonably be expected to choose character pairs which are stronger together over time. In addition, it is possible to calculate objectively optimal character pair selections in RPGLite [Kav21]; therefore, players should be expected to gravitate toward

³Many simulated decisions are random; this is because the model is designed to be naive, so it avoids informed decisions where possible. Informed decisions are expected to be woven later as aspects.

more optimal character pairs and away from poorer choices as they become more experienced in the game.

A few facets of gameplay are expected to be learned by players over time, and would theoretically also be suitable objects of simulated players' learning. One analog to character pair selection could be move selection: players would be expected to make increasingly optimal moves as they play. However, data can be muddled by players understanding the optimal use of some characters' moves better than others, as different characters invite different styles of play. Small changes to game states (the opposing team's composition, amount of health remaining, opponent play styles) could influence player behaviour also. Fundamentally, the state space of moves to be made and scenarios in which to make those moves is extremely large, making simulation computationally expensive. By comparison, character selection happens before gameplay starts, eliminating in-game factors which might influence player decisions.

Modelling the learning of character pair selection allows for modelling of players learning the metagame, rather than the game itself. Modelling learning of the metagame allows us to model players learning about the game and other players' interactions with it, rather than learning the game itself. An example of the distinction is the difference between learning how optimally make moves with a Barbarian in a team, and learning the mistakes other players commonly make when playing a Barbarian to exploit their weakness. As RPGLite is "solvable" — there is a provably correct way to play the game — players picking sub-optimal character pairs can be expected to observe that they are frequently beaten by better ones, increasingly favour pairings they identify as successful, and so converge on strong character pairs. Convergence en-masse would indicate a stable metagame. This was observed in practice in real-world RPGLite play [KM21]. However, while players made fewer mistakes over time when playing RPGLite, the average cost of a move to their probability of winning increased after many games (though decreased after a few). [KM21]. Player behaviour which converges on provably optimal play is simpler to represent than players' initial improvement followed by a seeming loss of performance. RPGLite Character selection is therefore a more convenient metric to model players' interaction with than move selection, as it presents a smaller state space for simulated players to explore while also providing a simpler expected end state for those simulated players to converge to.

Models of learning are therefore applied to players' character selections. Our primary goal in applying models of learning is to demonstrate whether players can be accurately modelled by augmenting

a naive model to represent them in particular, thereby answering the first research question. Different real-world players are expected to exhibit different styles of learning; thus, multiple models of learning will be applied. Different players might also learn differently, but in the same style; for example, two players might exhibit similar biases, but one could be quicker to learn from experience or another more cautious in the application of new knowledge.

To identify models of learning and parameters for those models which most realistically represent real-world players, we look to optimise the parameters for each model, for each player, by running multiple simulations of players learning in a particular manner (with particular parameters), identifying their preferred character pairings, and calculating the similarity of the simulated character pairings to that of a player in the real world. By doing so, we can anneal to a parameter which represents the optimally realistic model of a player learning in a given style. If the data produced by this optimally parameterised model of learning is not similar to empirically sourced data with statistical significance, we determine that the model of learning applied does not represent how a player learned in the real world: even the closest dataset the model produced fit its target dataset poorly. However, if a parameter can be found for which the model reliably produces character pair preferences which align with the real-world players' with statistical significance, then the model can be said to realistically represent that player's learning of the character pair metagame.

This section will first discuss models of learning generally in section 6.3.1, and will explain how aspects implementing these models can be produced in section 6.4. The strategy used to implement these models of learning is discussed in section 6.3.2. The design of experiments investigating the aforementioned research questions by applying these aspects is then described in section 6.5. **Rearrange either this paragraph of the sections it refers to so they're in order / make sense. At time of writing they refer to 6.3.1, then 6.4, then 6.3.2, then 6.5...!**

6.3.1 Models of Learning

Different people learn in different ways. Indeed, no universally-accepted definition of learning appears to exist. This is presumably because it is convenient to define what it means to learn differently in the context of different pieces of work. Cognitive models of learning can be useful when considering

mental processes specifically, for example, whereas functional models of learning could lend a more empirically applicable perspective. What it means to learn is clearly outwith the scope of this thesis; however, our experiments will include models of learning. To justify our model, we consider a functional approach to learning, as considering learning in this way appears more closely linked to the empirically focused work of modelling real-world behaviour than alternatives.

Lachman identify[Lac97] that standard definitions of learning along the lines of, “Learning refers to a relatively permanent change in behavior as a result of practice or experience” have practical shortcomings such as a focus on behavioural change (as learning may not change behaviour) or conflating learning’s process and its product (the process by which we learn is not obviously identical to its result, of which behavioural change is an example). They suggest learning might be better defined as:

“ (...)the process by which a relatively stable modification in stimulus-response relations is developed as a consequence of functional environmental interaction via the senses [...] rather than as a consequence of mere biological growth and development [Lac97].”

(...)

They note that their definition distinguishes learning from phenomena such as injury, changes to one’s maturity, or sensory adaptation, incorporates stimulus-response relationships the research community consider as learned, and differentiates learning’s process and product. Their model is inherently functional, making it useful for the purposes of simulation and modelling, although they offer only a definition of learning and a brief comparison to the standard textbook definition they introduce. The work presented is not intended to demonstrate its improved model of learning empirically, only to discuss its semantic merit. However, the models proposed in this thesis require only a theoretically informed, sound basis for their model of learning, and a lack of empirical justification is not a barrier to the relevance of the model Lachman have proposed.

De Houwer, Barnes-Holmes, and Moors propose a functional definition of learning which is primarily concerned with providing a definition of learning which is both accurate and useful for the purposes of cognitive learning research [DBM13]. Doing so attempts to provide a model around which some consensus can be reached; learning is a central concept in psychology, and they see their definition as

supportive of cognitive work. They introduce their definition as follows:

“Our definition consists of three components: (1) changes in the behavior of the organism, (2) a regularity in the environment of the organism, and (3) a causal relation between the regularity in the environment and the changes in behavior of the organism.”

This model of learning contains more nuance than the “textbook definitions” of learning they paraphrase as, “a change in behavior that is due to experience”, but does not stray far from the core concept: some environmental stimulus impacts behaviour in a causal fashion. Their introduction of “regularity” to their definition refers to the presence of the stimulus with some form of repetition, whether this be multiple instances of a stimulus at different times, or the same stimulus occurring concurrently. De Houwer, Barnes-Holmes, and Moors explain that such a model is straightforward without the sweeping inclusivity of the simple model mentioned earlier, and is easily verified (although, as in the work of Lachman [Lac97], empirical verification is omitted in favour of semantic analysis).

Aside from other benefits more particular to their research community, these benefits are especially useful from the perspective of modelling learning in our case. A simple, functional definition can be captured in a software model, and introduces few opportunities for misunderstanding or mis-application. It also introduces helpful concepts — such as regularity and causality — which the other definitions discussed do not. For the purposes of this thesis, we therefore adopt this definition as a basis for our model of learning.

6.3.2 Modelling Learning in RPGLite

We use De Houwer, Barnes-Holmes, and Moors’s definition of learning [DBM13] to arrive at a model of learning which is encodable in aspects that can be applied to our naive model of RPGLite. We are interested in modelling players learning a preference for character pairs over time. The model should therefore account for: how player behaviour is influenced in accordance with their learning, in line with the definition’s first criterion; repetition of experience in successive games influencing the direction of a player’s learning, in accordance with the definition’s second criterion; and do so in a causal manner, in accordance with the definition’s third criterion. We therefore look to model a causal

relationship between a player's observation of successful character pairs and their future choices of character pairs.

To fulfil these requirements, a model might draw on previously successful character pairs to determine future ones. There are many ways in which this could be done. For example, we could model learning as consistently playing the character pair which most recently was observed to win a game. Any game ending naturally determines a winning pair, and we can select this pair when playing future games, until a different pair is observed to win instead. However, this does not align with our expectations around how players would engage with a game in the real world. Having a strategy one is confident in and being unlucky with the game's random nature is unlikely to deter a player from what they believe is ideal. Indeed, we can expect players to understand that perfect play might not be winning play: in some games, the right moves might not lead to a successful outcome due to moves randomly missing opponent characters. Equally, players may take time to become confident in a strategy; we would expect a player to explore character choices before settling on a preferred pair early in their experience, and would expect very experienced players to choose characters based on what they have learned, rather than continuing to explore options for which they have sufficient information to reason about. We can infer that:

- There are scenarios where players can be expected to observe wins/losses without incurring behavioural change.
- Players' confidence in what they have learned can affect their inclination to draw on what they learn when making decisions.
- What players learn in successive games would have a small impact in their early experiences, but an increasingly significant impact proportional to their experience in the game.

The model of learning used to simulate players' improved play of RPLite can be explained following a similar structure: a component pertaining to observations players make about winning characters; a component pertaining to their inclination to use that knowledge when choosing characters (rather than exploring their options); and a component pertaining to the strength of that inclination in proportion to their experience.

The first — that players should make observations about winning players — can be implemented by modelling the choice as a probability mass function (PMF). The PMF maps character pairs to their chance of being selected by a player, and initially tracks all character pairs as having an equal chance of being selected. After every game, the chance of selecting the winning character pair increases, and the chance of selecting any other pair decreases. The sum of probabilities of being selected across all character pairs always sums to 1 (100%). This can be implemented as a record of the winning character pairs observed by a player: many ways of producing a PMF from a sequence of wins exist, but for the purposes of explanation, one such method is to take the proportion of wins for every character pair as their probability of being selected. This method produces a valid PMF because the probabilities must sum to 1 by virtue of the winning pairs being the sole determinant of the probabilities; every win contributing to the probability of a character pair being chosen means that 100% of character pairs are considered, and therefore the probabilities inferred from that data also sum to 1.

A model of learning where prior belief as to optimal character pair is causally related to those which were previously successful fulfils the first criterion of the model of learning proposed by Lachman [citelachman1997learning](#); however, players would be expected to explore a state space in the early stages of their experience. The model proposed by Lachman identifies that the experience of a learning agent draws from “regularity” in their environment; we therefore require that an experience affects behaviour not when it originally surfaces, but after repeated exposure to it. This fulfils the intuitive understanding of learning discussed earlier — players’ confidence in what they have learned affecting their inclination to act based on experience — our second criterion. We therefore require some model of confidence an agent has in what they have learned.

The third inference — that confidence should be low initially, and grow proportionally to experience — explains the “shape” of the confidence model. We require a monotonically increasing function mapping experience, quantifiable as games played, to confidence, being a percentage chance between 0% and 100% that a player is to determine their character choice based on what was learned, rather than on an exploration of their space of possible choices.

Such a model of confidence offers a model of learning when combined with a PMF representing historical observations of a character pair’s success. If the model of confidence indicates that the player is to explore their space of possible choices, they select character pairs randomly; if they are modelled

as confident enough that their behaviour is instead informed by their experience of their environment, they select a character pair according to the PMF their historical observations define. As the PMF affords higher probabilities to repeatedly winning character pairs, player behaviour is causally affected by the regularity of their experience, implementing a realistic model of learning in agreement with the functional model proposed by Lachman [Lac97].

Our model of learning is therefore complete, but a model of confidence remains to be defined.

6.3.3 Modelling Player Confidence in RPGLite

To model learning, we must find some function mapping experience (quantified by a count of games played) to some probability between 0 and 1, and fitting the criteria described in section 6.3.2.

A sigmoid curve fulfils both the second and third inferences in section 6.3.2. It also notionally conforms to expectations around “confidence” as an anthropomorphic trait: much like a sigmoid, confidence starts low, and remains so until it achieves some inflection point, after which it grows rapidly until tapering asymptotically toward some maximum. However, not all real-world players might express the same traits in their growth of confidence; the shape of their sigmoidal confidence may differ. In order to account for this while guaranteeing a monotonically increasing mapping of the number of games played to a player’s confidence, we select a sigmoid which can be parameterised to alter their shape.

A sigmoidal curve is suitable for this model of learning where other curves would not suffice, because we require a period where players lack confidence and explore their options. Sigmoid curves such as the logistic [Ver45] or Gompertz [Gom15] are widely used when modelling systems [WJ97], but while they fulfil the role of a monotonically increasing curve with asymptotically low and high initial and final states, the shape of such curves is not trivially modified to fit different players’ learning styles.

More flexible asymptotic curves were developed by RICHARDS [RIC59] drawing on growth curves developed by Von Bertalanffy [Von38], which afford a natural pattern of growth; RICHARDS amends this curve to offer a parameterised growth rate. This curve can be made equivalent to other curves, including the logistic and Gompertz [FT+84]. This curve allows for a parameterised rate of growth,

but lacks parameters controlling the points at which growth occurs most rapidly. The relative rate of confidence gain is a separate concern to the point at which such growth occurs: a player might cautiously grow in their confidence until they are already very experienced, or might bullishly grow in confidence yet plateau early, taking longer to reach complete confidence in themselves than they did to garner an initial increase, regardless of their relative growth in confidence.

The flexibility of a parameterised relative growth rate appeals to the notion that different players would gain confidence at different rates, but the point at which confidence accelerate most must also be controlled. We therefore employ the Birch curve, proposed by Birch [Bir99] for its increased flexibility as compared to the Richards curve combined with its additional parameter used to control the curve’s shape.⁴ Different players might exhibit different rates of growth in their confidence, and might grow maximally in their confidence at different points in their experience; the birch curve satisfies these properties as a model of confidence. Investigate: did we ever model confidence with anything other than $c=1$, i.e. a standard logistic curve?

6.4 Aspects Applied

Having a naive model of player interaction in RPGLite allows for the generation of a control dataset, which can be compared to aspect-applied datasets to examine whether the aspect-applied models are provably “closer” to empirically-sourced data than the naive dataset⁵. To generate the experimental dataset, the naive model producing our control is augmented through the application of aspects. These aspects fulfil different functions, and can broadly be grouped into three categories:

1. Aspects implementing behavioural models, changing the behaviour of simulated players;
2. Aspects instrumenting the naive model to perform observations necessary for the implementation of behavioural models (such as models of learning) on the naive model;
3. Aspects altering the behaviour of modelled players to simplify experimental observation, and to handle exceptions introduced by changes within the aspect-applied model.

⁴Birch refers to shape to mean the point of inflection of a curve. The point of inflection of an exponential rise to a limit is at its initial point; the point of inflection of the logistic curve is in the exact midpoint of the curve’s growth.

⁵“Closer” here refers to a similarity measurement which we will define later in this chapter.

The first set of aspects implement the behavioural change we anticipate will produce datasets closer in similarity to empirically sourced data than that sourced by the naive model. The second set lays the foundation for applying these changes. The third applies behavioural changes which experimental observations require. Necessary changes unrelated to our models of learning will first be explained in section 6.4.1, followed by an introduction of foundational observations for our models of learning in section 6.4.3. Finally, the models of learning built on those foundations are introduced in section 6.4.2. The ordering of this explanation of different uses of aspects is confusing; once the aspects are explained individually, reorder the sections, rework the introduction to the structure of this section, and ensure the outline of the section reflects the changes made.

A diagram of a game of RPGLite with all possible aspects applied⁶ is presented in fig. 6.2, by way of high-level overview of the aspects this section introduces.

6.4.1 Aspects for model improvement

Ensuring the Best Move is Played

As discussed in section 6.3, the goal within this experiment is to build a model of players' learning of character selection rather than move selection. However, randomly selected moves are liable to place players in unrealistically weak positions, as players are unlikely to make obviously poor moves such as skipping a turn with no clear reason. It would be unexpected to see true randomness in players' moves. This is a concern for the modelling of players' character selection, because the model of learning defined requires a causal relationship between what is observed (in this case characters which most reliably win games) and behavioural change (players choosing character pairs proportionally to their estimated chance of winning a game); if selecting random moves cause simulated players to lose games when they would have realistically won them, this would affect character selection by definition.

Character selection and move selection are therefore somewhat linked. However, move selection is most frequently optimal: in the majority of cases real-world players chose the best move available to them, as analysed by Kavanagh and Miller [KM21] :

⁶Including aspects which should not be woven in the same experimental run: some aspects implement different models of learning and are therefore conceptually incompatible.




Figure 6.2: A flowchart describing a simulated game of RPGLite, and all aspects woven into the game to implement the various models of learning. Some aspects should not be woven together in the same experimental run, as they implement different models of learning.

“ (...)the majority of actions taken were optimal, with a cost of 0.0. In total 73% of the player’s season 1 actions and 77.8% of their season 2 actions were optimal.^a”

^aDifferent seasons of RPGLite are discussed in more depth in [Make a reference to the chapter\(s\) where I explain different seasons of RPGLite.](#)

(...)

Therefore, a naive model of move selection which represents real-world move choices in $\frac{3}{4}$ of cases is to select the best move at every opportunity. We therefore apply an aspect to the function responsible for move selection which performs a lookup on the dataset of action costs defined by Kavanagh and Miller [KM21], and select the known-optimal move in every case.

One additional factor in move selection is handled by this aspect. As simulated players’ behaviours are augmented to select universally optimal moves, an anomaly in the dataset was identified. In some games, optimal play would result in infinite loops: players would arrive in a state where they would skip turns mutually and indefinitely. This is because it is possible that, in RPGLite, to make any move is disadvantageous.⁷ Both players can concurrently exist in this state. Therefore, both players’ optimal move is to skip their turn. Kavanagh and Miller note that such states were reached in 64 cases in real-world play, but real-world players never skipped their turn; therefore, this scenario is also an example where suboptimal play can be re-introduced. This aspect therefore returns random moves if the previous two moves in a game were to skip. Such a state indicates that the game reached a stable point where neither player would optimally deal damage to the other. Random moves destabilise this state and return to ordinary play.

A model of move selection more closely aligned with expectations of player behaviour is beyond the scope of this thesis: a demonstration that models of both move selection and character selection can be optimised concurrently to arrive at a realistic model of empirical play with multiple aspectually applied models of behavioural variance invites itself as future work. [Do I need code snippets for notes on aspects?](#)

⁷This is because barbarians deal more damage once they lose a sufficient number of hit points. As a result, to deal damage to a barbarian can result in their having an opportunity to win in one move rather than two in certain health states.

Handle Game States with no Viable Moves

A consequence of playing games using only optimal moves was that unexpected states could occur. One might assume that, choosing moves optimally from a dataset of available moves in every game state it would be infeasible to arrive at states not present in this dataset. However: the dataset was produced through model checking, with the aim of identifying the cost of an action regards its impact on a player's chance of winning. If a loss is guaranteed, all moves have a 0% chance of winning; the model checker producing the dataset of available moves therefore identifies that the game is a foregone conclusion, and contains no valid moves.

In such a situation, identifying the winning player is as simple as identifying the player who made the previous move. We therefore apply an aspect which handles exceptions raised by move selection (as lookups will fail due to the game's current move state not existing in the dataset of action costs) and, after identifying that the exception indicates that the state does not exist (in this case, a `KeyError`), the aspect handles the exception by assigning the losing player's characters 0 health and swapping active players so that the simulated game proceeds with the losing player taking the following turn. As this turn starts with the active player having 0 health, the game ends as expected.

6.4.2 Aspects for Instrumentation

As explained in section 6.4, the experiment being discussed requires observational apparatus to be implemented in order to collect data about the experiment for later analysis. Ordinary experimental setups might require this apparatus to be tangled through the experimental codebase; however, aspect orientation provides a technique for augmenting the simulation to make observations before, during, or after any step of within the model. This demonstrates observational scaffolding as a cross-cutting concern as suggested by Gulyás and Kozsik [GK99]. Aspects realising this concept and supplying observations and additional models in support of our model of learning are discussed here.

Update Model of Confidence

The models of learning described in section 6.4.3 have players make observations as to the character pairs which most frequently win games, and have those players alter their behaviour in response to these observations. However, as discussed in section 6.3.3, a model of confidence is required to complete our model of learning. This model of confidence is a sigmoid curve, and the Birch curve was selected for our model of confidence. The role of this aspect is to update the birch curve when a game ends.

An aspect updating our model of confidence is applied to act after a game ends (an “encore” in PyDySoFu’s parlance), and updates a value in the model’s environment for confidence models. Confidence is represented as a mapping of player objects to floating point values between 0 and 1. The aspect supports modelling confidence in many ways, identified by a `sigmoid_type` parameter within the model’s environment. OH NO. I just realised our models don’t use birch at all, they use a logistic curve. They also use a relative growth rate on the logistic curve that I’m unsure of — not sure why I implemented it the way I did. Getting increasingly close to just re-running my models. I could run them again using birch, and maybe anneal to a value for the curve shape for different players too...?! THIS NEEDS TO BE FIXED ASAP.

1. A logistic curve

Record Prior Distribution of Character Preferences

Keeping this in because I might need to explain the aspect, but it’s only used in the implementation of choosing based on prior distribution, which I don’t believe is ever applied. So! Decide whether to get rid of the writeup of that model. If I am, then get rid of this too. If not, add it to the learning models writeup below and explain it here.

Record Character Pair Choices

It is important to keep track of the characters chosen by simulated players, as this data is required to analyse the outcome of the experiment Do we describe an experiment singular or experiments plural in chapter 6? outlined in this chapter. However, the naive model makes no special accommodation for

the collection of character choices made by players. Character choices could be calculated by iterating through the list of completed games after a simulation is complete, filtering for games relating to a specific player. However, the collection of character selection data can be simplified by tracking what was played at the end of a game. This is achieved with with an aspect which observes the character pairs played in a completed game, and appends each players' chosen pairs to their own list. This aspect is applied after its join point executes (an “encore” in PyDySoFu parlance [reference pdsf's explanatory chapter / re-engineering chapter here?](#) [unsure which is most appropriate / introduces the concepts first.](#)).

There are many ways to collect this data, aspectually or otherwise; however, the simplicity of collecting information mid-process without requiring modification of a base model demonstrates the flexibility of an aspectual augmentation of models and simulations as an approach. Collection of chosen character pairs mid-simulation is an opportune example of this convenience.

Track Detailed Outcomes of Games

Similarly to the recording of character pair choices, the outcomes of games must be tracked. This information is important as an input of our model of learning, as successes and failures observed by players for each character pair can inform their models of learning. Some models of learning implemented, such as one including a hyperbolic decay bias (to be explained later, in section 6.4.3), require not only knowledge of the proportion of wins each character pair has achieved, but also the history of wins and losses: players might be biased toward recent win/loss observations, discounting older experiences in favour of the new. Detailed historical win/loss information must therefore be recorded.

As with character pair choices, the naive model was not engineered with the intention of providing this information specifically. However, the model is easily instrumented to collect such information at many suitable points. As with recording a player's chosen character pairs, an “encore” aspect was implemented which records wins and losses observed for character pairs on game end. Specifically, the aspect models players observing character pairs which won at the end of a game, lost at the end of a game, and also pairs which were uninvolved in the game. This requires a list of outcomes (wins, losses,

and neither) for every player, for every character pair. The lists of observed outcomes for a given character pair record True for a winning pair, False for a losing pair, and None for a character which was not involved in the game's outcome. Each player observes the relevant state for every character at the end of any game they play. Provisions were made in the implementation of this aspect for a player only paying attention to their own outcomes (i.e. only recording whether their own pair won or lost, without consideration for their opponent's outcome), though in practice all simulations were performed with players observing both their own outcomes and those of their opponents.

Record Winning Pair observed by Players on Game End

Unlike the detailed game outcome data recorded by the aspect described in section 6.4.2 and used in advanced learning models such as that of hyperbolic decay (to be explained later, in section 6.4.3), other learning models required only simple outcome data. The simpler model of learning making use of a distribution of winning teams to be explained in section 6.4.3 is more easily implemented against a simple list of winning teams. Rather than coalescing the more complex dataset collected as described in section 6.4.2 to achieve the desired format, another aspect can be applied to collect the data in a simpler format directly.

An aspect was implemented which also models players observing winning character pairs on game end, but collects the data in a simple list of character pairs observed to win games, rather than a mapping of character pairs to game end states as described in section 6.4.2. Character pairs which lose games or are not involved are not recorded for the purpose of the alternative model of learning.

6.4.3 Aspects Implementing Models of Learning

The aspects described above implement setup for the models of learning applied in this thesis, both instrumenting the underlying RPGLite model for observation and tweaking behaviours of the naive model. Two aspects remain unexplained, implementing the learning models themselves.

The first model of learning left to be described is implemented as described in section 6.3.1, drawing from a record of character pairs observed to win in earlier games to select characters in future games, in accordance with a model of confidence which activates either this behaviour or a randomly selected

pair. A detailed description follows in section 6.4.3.

The second model of learning left to be described implements a similar logic, but introduces a weighting on the history of winning pairs a player observes, introducing a hyperbolic discounting bias to historical observations. In this model, players increasingly discount old observations in favour of recent ones. A model of learning with a hyperbolic discounting bias is described in section 6.5.

Why didn't we make a bayesian model of learning? Should we have? Would this be difficult at all?

Character Selection using confidence sigmoid

The first model of learning applied through aspect orientation draws on character pairs a simulated player observed to have won games (as recorded by the relevant instrumenting aspect, discussed in section 6.4.2). The record of previously winning character pairs defines a probability mass function by selecting a character with equal probability to their rate of appearance in the history of winning characters.

Selecting a character based on this distribution is gated by a confidence model as described in section 6.3.3. If a simulated player's confidence model indicates insufficient experience to found decisions on, players will instead select characters randomly, effectively exploring their space of possible choices. By doing so, players have the time and opportunity to observe many matchups between different character pairs. Time to observe matchups is important because some character pairs may only present a strong choice if played against specific alternatives. We can imagine a character pair which is extremely effective against 50% of pairs, but extremely likely to lose games played against the remaining 50% of possible opponents. Without exploring possible matchups, a player may lack observations which would inform them about whether a character pair is effective in general, or in a narrow set of circumstances. A confidence model encouraging early exploration promotes the experience of a wide variety of matchups, avoiding this issue.

Another concern early exploration mitigates is that simulated players need opportunities to build well-rounded priors: if a character pair is never selected by any player, experienced players basing their character pair choices on those which previously win games cannot include pairs which were never selected, because they will never have had the opportunity to win games at all.

$$score(D) = \frac{1}{1+kD}$$

Figure 6.3: Hyperbolic discounting score calculation

This aspect was implemented as an “around” aspect in PyDySoFu’s parlance, allowing it to apply additional logic before and after its join point. It could equally appropriately have been implemented as an “encore”, applying logic only after its join point had executed; the aspect discards the random selection made by the base model unless the simulated player lacks confidence. Either implementation can return a different character pair choice to the join point’s caller, allowing the chosen character pair to effectively be overridden, and so allowing the aspect-applied RPGLite model to proceed executing as it would otherwise; the only difference being the pair of characters “chosen” by the simulated player whose behaviour was augmented.

Character Selection exhibiting hyperbolic decay

Character selection with hyperbolic decay works similarly to the simple model of learning as described in section 6.4.3. Character pairs are selected randomly if a player lacks confidence or, if the confidence model applied to the player indicates that character pairs are to be chosen according to their historical observations of winning pairs, their record of winning pairs in previous games informs their choices instead. However, this model also applies a hyperbolic discounting bias which weights the player’s historical observations when making choices. *Maybe we should have also produced models of other discounting mechanisms, like exponential discounting (which I believe is actually the logically sound way to discount). Would be interesting to see whether some players are well-modelled by other discounting mechanisms. Investigate this!*

Hyperbolic discounting is a bias studied in behavioural economics. *Find citations explaining hyperbolic discounting*

A hyperbolically discounted score for an event after a delay D with discounting factor k is calculated by:

This score can be applied to a history of games played with different character pairs to weight that

history; this weighting can then be used to create a new PMF to draw character pair choices from, by allocating each pair a probability of being chosen according to its share of the weighted history of winning pairs. Where the previous method for selecting character pairs selected each pair with the frequency they appear in the record of winning pairs, this method weights the record of winning pairs; the sum weighted history of a character pair's wins as a proportion of the total weighted history for all pairs becomes the PMF for character pair selection. *Maybe there's need here for a diagram or some extra equations to make this clearer...?*

With this scoring mechanism, our model of learning incorporates a behavioural bias which some real-world players may exhibit. Player character choices can now be modelled through different behavioural variations applied to the naive model. The resulting behaviours from each model of learning can be compared to that exhibited by the naive model to discern which produces simulated RPGLite play more similar to a given real-world player's; methods for doing so are explored in the remainder of this chapter.

6.5 Experimental Design using Aspectual Learning Models

The models of learning described in section 6.4.3 can be applied to augment the naive model described in section 6.2 with learning behaviours. However, answering the research questions proposed in this thesis requires assessing whether the aspect-applied model more closely resembles real-world play than the naive model.

This section presents the strategy for comparing simulated and real-world play, and discusses the measures taken to ensure an integral analysis of the datasets produced, both in terms of mitigation of the impact of RPGLite's inherent randomness, and in terms of the statistical significance of the results obtained.

6.5.1 Quantifying Similarity of Character Pair Selection

The experiments presented in this chapter make comparisons of many synthetic datasets to a single empirically sourced dataset; that of a player whose learning of RPGLite we aim to simulate.

The comparisons are required to determine which of the synthetic datasets most closely matches the empirical dataset. We are specifically interested in determining whether players played with preference toward similar character pairs. The distribution of their choices describes a bias toward some pairs, and away from others. The distribution of the naive model, being randomly chosen, should exhibit no bias; a simple null hypothesis is therefore that a simulated dataset is as similar to real-world play as a randomly generated one; there exists no bias comparable to that exhibited by the empirical dataset in this case. Our hypothesis in answering the first research question posed is that, with models of learning applied through aspect orientation, the naive model can exhibit biases toward character pairs similar to that of a real-world dataset. The simulated play would therefore exhibit a more “realistic” behaviour in character selection than the naive model as a result of the aspects applied.

We expect discrepancies in the specific character pairs selected in games (such as the order of character selection) because of the random nature of character selection. However, it is feasible to assess the similarity between the distribution of character pairs chosen by simulated and real-world players. The distribution of a player’s chosen character pairs describes the preferences of that player by ranking their selections. We can therefore examine the similarities between a player’s character pair preferences, assessed by their rankings, and the preferences of simulated players to arrive at a measurement of similarity between the datasets. We therefore seek to quantify similarity of character pair selection by measuring the distribution’s rank correlation. *Should I explain rank correlation more where it’s first introduced...?*

Many rank correlation measurements exist; *Briefly discuss the different rank correlation measurements I could have chosen.*

The strength of a player’s preference toward different character pairs is a less useful measure of similarity for the purposes of this experiment than the concordance of that ordering. Players having a strong preference for a subset of character pairs will leave others rarely played, if ever. With a degree of randomness to the choices of character pairs, this may mean that a random selection of a character pair slightly more frequently in one dataset than another might — with a slight boost to its rate of choice — significantly increase its ranking in the distribution of character pair choices. This is because rarely chosen pairs will have their rates of choice dwarfed by a player’s preference, and so their relative rates of choice are likely to be similar. This impacts our choice of rank correlation.

Stopped discussing rank correlation here because I may re-run my experiments; if I do, I don't want to have to re-write this, because frankly it's been pretty tricky to write so far and I'm keen not to have to redo it. Maybe there are other correlation metrics we could pick instead; Somers' D looks promising...? Are there citations I could use that compare different measurements, so I could justify my choice, and maybe pick one that's better than τ if it exists?

6.5.2 Annealing toward Player-Specific Learning Parameters

Discussion required here around the mitigation of randomness — it's required because of RPGLite's inherent RNG — and also explaining how we make use of a grid search in combination with k-fold validation to anneal toward well-fitted parameters of our models of learning for each individual player.

6.6 Experimental Results

Presentation of the results owing from the experiment as described in section 6.5, and evaluation of the research question with respect to these findings.

Again, no subsec likely required here.

6.7 Discussion

A closing discussion on what we found, and how the research question was answered.

7 | A chapter title here for the experiment moving aspects to new systems, or systems with some changes

This is a relatively short chapter; a lot of the building blocks for it exist in the previous chapter, so there's less ground to cover. If it ends up quite lop-sided, I'd chop the earlier chapter in two rather than artificially making this chapter beefier; I think it'd flow better.

7.1 (Reword) Experimental motivations / motivation of research question

This section should describe why it's interesting to move an aspect trained on some actors' behaviours to a new system, and discuss what investigating this can teach us that we don't already know from the previous experimental chapter. We've got aspects which represent behavioural variance. If the system changes, can we expect that these aspects still apply to the new system? Is the representation of behavioural variance in this model separable from the system the behaviour occurs within?

7.1.1 Coupling of Model and Behaviour

Explaining the issues of behaviour coupled to models. Highly related to some of the reviewed literature — I forget exactly what — which discussed whether aspects which were designed to be applied to one system could feasibly be transferred to other systems, or whether they're inherently aware of the system they're originally designed for. The core concept is that to be making changes to some underlying codebase, you probably have to know what that codebase is, in most cases at least.

7.1.2 (Reword) Gaps/opportunities left by previous experiment

In the previous experiment we demonstrated that behavioural variance can be plausibly realistic. Are those variations separable from their underlying model, or — when trained i.e. made realistic — do they suffer from the coupling discussed in the previous subsection?

7.1.3 Research Question

Whatever the specific research question's wording for this was. Something about decoupling realistic aspects from a given model maybe? It should be in an earlier chapter somewhere.

7.2 Experimental Design

We took old player data and trained aspects on them, and in the previous experiment we found they were statistically significantly accurate. We ran a second season of RPGLite with slightly different parameters on player data. We modelled how players learned and variations on their learning patterns, so in theory, we should be able to apply the new learning patterns to the other system too. Do we have to re-train the aspects? How portable are they? This section lays out the design of this experiment.

7.2.1 Changes to RPGLite

What changes did we make to our system?

7.2.2 Applying Behavioural Variations to New System

Layout of new experiment, how it'll work, what's measured, why it should answer the RQ.

7.3 Applying Aspects from a Control System to a New System

Our implementation & results from the experiment described above. Discusses how the above experiment was realised, lays out the results we found, and relates those results to the research question we started with.

7.3.1 Implementation

Implementation of the experiment

7.3.2 Results

Presentation of the results of the experiment, analysis, some discussion (more in next section too)

7.4 Discussion

Some notes discussing the outcome of this research question (that trained aspects representing behavioural variance aren't separable from the system they're trained around)

8 | Future Work

The focus of this thesis is to develop a state-of-the-art aspect-oriented framework, to produce a suitable experimental environment to demonstrate its effectiveness, and to use that environment to investigate whether aspect orientation can be used to augment models with behavioural variance. *is suitable for simulation purposes? It's about showing that we can use aspect orientation appropriately in simulation environments, and that aspect orientation can also lead us to realistic and nuanced simulations, too. Go back through the pdf and lit review chapters to make this argument properly.*

. We have found that aspect orientation can be used to create realistic and nuanced simulations, have successfully produced a well-constrained environment to simulate for testing purposes, and have produced a novel aspect orientation framework which demonstrates novel and powerful weaving concepts; lots of opportunities for research outwith this thesis' scope present themselves.

This chapter describes some possibilities for the presented research to be extended in the future.

Add a RQ at the end of each research opportunity; this can and should be a contribution of the thesis, and the RQs would help to sell it as such.

8.1 Aspect-Oriented Metaprogramming in real-world Software Engineering

The combination of metaprogramming and aspect-orientation introduces powerful new possibilities in the realm of aspect-orientation. In traditional aspect-oriented work, aspects treat their targets as black boxes; this leads to some limitations, which aspect-oriented metaprogramming is well positioned

to address.

Traditional aspects cannot add their behavioural modifications interspersed within the work being done by their target. The “textbook” use-case for aspect-orientation is logging: aspects can separate logging from the business logic they are applied to. However, a programmer in mainstream programming paradigms may wish to insert logging behaviour within their business logic, rather than around it. Aspect-oriented metaprogramming makes this possible, as the target can have logging logic interspersed through otherwise decoupled business logic when woven. To achieve this end goal without “within”-style aspect application would require a refactoring of business logic to create join points which traditional aspects could apply against; this breaks the “obliviousness” property of the aspect orientated philosophy as discussed in chapter 2.

8.2 Augmentation of pre-existing models

Myriad models exist which have been provided accurate results in past research, but where unexpected events such as financial collapses, pandemics, or unpredictable weather events cause one-off shifts in real-world data which cannot be accounted for. For example, models of world health over time could not account for the Covid19 pandemic, and models of the world economy could not incorporate real-world data from the recession caused by responses to the pandemic, or the 2008 financial crisis. The World3 model is an example of one which has provided accurate predictions for decades [Find citations for World3 model](#), but could not account for unexpected incidents when constructed.

Models are often initialised with data describing a system’s initial state, which a model computes on to produce predictions. Any output produced could only represent the impact of a large, unpredictable event within the system as changes to data or process; changing either for the entire model risks compromising a model’s predictive quality for the timespan prior to an event. The modification must therefore be represented as a special case within the model, leading to ancillary code inline with the model’s core logic in a similar manner to the scattered and tangled code discussed in section 2.1, which aspect orientation was initially designed to alleviate. It therefore has similar properties to a cross-cutting concern [\[Kic+97\]](#) and may be suitable for factoring out of a model’s logic and into advice.

Intuitively, this research would aim to show that an aspect could “nudge” a system’s state in-line with real-world data when exceptional real-world circumstances affect a model, but cannot elegantly be represented within the model’s logic. A relevant research question for future work to investigate is:

“ Can aspect orientation be used to introduce special cases to real-world systems to correct a model’s predictions in when a system under research is altered by a freak event?”

lorum ipsum

“Freak event”
might be too
informal here
but I’m not
sure what
to replace it
with.

8.3 Aspect Orientation’s utility for Research Software Engineers

As a corollary of the research opportunities for aspect orientation’s use in software engineering discussed in section 8.1, there is an opportunity for research software engineers to also benefit from the adoption of aspect oriented programming. However, while aspect orientation’s use in industrial software engineering has drawn criticism [Ste06; Prz10; CSS04]¹ its use within research codebases is a special case where it may be more suitable, particularly considering the possibilities discussed through the rest of this chapter. Revisit this paragraph. Is this too wordy / ramby / unclear? Can I tighten it up at all?

The resource constraints and stringent requirements for accuracy in research codebases present challenges which augmentation through aspects could assist with. Compensating for a special case in a pre-existing model would require maintenance of the codebase, which takes time and could inadvertently alter its behaviour. Time is a scarce resource in research environments, and undesired changes to a model’s behaviour can invalidate research results. An alternative to adjusting the models directly is to construct aspects which represent large and unpredictable events in real-world systems such as pandemics, economic crises, war and famine. These can be modelled on real-world data for accuracy, which can produce realistic simulations as this thesis demonstrates. This use of aspect orientation in simulation and modelling can be investigated by creating a proof-of-concept of the approach as applied to pre-existing models.² Studies can also be conducted to investigate whether an

¹A review of aspect orientation’s critique is presented in ??.

²This approach is similar to PyDySoFu’s initial proof-of-concept study [WS18a] as discussed in chapter 3.

aspectually-augmented model is quicker to construct and easier to maintain in future than a codebase with “patches” written into its original logic.

Researchers investigating this technique’s application to existing models could also investigate the difficulties of augmenting a model created with no intention of weaving advice in the future. The construction of advice requires appropriate join points to be specified, and codebases which are structured in a way which doesn’t yield convenient join points might be more complex to augment aspectually. These cases raise another use-case of PyDySoFu’s “within”-style weaving through runtime metaprogramming: where other aspect orientation frameworks force aspects to treat the targets they are invoked on as black-boxes, PyDySoFu can make modifications within them. PyDySoFu’s new kinds of aspects may therefore be more useful than those in other frameworks for researchers responsible for maintaining codebases written with no consideration for join-point specification, and would therefore be more universally applicable; this possibility requires future investigation.

If the research described successfully shows that aspectually-augmented simulations are easier and quicker for an RSE to maintain and deliver than direct maintenance of the model’s codebase, then

augmentation of existing models to improve their accuracy can follow in the community.

If I’ve got time, I’d love to push my codebase augmenting a World3 model to achieve just this with predictive implications for the 2008 financial crisis and the covid pandemic.

8.4 Hypothesising Possible System Dynamics via Aspects

Unpredictable events can cause discrepancies between simulated system state at a given time and the real-world system it models, as discussed in ???. While this technique presents promising research opportunities, researchers face other kind of model uncertainty which aspect orientation could also counter.

In some scenarios, it is difficult or impossible to make predictions about a system’s future states because its dynamics are actively being researched. Standard scientific practice is to create a model — mathematical, computational, or otherwise — to create synthetic datasets which indicate accuracy if their predictions align with what is empirically observed. **Cite popper, any point including studies that do this too? It’s fairly table stakes as a concept...** Some aspects of the system under study may be well-understood.

Rather than creating models of an existing system which encode its hypothesised behaviour, a naive model can be created which operates as the scientific consensus understands it. Hypothesised behaviour takes the form of aspects altering behaviour in any manner the hypothesis requires. Within-style aspects allow arbitrary modifications, to simulated behaviour, increasing the flexibility of the technique. Data produced by each model can be compared to empirically sourced data, and their similarity quantified, as demonstrated in chapter 6 and chapter 7. Our null hypothesis is that the naive model’s similarity to empirically sourced data is greater than that of the aspectually augmented model; our experiment’s hypothesis is that the behavioural change applied as aspects is more representative of the system under study than that of the community consensus.

This technique has a satisfying property: the hypothesis in a given experiment is clearly described by its aspectual representation, and if an experiment is successful — meaning its hypothesised behaviour accurately represents the real-world system — then the model adopted in future research as that of the community consensus can be the aspectually-augmented one produced by the original study. In this way, the scientific process is directly represented in the structure of the codebase, and the community’s progression to increasingly accurate models of a system is represented by the progressive adoption of “patches” to an original theory.

Maybe too informal, but I don’t know, maybe live a little...?

Hypotheses can also be created compositionally in this model. Researchers might develop a series of potential system properties or behaviours, but are unable to investigate all reasonable combinations in a timely manner. However, sets of aspects representing each can be composed to produce, for N hypothesised behaviours, 2^N potential behaviours which can be compared to empirical datasets to identify which combination of possible behaviours most closely resembles that of the real-world system under study. The technique is similar to modern pharmaceutical development, where a series of algorithmically identified compounds can be produced and tested to identify which are most suitable to treat a medical issue [find a citation for modern algorithmic / robotic drug development](#), but can be applied to any models where aspectual augmentation is appropriate.

This technique for developing experimental model codebases has another desirable property: it unifies seemingly incompatible philosophies of the scientific process. Kuhn [Kuh12] explains the scientific process as inherently social: it starts with a paradigm which is accepted as broadly true, and accumulates an increasing number of exceptions until the paradigm itself is deemed unfit, and a new

This sub-sec is already quite long, and including this point makes it significantly longer. Maybe that’s not a prob; maybe this

basis for a field’s research is adopted. Popper [Pop13] explains the scientific process as an approximation towards truth, with incremental progress made with each result achieved by a research community. The proposed technique for developing experimental model codebases demonstrates features of both. Consider the original model a paradigm initially selected by community consensus, and the application of aspects Popper’s incremental movements toward truth (as successful experiments are conducted) or Kuhn’s exceptions to the agreed model (as experiments identify weaknesses in the original model). In this case, each successive new model adopted by a community is adopted in a Popplarian manner: improvements are objectively measured, incremental, and would trend towards truth as a model’s behaviour fits empirical observations increasingly closely. However: over a sufficient period of time, the incremental patching of an original model would produce an accepted community model which contains a relatively large amount of discovery and complexity encoded in aspects, as compared to the original model they are applied to. One would expect the research community to rewrite the base model to simplify future aspect application and to more elegantly encode recent research findings; effectively discarding the original paradigm in favour of a new one. This process is Kuhn’s “paradigm shift”, where paradigms are dropped once a generation of researchers determine that an originally accepted theory on a topic is unfit for purpose as evidenced by mounting exceptions in the literature; a new paradigm is to be accepted by the community, as a new base model would have to be written and adopted.

The relevant philosophy of science is more nuanced than its brief explanation here, and the suitability of the approach for the development of research codebases is to be investigated; the work involved is outwith the scope of this thesis, but is suggested as future work. A basis for the incremental improvement of models via aspects is effectively demonstrated in chapter 6 and chapter 7, but the feasibility of the approach as a basis of a community’s scientific process and relation to philosophy of science warrants further investigation.

8.5 Standards for Aspect Orientation in Research Codebases

The possibility of a research community sharing models containing aspectual augmentation and possibly developing additional aspects to augment a model further involves a considerable amount of

model logic written within as advice. The community developing these aspects have the responsibilities of maintaining a codebase, but the added complexity that research software engineering introduces. These codebases may be used for many years, and may be iterated in in a series of future experiments. The legibility of these codebases and their long-term maintenance are areas of criticism in the software engineering community [Ste06; Prz10; CSS04]. The research community must therefore mitigate these weaknesses of the aspect-oriented paradigm when adopting the techniques discussed in this chapter for simulation and modelling.

To address the concern of the visibility of advice being woven, researchers may already take advantage of :

1. Improvements to tooling produced by the aspect orientation research community, including IDE integration [CCK03] and runtime inspection [MR02], should make clear to engineers what advice is being woven in a codebase and assist with debugging aspect-oriented programs respectively.
2. Aspect-orientation frameworks specifically designed to clarify to an engineer the aspects being woven should allow for less friction on the part of a maintainer who inherits a codebase and must reason about its behaviour. This is particularly important if the maintainer aims to weave more aspects into the codebase, and so must understand its existing behaviour before augmenting it further. Adopting weaving patterns such as import hook weaving — described in chapter 4 — should make a program clearer to a developer regardless of the tooling they have access to.

I've made this an enumerated list; is it better as a paragraph?

The impact of framework design on a codebase's maintenance should assist a developer even in the absence of tooling, but the success of import hook weaving in this regard is untested. An appropriate research question which arises is therefore:

“ Do aspect orientation frameworks with weaving techniques designed to simplify a developer's understanding of a program affect a codebase's long-term maintainability?”

8.6 Standard Aspect-Oriented Model Features

Researchers who build aspect-oriented models and extend others’ aspect-oriented codebases must be able to collaborate at least as easily as they currently do in a culture without aspect orientation. One way aspect orientation might improve researchers’ ease of collaboration is with standardised libraries for aspect construction. Similar libraries were developed when developing a case study for PyDySoFu’s viability [WS18a; SW16], but development was left incomplete as discussed in section 3.3.1.

The original library with this aim, Fuzzi-Moss, was originally designed to provide standardised aspects to represent behavioural variance in socio-technical systems. The aspects developed in Fuzzi-Moss have no notion of being fitted to real-world data. Instead, they are simple models of behavioural variances such as distraction, which are parameterised to allow users of the library to use these simple models in whatever manner is appropriate for their use-case. A broader collection of these behavioural variations could simplify the use of aspect-oriented behavioural variation in the research community writ large, by removing researchers’ burden to develop these themselves. Early construction of a library with Fuzzi-Moss’ goals would also support researchers in sharing models or extending others’, as they would be familiar with a common set of tools. The development of such a library and the production of case studies demonstrating its effectiveness would answer the research question:

“ Can researchers using aspect-oriented behavioural variance share a common set of tools to support and simplify the use of the technique in their codebases?”

Other tooling to support researchers in the development of aspectually augmented simulations and models could also be developed. For example, a library of fuzzers which make changes to an abstract syntax tree could be constructed. Such a library would not model specific behaviours, but would allow researchers to build models performing within-style aspect weaving without writing code which contained no metaprogramming logic. Instead, this logic would be encapsulated in utility functions provided by the proposed library. This would also reduce the work required of researchers looking to use the techniques demonstrated in this thesis. The library could also support the development of a Fuzzi-Moss-like set of socio-technical behavioural variances. However, such a library does not currently exist. A summary the contribution to the research community which the proposed library would make

is:

“ The development of a library of metaprogramming operations which simplify the construction of within-style aspects, supporting its research use and demonstrated in case studies.”

The proposed library need not be a separate codebase to the aspect orientation framework it is designed alongside. PyDySoFu could contain this collection of metaprogramming operations; however, there is no obvious requirement for the library to be contained within PyDySoFu, and separating the proposed library from the framework it is built to interact with removes any dependency on a particular aspect orientation framework; future researchers could develop alternative frameworks with within-style aspects which the proposed library could also be compatible with. We therefore suggest that the development of this library is a separate project from the development of PyDySoFu itself.

“obvious” is a little too informal, not sure what to replace it with here.

This might be a little rambly — consider omitting, rewording, or factoring into a footnote.

8.7 Testing Frameworks to detect Unrealistic Behavioural Variances

Many of the uses of aspect orientation in simulation & modelling research discussed in this chapter affect the construction and maintenance of a codebase, aspect orientation also has potential in the instrumentation of an experiment, as discussed in ??.³ Experimental instrumentation was demonstrated for data collection in this thesis’ experiments, as discussed in chapter 6 and chapter 7. **We can get more precise than chapter references; find the sections / subsections.** Instrumentation also has potential in other aspects of research software engineering; one possibility to be investigated is the use of aspects to alert researchers to impossible states visited by a model during runtime.

Many models construct a version of a real-world system which cannot achieve certain states. For example, a model of collisions might never be expected to see an increase in the total energy present across all modelled entities; simulations of socio-technical systems might expect bounds on the hours worked by a simulated workforce, or the amount of output the workforce creates; researchers might study a system to observe a n emergent property which is bounded by laws around its growth (for example, a linear increase in some property over time as opposed to exponential growth). It is important

³In particular, see relevant notes on “The use of aspect-oriented programming in scientific simulations” [GK99].

that any model of a real-world system accurately reflects the system’s physical limits and underlying mechanics. Models may fail to reflect these mechanics for many reasons, including :

Are there any other reasons as to why models might be incorrect (and instrumentation could identify this by boundary checking)?

- The model may be developed with bugs which are not detected, i.e. it is conceptually incorrect;
- The model may be verified as correct but later maintenance introduces undetected errors, i.e. it contains bugs;
- The model may be correctly written, but given unexpected inputs.

Software engineering practices such as unit testing are common practice and may catch these bugs through identifying incorrect behaviour in components. Thorough test suites may also include integration tests which observe the behaviour of many model components working together. However, there are limitations of both testing techniques:

- Test suites as described may fail to identify non-deterministic errors in a model, which arise only in a small number of cases;
- Traditional test suites also may fail to identify emergent state errors, where individual components operate correctly but their repeated integration results in an incorrect emergent property of the system ⁴;
- Test suites usually fix values such as random seeds or input data, which may hide non-deterministic state errors or may not exhibit the properties of real-world data (and so exhibit different behaviour under test than when an experiment is conducted).

We propose the use of aspects to instrument models to assert a system state within expected bounds, both as components of a test suite and as assurance that a model was correct at any time it is used as part of an experiment. All of these limitations can be alleviated by a test suite which is able to assert the correctness of model state at different points in a program’s execution. Models which non-deterministically reach erroneous states are primarily an issue during the execution of a model; aspectual state assertions need not alter a model’s program flow, so are appropriate for use during its experimental use should any unexpected state arise. These states may also be deterministically caused

⁴Emergent properties are often the state under study, as in PyDySoFu’s original case study [WS18a]

by experimental input data which is not reflected by the inputs in a test suite; weaving these aspects during experimental use also addresses this limitation. Emergent properties of a system may also be tracked over time, and aspects may measure emergent properties to ensure they are within appropriate bounds at all times in an experiment and have correct properties of their own (for example, these aspects could ensure correct growth curves for the property in question).

Aspects can also be used as a regression test to detect the recurrence of bugs after they have been identified and fixed. If a simulated system enters an incorrect state during development, and the underlying bug is fixed, an aspect observing the system's erroneous property can be constructed which alerts developers if the bug is re-introduced.

Aspects instrumenting a model throughout its use in an experiment would also serve to guarantee peer reviewers of a study that the model was correct regarding the properties observed by the aspects. A suite of aspects which observe bounds on and characteristics of system properties is, functionally speaking, a declarative set of known-good behaviours of a model. Peer reviewers who sought to check that a model was correct should be able to extend this list of properties and observe any assertions during reproduction of a study's findings.

Within-style aspects are particularly appropriate for this task, as the measurement desired of a program might exist inside an existing function or block of code. It is conceivable that observing system state before and after a function executes would not be sufficiently granular to assert correct system state in some cases. Aspectual instrumentation of research codebases is therefore particularly useful in light of new weaving techniques in PyDySoFu, traditional aspect orientation frameworks would be sufficient in instrumenting at least some models for observation.

To demonstrate the efficacy of this approach, a library of tools to observe boundaries on system states could be developed, similarly to the library of tools suggested in section 8.6. Case studies augmenting research codebases using either this observational tooling or ordinary aspects can then be constructed. These should include new codebases, to demonstrate the technique's utility during development, but should also include the instrumentation of old models and the detection of any irregularities in existing, published work. Such a project would answer the research question:

“ Can aspect orientation be used to instrument a model and assert correctness by observing that the model’s properties remain within expected bounds?”

8.8 Discovery of System Features by Fitting Aspect Parameters

8.9 Optimisation of Multiple Aspects

Experiments in chapter 6 and chapter 7 made use of models which were fitted to real-world data. However, they only fitted parameters of learning models, and only some parameters were fitted. This was partly due to the computational cost involved in fitting these parameters and time constraints on this research.

Is this point worthwhile? I cited it in IIRC an experimental chapter, but I’m unsure it’s worth including. For future Tom to revisit.

Optimisation is an independent research field with a broad range of techniques outwith the scope of this thesis to explore; however, a study of optimisation algorithms and technologies as applied to the fitting of aspects to real-world system models is a topic for future research. Models of real-world systems can be computationally expensive to complete, meaning that measuring the accuracy of a model with a variety of parameters could be computationally expensive. A survey of optimisation techniques suitable for aspectually augmented simulations and models would answer the research question:

“ Which optimisation techniques are best suited to fitting parameters of aspects woven into models of real-world systems, with regards the accuracy of the aspectually augmented model produced?”

8.10 Future Work pertaining to RPGLite

RPGLite’s dataset was analysed for the purposes it was collected for in this thesis: to aid in the realistic simulation of a well-controlled socio-technical system. However, some pieces of analysis have not been completed; this section discusses future work enabled by the dataset of real-world play which this thesis contributes.

8.10.1 Causes of Game Abandonment

The games included in the dataset produced by RPGLite finish in varying states: some are incomplete (as some games were ongoing when the snapshot of the game’s database was taken), and others are finished. A third category of games are ones where a player opted to abandon a game they were participating in. This feature of RPGLite was developed to allow players who were playing inactive opponents to free a game slot. It also allowed players who felt they had no hope of winning, or were bored with a particular game to forfeit.

As many datapoints are included in the published dataset, it is possible to observe how active players were (by measuring the frequency with which they would open RPGLite), and engagement with the app might be quantified by players’ exploration of (or frequent use of) leaderboards and customisation features. Correlations might exist between players’ activity, engagement, or experience with their propensity to abandon a game in different states. For example, one hypothesis could be that particularly active players frequently forfeited games against inactive players out of boredom, or that players who were unlikely to win a game (or were inexperienced players of RPGLite paired against competent players) might be more likely to forfeit. These results might be of interest to the broader game design community. Research investigating reasons for abandonment would answer the research question:

“ Can players who are likely to forfeit a game be identified through patterns of RPGLite play?”

8.10.2 Patterns of Play within Cliques of Players

RPGLite’s application implemented matchmaking systems which allowed players to advertise themselves as open to new games, challenge players discovered on a leaderboard, or rematch against players they had previously played against through a game history. However, two factors might have impacted how players chose their opponents:

1. Recruiting methods means people might have known each other

2. Players who find others they are well-matched against might repeatedly re-match rather than continuing to discover players they might have less satisfying games with

These factors would imply the formation of cliques of players: small groups who often created games against each other much more frequently than they did with the game’s broader community. If so, this could have implications for uses of the RPGLite play dataset. For example, in the experiments discussed in chapter 6 and chapter 7 no structure is imposed on the matching of simulated players, so that players would be uniformly exposed to the game’s community. As these experiments simulate learning from the outcomes of games, exposure to the entire community would expose players to learnings from all games played; cliques of players have a limited set of games to learn from, and so could affect what is learned by that clique. The broader community would also be insulated from what was learned within the clique, meaning that a loss of information affects all players.

Future research making use of the RPGLite dataset should be aware of any underlying patterns which could affect the dataset’s emergent properties, such as learning styles. Clique formation is an example of such a pattern; others might yet present themselves. Future research making use of the dataset of RPGLite play may look to address research questions such as:

“ Are patterns such as cliques of players present in the RPGLite dataset? If so, what patterns exist, and what is likely to have driven the formation of those patterns?”

Relatedly, the dataset presents opportunities for research in aspectually augmented simulation & modelling. A research question presents itself related to the future work described in section 8.8 related to the discovery of system features:

“ Can an aspectually augmented model be trained to identify features of RPGLite play in an explainable manner?”

8.10.3 Additional Game Features Driving Engagement

RPGLite’s dataset contains information about players’ interactions with the application itself; as the game made available some features typical of modern games (leaderboards, matchmaking,

achievements, graphical customisation), an analysis of the features most commonly used can shed light on the more effective aspects of modern game design in both the general playerbase and more dedicated players⁵⁶.

8.10.4 Larger-scale Data Collection

RPGLite’s playerbase was recruited informally and there is scope for a larger and longer-term data collection effort to be made. A re-release of the application in major mobile app stores with a concerted effort to release new seasons of the game and maintain player interest for an extended period of time — perhaps with additional features, such as in-game chat, favourites lists –of previous opponents, or match replay and analysis (with suggestions for improved play backed by the formal methods inherent in RPGLite’s design) would enable a richer analysis, and broader utility to the games research community.

While investigations into these questions warrant further study, they remain outwith the scope of this thesis, which focuses on simulation technologies more than it does game design. There are many opportunities available for the game design research community to investigate. Publications in the field from co-creators of RPGLite reflect further on the design and future improvements of the game; see [cite William’s PhD here](#) .

A large-scale user study of RPGLite could also be used to verify other research questions suggested in this section. For example, during data collection, players could be classified by their likelihood of abandoning games using the patterns identified in section 8.10.1 to verify the patterns empirically.

8.11 Discussion

This section is not intentionally left blank.

⁵For example, are some features heavily used, but only by a dedicated subset? Do all players use other features a moderate amount, showing mild but general appeal?

⁶Find related research on engagement in mobile games, which must be plentiful.

9 | Closing Discussion

A chapter rounding out the thesis with a brief discussion.

Bibliography

- [+06] hridesht rajan hridesht et al. nu: towards an aspect-oriented invocation mechanism. Tech. rep. technical report 414, iowa state university, department of computer science, 2006.
- [ABZ08] AU Aksu, Faruk Belet, and B Zdemir. “Developing aspects for a discrete event simulation system”. In: Proceedings of the 3rd Turkish Aspect-Oriented Software Development Workshop. Bilkent University. 2008, pp. 84–93.
- [AGL98] Rakesh Agrawal, Dimitrios Gunopulos, and Frank Leymann. “Mining process models from workflow logs”. In: Advances in Database Technology—EDBT’98: 6th International Conference on Extending Database Technology Valencia, Spain, March 23–27, 1998 Proceedings 6. Springer. 1998, pp. 467–483.
- [BH02a] Jason Baker and Wilson Hsieh. “Runtime aspect weaving through metaprogramming”. In: Proceedings of the 1st international conference on Aspect-oriented software development - AOSD ’02. ACM Press, 2002. DOI: [10.1145/508386.508396](https://doi.org/10.1145/508386.508396). URL: <https://doi.org/10.1145/508386.508396>.
- [BH02b] Jason Baker and Wilson C. Hsieh. “Maya: Multiple-Dispatch Syntax Extension in Java”. In: Proceedings of the ACM SIGPLAN 2002 Conference on Programming Language Design and Implementation. PLDI ’02. Berlin, Germany: Association for Computing Machinery, 2002, pp. 270–281. ISBN: 1581134630. DOI: [10.1145/512529.512562](https://doi.org/10.1145/512529.512562). URL: <https://doi.org/10.1145/512529.512562>.
- [Bir99] Colin PD Birch. “A new generalized logistic sigmoid growth equation compared with the Richards growth equation”. In: Annals of botany 83.6 (1999), pp. 713–723.

- [Bui14] JCAM Buijs. “Flexible evolutionary algorithms for mining structured process models”. In: Technische Universiteit Eindhoven 57 (2014).
- [Cap+09] Claudia Cappelli et al. “An aspect-oriented approach to business process modeling”. In: Proceedings of the 15th workshop on Early aspects - EA '09. ACM Press, 2009. DOI: [10.1145/1509825.1509828](https://doi.org/10.1145/1509825.1509828). URL: [https://doi.org/10.1145%2F1509825.1509828](https://doi.org/10.1145/2F1509825.1509828).
- [CBB13] Meriem Chibani, Brahim Belattar, and Abdelhabib Bourouis. “Toward an aspect-oriented simulation”. In: International Journal of New Computer Architectures and their Applications (IJNCAA) 3.1 (2013), pp. 1–10.
- [CBB14] Meriem Chibani, Brahim Belattar, and Abdelhabib Bourouis. “Practical benefits of aspect-oriented programming paradigm in discrete event simulation”. In: Modelling and Simulation in Engineering 2014 (2014).
- [CBB19] Meriem Chibani, Brahim Belattar, and Abdelhabib Bourouis. “Using aop in discrete event simulation: A case study with japrosim”. In: International Journal of Applied Mathematics, Computational Science and Systems Engineering 1 (2019).
- [CCK03] Andy Clement, Adrian Colyer, and Mik Kersten. “Aspect-oriented programming with AJDT”. In: ECOOP Workshop on Analysis of Aspect-Oriented Software. Vol. 10. 2003.
- [Cla73] R Lawrence Clark. “LINGUISTIC CONTRIBUTION TO GOTO-LESS PROGRAMMING”. In: Datamation 19.12 (1973), pp. 62–63.
- [CM06] Anis Charfi and Mira Mezini. “Aspect-Oriented Workflow Languages”. In: OTM Conferences. 2006.
- [CM07] Anis Charfi and Mira Mezini. “Ao4bpel: An aspect-oriented extension to bpel”. In: World wide web 10.3 (2007), pp. 309–344.
- [CMM10] Anis Charfi, Heiko Müller, and Mira Mezini. “Aspect-Oriented Business Process Modeling with AO4BPMN”. In: Modelling Foundations and Applications. Ed. by Thomas Kühne et al. Berlin, Heidelberg: Springer Berlin Heidelberg, 2010, pp. 48–61. ISBN: 978-3-642-13595-8.
- [con16] RafeKettler & github contributors. magicmethods (A guide to Python’s magic methods). <https://github.com/RafeKettler/magicmethods/tree/65cc4a7bf4e72ba18df1ad17a377d879c9e7fe41>. 2011 — 2016.

- [con21] Clarete & github contributors. ForbiddenFruit. <https://web.archive.org/web/20210515092416/https://github.com/clarete/forbiddenfruit>. 2021.
- [CS04] Ruzanna Chitchyan and Ian Sommerville. “Comparing dynamic AO systems”. In: 2004.
- [CSS04] Constantinos Constantinides, Therapon Skotiniotis, and Maximilian Stoerzer. “AOP Considered Harmful”. In: In Proceedings of European Interactive Workshop on Aspects in Software (EIWAS). 2004.
- [DBM13] Jan De Houwer, Dermot Barnes-Holmes, and Agnes Moors. “What is learning? On the nature and merits of a functional definition of learning”. In: Psychonomic bulletin & review 20 (2013), pp. 631–642.
- [DD82] Edsger W Dijkstra and Edsger W Dijkstra. “On the role of scientific thought”. In: Selected writings on computing: a personal perspective (1982), pp. 60–66.
- [dev21] SimPy developers. SimPy 4.0.1 documentation. <https://simpy.readthedocs.io/en/4.0.1/>. 2021.
- [DGM08] Adrian Dozsa, Tudor Girba, and Radu Marinescu. “How lisp systems look different”. In: 2008 12th European Conference on Software Maintenance and Reengineering. IEEE. 2008, pp. 223–232.
- [Dij68] Edsger W Dijkstra. “Letters to the editor: go to statement considered harmful”. In: Communications of the ACM 11.3 (1968), pp. 147–148.
- [Don20] Boudewijn van Dongen. “BPI Challenges: 10 years of real-life datasets”. In: IEEE Task Force on Process Mining Newsletter #2 (May 14, 2020). URL: <https://www.tf-pm.org/newsletter/newsletter-stream-2-05-2020/bpi-challenges-10-years-of-real-life-datasets> (visited on 05/14/2020).
- [DOR95] DOV DORI. “Object-process Analysis: Maintaining the Balance Between System Structure and Behaviour”. In: Journal of Logic and Computation 5.2 (Apr. 1995), pp. 227–249. ISSN: 0955-792X. DOI: 10.1093/logcom/5.2.227. eprint: <https://academic.oup.com/logcom/article-pdf/5/2/227/6244720/5-2-227.pdf>. URL: <https://doi.org/10.1093/logcom/5.2.227>.
- [DR10] Robert Dyer and Hriday Rajan. “Supporting dynamic aspect-oriented features”. In: ACM Transactions on Software Engineering and Methodology 20.2 (Aug. 2010), pp. 1–34. DOI: 10.1145/1824760.1824764. URL: <https://doi.org/10.1145/1824760.1824764>.

- [Esp04] Miklós Espák. “Aspect-Oriented Programming On Lisp”. In: International Conference on Applied Informatics. 2004.
- [FBS17] Daniel Friesel, Markus Buschhoff, and Olaf Spinczyk. “Annotations in Operating Systems with Custom AspectC++ Attributes”. In: Proceedings of the 9th Workshop on Programming Languages and Operating Systems. 2017, pp. 36–42.
- [FFN00] Robert E Filman, Daniel P Friedman, and Peter Norvig. “Aspect-oriented programming is quantification and obliviousness”. In: (2000).
- [FT+84] James France, John HM Thornley, et al. Mathematical models in agriculture. Butterworths, 1984.
- [GK99] László Gulyás and Tamás Kozsik. “The use of aspect-oriented programming in scientific simulations”. In: Proceedings of Sixth Fenno-Ugric Symposium on Software Technology, Estonia. 1999.
- [Gom15] B. Gompertz. On the Nature of the Function Expressive of the Law of Human Mortality, and on a New Mode of Determining the Value of Life Contingencies. Jan. 1815. DOI: [10.1098/rspl.1815.0271](https://doi.org/10.1098/rspl.1815.0271). URL: <https://doi.org/10.1098/rspl.1815.0271>.
- [GS04] Wasif Gilani and Olaf Spinczyk. “A family of aspect dynamic weavers”. In: Dynamic Aspects Workshop (DAW04). 2004.
- [Hil+00] Erik Hilsdale et al. “AspectJ: The Language and Support Tools”. In: Addendum to the 2000 Proceedings of the Conference on Object-Oriented Programming, Systems, Languages, and Applications (Addendum). OOPSLA ’00. Minneapolis, Minnesota, USA: Association for Computing Machinery, 2000, p. 163. ISBN: 1581133073. DOI: [10.1145/367845.368070](https://doi.org/10.1145/367845.368070). URL: <https://doi.org/10.1145/367845.368070>.
- [Ion+09] Tudor B Ionescu et al. “An Aspect-Oriented Approach for Disaster Prevention Simulation Workflows on Supercomputers, Clusters, and Grids”. In: 2009 13th IEEE/ACM International Symposium on Distributed Simulation and Real Time Applications. IEEE. 2009, pp. 21–30.
- [JP14] Paul Johannesson and Erik Perjons. An introduction to design science. Vol. 10. Springer, 2014.

- [JS07] Chris W Johnson and Christine Shea. “A comparison of the role of degraded modes of operation in the causes of accidents in rail and air traffic management”. In: 2007 2nd Institution of Engineering and Technology International Conference on System Safety. IET. 2007, pp. 89–94.
- [JWO12] Amin Jalali, Petia Wohed, and Chun Ouyang. “Aspect oriented business process modelling with precedence”. In: International Workshop on Business Process Modeling Notation. Springer. 2012, pp. 23–37.
- [Kav+19] William Kavanagh et al. “Balancing turn-based games with chained strategy generation”. In: IEEE Transactions on Games 13.2 (2019), pp. 113–122.
- [Kav21] William Kavanagh. “Using probabilistic model checking to balance games”. PhD thesis. University of Glasgow, 2021.
- [KDB91] Gregor Kiczales, Jim Des Rivieres, and Daniel G Bobrow. The art of the metaobject protocol. MIT press, 1991.
- [Kel08] Stephen Kell. “A survey of practical software adaptation techniques.” In: J. Univers. Comput. Sci. 14.13 (2008), pp. 2110–2157.
- [KH98] Ralph Keller and Urs Hölzle. “Binary component adaptation”. In: European Conference on Object-Oriented Programming. Springer. 1998, pp. 307–329.
- [Kic+01] G. Kiczales et al. “An Overview of AspectJ”. In: ECOOP. 2001.
- [Kic+97] Gregor Kiczales et al. “Aspect-oriented programming”. In: European conference on object-oriented programming. Springer. 1997, pp. 220–242.
- [KM20] William Kavanagh and Alice Miller. “Gameplay Analysis of Multiplayer Games with Verified Action-Costs”. In: The Computer Games Journal 10.1-4 (Dec. 2020), pp. 89–110. DOI: [10.1007/s40869-020-00121-5](https://doi.org/10.1007/s40869-020-00121-5). URL: <https://doi.org/10.1007%2Fs40869-020-00121-5>.
- [KM21] William Kavanagh and Alice Miller. “Gameplay analysis of multiplayer games with verified action-costs”. In: The Computer Games Journal 10 (2021), pp. 89–110.
- [Kuh12] Thomas S Kuhn. The structure of scientific revolutions. University of Chicago press, 2012.
- [KWM20] William Kavanagh, Tom Wallis, and Alice Miller. “RPGLite player data and lookup tables”. In: (2020).

- [Lac97] Sheldon J Lachman. “Learning is a process: Toward an improved definition of learning”. In: *The Journal of psychology* 131.5 (1997), pp. 477–480.
- [LC07] Gary T Leavens and Curtis Clifton. “Multiple concerns in aspect-oriented language design: a language engineering approach to balancing benefits, with examples”. In: *Proceedings of the 5th workshop on Software engineering properties of languages and aspect technologies*. 2007, 6–es.
- [Li10] Chen Li. “Mining process model variants: Challenges, techniques, examples”. PhD thesis. University of Twente, The Netherlands, 2010.
- [Mac+11] Idarlan Machado et al. “Managing variability in business processes”. In: *Proceedings of the 2011 international workshop on Early aspects - EA ’11*. ACM Press, 2011. DOI: [10.1145/1960502.1960508](https://doi.org/10.1145/1960502.1960508). URL: <https://doi.org/10.1145/1960502.1960508>.
- [MAG12] Giselle Machado, Vander Alves, and Rohit Gheyi. “Formal Specification and Verification of Well-formedness in Business Process Product Lines”. In: *6th Latin American Workshop on Aspect-Oriented Software Development: Advanced Modularization Techniques, LAWASP*. Vol. 12. 2012.
- [Mar94] Robert Martin. “OO design quality metrics”. In: *An analysis of dependencies* 12.1 (1994), pp. 151–170.
- [Mat08] Norm Matloff. “Introduction to discrete-event simulation and the simpy language”. In: Davis, CA. Dept of Computer Science. University of California at Davis. Retrieved on August 2.2009 (2008), pp. 1–33.
- [MD11] Richard Millham and Evans Dogbe. “Aspect-Oriented Security and Exception Handling within an Object Oriented System”. In: *2011 IEEE 35th Annual Computer Software and Applications Conference Workshops*. July 2011, pp. 321–326. DOI: [10.1109/COMPSACW.2011.60](https://doi.org/10.1109/COMPSACW.2011.60).
- [Mit+17] Alexey A. Mitsyuk et al. “Generating event logs for high-level process models”. In: *Simulation Modelling Practice and Theory* 74 (May 2017), pp. 1–16. DOI: [10.1016/j.simpat.2017.01.003](https://doi.org/10.1016/j.simpat.2017.01.003). URL: <https://doi.org/10.1016/j.simpat.2017.01.003>.
- [MK03] Hidehiko Masuhara and Gregor Kiczales. “A modeling framework for aspect-oriented mechanisms”. In: *Proceeding ECOOP’03*. 2003.

- [MR02] Katharina Mehner and Awais Rashid. “Towards a standard interface for runtime inspection in AOP environments”. In: Workshop on Tools for Aspect-Oriented Software Development at OOPSLA. Vol. 2. 2002.
- [Nie77] Jurg Nievergelt. “Information Content of Chess Positions”. In: SIGART Bull. 62 (Apr. 1977), pp. 13–15. ISSN: 0163-5719. DOI: [10.1145/1045398.1045400](https://doi.org/10.1145/1045398.1045400). URL: <https://doi.org/10.1145/1045398.1045400>.
- [Obj11] Object Management Group (OMG). Business Process Model and Notation, Version 2.0. OMG Document Number formal/2011-01-03 (<https://www.omg.org/spec/BPMN/2.0>). 2011.
- [Obj15] Object Management Group (OMG). Unified Modelling Language, Version 2.5. OMG Document Number formal/2015-03-01 (<http://www.omg.org/spec/UML/2.5>). 2015.
- [OLe20] Aran O’Leary. “Simulating Human Behaviour in a Micro-epidemic using an Agent-based Model with Fuzzing Aspects”. In: Unpublished masters thesis (2020).
- [PAG03] Andrei Popovici, Gustavo Alonso, and Thomas Gross. “Just-in-Time Aspects: Efficient Dynamic Weaving for Java”. In: Proceedings of the 2nd International Conference on Aspect-Oriented Software Development. Association for Computing Machinery, 2003, pp. 100–109.
- [Par72] D. L. Parnas. “On the criteria to be used in decomposing systems into modules”. In: Communications of the ACM 15.12 (Dec. 1972), pp. 1053–1058. DOI: [10.1145/361598.361623](https://doi.org/10.1145/361598.361623). URL: <https://doi.org/10.1145/361598.361623>.
- [PGA02] Andrei Popovici, Thomas Gross, and Gustavo Alonso. “Dynamic Weaving for Aspect-Oriented Programming”. In: Proceedings of the 1st International Conference on Aspect-Oriented Software Development. AOSD ’02. Enschede, The Netherlands: Association for Computing Machinery, 2002, pp. 141–147. ISBN: 158113469X. DOI: [10.1145/508386.508404](https://doi.org/10.1145/508386.508404). URL: <https://doi.org/10.1145/508386.508404>.
- [Pop13] Karl Popper. “The Logic and Evolution of Scientific Theory”. In: Routledge, 2013, pp. 3–22.
- [Pou+15] Asef Pourmasoumi et al. “On Business Process Variants Generation.” In: CAiSE Forum. 2015, pp. 179–188.

- [Prz10] Adam Przybylek. “What is Wrong with AOP?” In: ICSoft (2). Citeseer. 2010, pp. 125–130.
- [RIC59] F. J. RICHARDS. “A Flexible Growth Function for Empirical Use”. In: Journal of Experimental Botany 10.2 (June 1959), pp. 290–301. ISSN: 0022-0957. DOI: [10.1093/jxb/10.2.290](https://doi.org/10.1093/jxb/10.2.290). eprint: <https://academic.oup.com/jxb/article-pdf/10/2/290/1411755/10-2-290.pdf>. URL: <https://doi.org/10.1093/jxb/10.2.290>.
- [Ros97] Jarrett Rosenberg. “Some misconceptions about lines of code”. In: Proceedings fourth international software metrics symposium. IEEE. 1997, pp. 137–142.
- [SA13] Thomas Stocker and Rafael Accorsi. “Secsy: Security-aware synthesis of process event logs”. In: Proceedings of the 5th International Workshop on Enterprise Modelling and Information Systems Architectures, St. Gallen, Switzerland. Citeseer. 2013.
- [SA14] Thomas Stocker and Rafael Accorsi. “SecSy: A Security-oriented Tool for Synthesizing Process Event Logs.” In: BPM (Demos). 2014, p. 71.
- [Seo11] Simon Seow. “OBASHI White Paper”. In: APMG-International, High Wycombe, UK (2011).
- [Sil+20] Hercules Sant Ana da Silva Jose et al. “Implementation of Aspect-oriented Business Process Models with Web Services.” In: Bus. Inf. Syst. Eng. 62.6 (2020), pp. 561–584.
- [SL07] Olaf Spinczyk and Daniel Lohmann. “The design and implementation of AspectC++”. In: Knowledge-Based Systems 20.7 (2007), pp. 636–651.
- [SM14] Ivan Shugurov and Alexey A Mitsyuk. “Generation of a set of event logs with noise”. In: Proceedings of the Spring/Summer Young Researchers’ Colloquium on Software Engineering. 8. ... 2014.
- [Ste06] Friedrich Steimann. “The paradoxical success of aspect-oriented programming”. In: Proceedings of the 21st annual ACM SIGPLAN conference on Object-oriented programming systems, languages, and applications - OOPSLA ’06. ACM Press, 2006. DOI: [10.1145/1167473.1167514](https://doi.org/10.1145/1167473.1167514). URL: <https://doi.org/10.1145%2F1167473.1167514>.
- [Sto23] Tim Storer. Theatre_AG. https://web.archive.org/web/20230728153002/https://github.com/twsswt/theatre_ag. 2023.

- [Sul01] Gregory T Sullivan. “Aspect-oriented programming using reflection and metaobject protocols”. In: Communications of the ACM 44.10 (2001), pp. 95–97.
- [SW16] Tim Storer and William Wallis. Fuzzi-Moss. <https://web.archive.org/web/20201018121800/https://github.com/twsswt/fuzzi-moss>. 2016.
- [Tha+17] Dai Hai Ton That et al. “Sciunits: Reusable Research Objects”. In: 2017 IEEE 13th International Conference on e-Science (e-Science). IEEE, Oct. 2017. DOI: [10.1109/escience.2017.51](https://doi.org/10.1109/escience.2017.51). URL: <https://doi.org/10.1109%2Fescience.2017.51>.
- [Van+05] Boudewijn F Van Dongen et al. “The ProM framework: A new era in process mining tool support”. In: Applications and Theory of Petri Nets 2005: 26th International Conference, ICATPN 2005, Miami, USA, June 20-25, 2005. Proceedings 26. Springer. 2005, pp. 444–454.
- [VD09] Guido Van Rossum and Fred L. Drake. Python 3 Reference Manual. Scotts Valley, CA: CreateSpace, 2009. ISBN: 1441412697.
- [Ver45] PF Verhulst. “La loi d’accroissement de la population”. In: Nouveaux Memories de l’Académie Royale des Sciences et Belles-Lettres de Bruxelles 18 (1845), pp. 14–54.
- [Von38] Ludwig Von Bertalanffy. “A quantitative theory of organic growth (inquiries on growth laws. II)”. In: Human biology 10.2 (1938), pp. 181–213.
- [VVK09] Jussi Vanhatalo, Hagen Völzer, and Jana Koehler. “The refined process structure tree”. In: Data & Knowledge Engineering 68.9 (2009), pp. 793–818.
- [VW04] Wil MP Van der Aalst and Anton JMM Weijters. “Process mining: a research agenda”. In: Computers in industry 53.3 (2004), pp. 231–244.
- [Wal+20] William Wallis et al. “Designing a Mobile Game to Generate Player Data-Lessons Learned”. In: (2020).
- [Wal+21] William Wallis et al. “Designing a mobile game to generate player data - lessons learned THIS IS A DUPLICATE FIXME THIS IS A DUPLICATE FIXMETHIS IS A DUPLICATE FIXMETHIS IS A DUPLICATE FIXME THIS IS A DUPLICATE FIXME THIS IS A DUPLICATE FIXME THIS IS A DUPLICATE FIXME THIS IS A DUPLICATE FIXME Duplicate citation to clean up ”. In: CoRR abs/2101.07144 (2021). arXiv: [2101.07144](https://arxiv.org/abs/2101.07144). URL: <https://arxiv.org/abs/2101.07144>.
- [WC10] Paul Wallis and Fergus Cloughley. The OBASHI Methodology. The Stationery Office, 2010.

- [WJ97] AR Werker and KW Jaggard. “Modelling asymmetrical growth curves that rise and then fall: Applications to foliage dynamics of sugar beet (*Beta vulgaris*L.)” In: *Annals of Botany* 79.6 (1997), pp. 657–665.
- [WS18a] Tom Wallis and Tim Storer. “Modelling realistic user behaviour in information systems simulations as fuzzing aspects”. In: *International Conference on Advanced Information Systems Engineering*. Springer. 2018, pp. 254–268.
- [WS18b] Tom Wallis and Tim Storer. “Process Fuzzing as an Approach to Genetic Programming”. In: *Proceedings of the SICSA Workshop on Reasoning, Learning and Explainability*. Ed. by Kyle Martin, Nirmalie Wiratunga, and Leslie S. Smith. Vol. 2151. *CEUR Workshop Proceedings*. Aberdeen, UK: CEUR-WS.org, June 2018. URL: http://ceur-ws.org/Vol-2151/Paper_S3.pdf.
- [WS18c] Tom Wallis and Tim Storer. *PyDySoFu*. <https://github.com/twsswt/pydysofu>. 2018.
- [WS23] Tom Wallis and Tim Storer. *ASP*. <https://web.archive.org/web/20230730164214/https://github.com/probablytom/asp>. 2023.
- [YC79] Edward Yourdon and Larry L Constantine. “Structured design. Fundamentals of a discipline of computer program and systems design”. In: *Englewood Cliffs: Yourdon Press* (1979).