

Plausibly Realistic Sociotechnical Simulation with Aspect Orientation

Tom Wallis

Submitted in fulfilment of the requirements for the Degree of Doctor of
Philosophy
School of Computing Science
College of Science and Engineering

Table of Contents

1	Introduction	7
1.1	The first section header	7
2	Relevant Literature	9
2.1	Early work on PyDySoFu	9
2.1.1	PyDySoFu's implementation and features	10
2.1.2	Aspect Orientation & PyDySoFu	11
2.1.3	Opportunities presented by PyDySoFu	12
2.2	Dynamism in AOP	13
2.2.1	Dynamic and static weaving	14
2.2.2	PROSE	15
2.2.3	Handi-Wrap	16
2.2.4	Nu	16
2.2.5	Binary Component Adaptation	17
2.3	Aspect Orientation & Modelling	17
2.3.1	Aspect Orientation in Business Process Modelling	18
2.3.2	MAML & SWARM	18
2.3.3	18
2.4	Aspect Orientation & Simulation	18
2.4.1	Aspect-oriented L-Systems	19
2.5	Variable Behaviour in Simulation and Modelling	20
2.5.1	Business Process Modelling & variation in behaviour	20
2.5.2	Variations in Process Models	21
2.5.3	Process mining & variation in data	21
2.6	Research Opportunities in the Literature	24

Todo list

I assume we discuss PDSF existing in the intro	9
Make a crossreference to the discussion of PDSF's rewrite.	11
CITECITECITE	12
Do we need a brief explainer of what aspect orientation is before jumping into outside lit? Or will this go in the introduction? already a little in the earlier litrev subsections.	13
CITECITECITE	14
surely this isn't dynamic, Ian...?!	14
Find more citations for both dynamic weavers with aspect monitors and without. Nanning aspects? Nu?	15
Do I want a citation for this? Probably not, but worth revisiting.	16
Revisit this inline quote format	16
Some extra things here about Nu, such as [13].	16
Write more NU writeup — requires more citations etc.	16
which?!	17
find more citations for AOBPM	18
Is MAML/SWARM really modelling, or simulation? Simulation, right?	18
The simulation section <i>badly</i> needs revisiting.	18
surely one more here to round this out.	18
does hawthorne effect need a citation?	19
throw more in here from excel	19
throw more in here from excel	19
Rethink the argument around AOP and sim in the lit. The angle is, "much of the lit strays from any kind of real-world testing of sim tech — why?"	19

include more!	19
Is this going from data to a model? Models to data? Potentially multiple models either way depending on their use? Consider this and possibly restructure.	20
Write a short subsection on the trend of sanitising data / coping with noisy datasets. Sometimes behavioural variance isn't desirable.	21
Consider making noisy data a subsubsection of ??	21
So, there's some stuff to cite on sanitisation and mining in the presence of noise — see On process model synthesis based on event logs with noise, [28].	21
Add more stuff to be cited here, at the very least...	21
Shugurov paper from 2014 should be checked for citing here too. There's a useful summary in On Business Process Variants Generation (Pourmasoumi 2015).	22
In [18], we can find good things I should write about.	23
In [38], we can find good things I should write about.	23
In [39], we can find good things I should write about.	23
cite work that focuses on variations specifically	24
what are our research opportunities given the above...?	24

Chapter 1

Introduction

Some good things here
also some text

1.1 The first section header

this thesis has many features, such as content and citations, such as [1]

Chapter 2

Relevant Literature

The work presented in this thesis revolves around the combination of simulation & modelling and aspect orientation. As discussed, some research on the topic was also done prior to this body of work **I assume we discuss PDSF existing in the intro**. There are therefore several things to discuss: work that was done on this project prior to this thesis must be discussed so as to make clear what is *new* in this thesis, related work by others should make clear the context in which this research was done, and opportunities found in this related work for which PyDySoFu is well-suited must be identified to set the scene for the rest of this thesis.

With this in mind, section 2.1 discusses the earlier work contextualising this thesis, summarised in [2]. As PyDySoFu has some unusual features as an aspect orientation library, section 2.2 follows this with a review of similar approaches in the literature. section 2.3 and section 2.4 then review aspect orientation as applied in simulation and modelling respectively, and section 2.5 reviews literature on variable behaviour in simulation and modelling, a key strength of PyDySoFu. Finally, research opportunities identified in these various bodies of work are discussed in section 2.6, which pave the way for the research composing the body of this thesis.

2.1 Early work on PyDySoFu

PyDySoFu¹ is a Python library [3] built for making changes to the source code of a Python function as it is called, and before it is executed, while the original function definition remains oblivious to the changes being made. It was originally developed as an honours-level dissertation, which was built upon

¹Or “PDSF” for short.

and detailed in a subsequent paper [2]. This thesis furthers that original work. To be clear about the work this thesis contains, the state of the project *before* this work began is briefly discussed here.

2.1.1 PyDySoFu's implementation and features

The original version of PyDySoFu² patched Python classes with additional handlers. Attributes of Python objects are usually retrieved using dot notation (i.e. `object_id.attr_id`), which evaluates internally to a call to `object_class.__getattr__()`. PyDySoFu replaces a class' built-in `__getattr__()` method with a new one, which calls the original to acquire the required attribute.

In the case where the required attribute is not callable, the value is returned as normal. Callable attributes are modified, however. In this case, the replacement `__getattr__()` also checks for a set of manipulations to make to the original code. These can be applied before or after the original code is run, as well as around it. A new function is returned containing a reference to the originally sought attribute, but which will search for these additional pieces of work before executing it, and can execute this work before or after the call (or both). These pieces of work are referred to as “advice”, adopting aspect orientation terminology.

As discussed further in section 2.1.2, this approach is effectively an implementation of a traditional aspect orientation framework. However, unlike existing frameworks, PyDySoFu also supports a special kind of “around” advice: before a function is called, it can be rewritten. This is done by applying “before” advice which retrieves the abstract syntax tree of the target callable attribute using Python's `inspect` module (its built-in reflection), applying arbitrary transformations to the tree, and recompiling it into a Python code object (its representation of its internal bytecode). At this point, many things are possible: the transformation can be cached for later use, can replace the original callable's code object to make the transformation persistent, or can be discarded after use. This transformed code is run in lieu of the original, effectively enabling aspect orientation which can make adaptations *inside* a procedure as well as before and after its execution.

This approach also had some limitations:

→ Traditional pointcuts cannot target points inside a procedure, meaning

²Further improvements have been made through this research which improve on the design, but this section is to discuss the state of the project before this work began, and the general principles around it, which remain unchanged.

that an aspect applied “inside” its target must manage the points where its transformation is applied manually.

- Importantly, a callable object’s internal bytecode cannot be replaced in Python3, leading to a rewrite discussed in [Make a crossreference to the discussion of PDSF’s rewrite](#).
- This method is significantly slower than other aspect orientation approaches, as rewriting a class’ `__getattribute__` method means that *every* resolution of an object’s attributes — whether they are methods or values, and including a class’ built-in “magic” methods — incurs an overhead from the replaced `__getattribute__` implementation. However slight this overhead can be made, affecting Python’s built-in methods on classes means that rewriting the `__getattribute__` method is unavoidably expensive due to the scale of these methods’ use.

However, the goal of the original research was to develop a flexible “proof-of-concept” of aspect orientation adapting procedure definition at runtime, which was successfully achieved[2], [4].

2.1.2 Aspect Orientation & PyDySoFu

The goals of “changing a function’s behaviour” and maintaining “obliviousness” in the original definition of that function speak to the goals of the aspect oriented programming paradigm[5]. Quoting their original definitions:

“Components are properties of a system, for which the implementation can be cleanly encapsulated in a generalized procedure. Aspects are properties for which the implementation cannot be cleanly encapsulated in a generalized procedure. Aspects and cross-cut components cross-cut each other in a system’s implementation. [...] The key difference between AOP and other approaches is that AOP provides component and aspect languages with different abstraction and composition mechanisms.”

Generally, aspect orientation is perceived to be a technique for separation of concerns. Any cross-cutting concerns can be separated from their components into aspects applied where that concern arises. The strength of aspect orientation lies in its compositional nature: developers can write short, maintainable implementations of a procedure’s core purpose (for example, business logic) and ancillary concerns such as logging or security can be woven into this

implementation as preprocessing, compilation, or at runtime. This compositional nature is what gives rise to aspect orientation’s “obliviousness”, as the procedure targetted by a piece of advice is written without regard to that fact.

The original PyDySoFu implementation was an aspect orientation library focusing on separating a function’s definition from *potential changes to it*. This was used to model “contingent behaviour” — behaviour sensitive to some condition — as an original, “idealised” definition of that behaviour, plus some possible alterations. These changes might apply to many different behaviours in the same manner, and therefore represent concerns which separate cleanly into an aspect. An example would be the behaviour of a worker whose job requires focus on allocated tasks. A lack of focus could be represented as steps of the worker’s tasks being executed in duplicate, out-of-order, or skipped. Assuming aspects as described by Kiczales, Lamping, Mendhekar, *et al.* are able to edit the definition or execution of a procedure³, such contingent behaviours are well modelled as aspects.

To achieve this, a model was presented in [2] wherein aspects were developed which could change function *definitions* on each invocation of that function, contingent on program state. This allowed behavioural adaptation to be simulated in an aspect-oriented fashion. In addition, a library of behavioural adaptations called FUZZI-MOSSCITECITECITE was developed which implemented many cross-cutting, contingent behaviours in procedural simulations of socio-technical systems.

One important contribution of this work is that PDSF aspects are effectively able to operate *inside* a target. In typical aspect orientation frameworks such as AspectJ[6], aspects operate by effectively prepending or appending work to a target, referred to as “before” or “after” pointcuts respectively. To do both is referred to as “around”. By manipulating procedures within Python directly, PDSF is able to manipulate its target from a new perspective, adding (or removing) work during the target’s execution⁴. Moreover, because weaving is performed dynamically, every execution of a function may perform different operations.

2.1.3 Opportunities presented by PyDySoFu

PyDySoFu presented several opportunities for future research. Some salient properties of the original work include:

- It provided an aspect orientation library which could weave and unweave aspects during program execution, without relying on anything other than

³As opposed to simply wrapping it with additional behaviour before and/or after execution

⁴Similarly to ??, but in an aspect oriented manner.

Python’s built-in language features. As discussed in section 2.2, this is supported by some early aspect orientation frameworks also, but AspectJ dominates in the world of aspect orientation frameworks and does not support weaving during program execution.

- It provided the capacity to weave aspects *inside* targets, as opposed to around them, or at either end of their execution. So far as we are aware, no aspect orientation framework in research or industry has offered this feature, and its applications and potential are yet to be explored.
- Relatedly, PyDySoFu was used in the context of simulating behaviour which may change over time. Contingent behaviour being a cross-cutting concern is an innovation of the early research which suggests aspect orientation may have strong applications in socio-technical simulation & modelling.

Do we need a brief explainer of what aspect orientation is before jumping into outside lit? Or will this go in the introduction? already a little in the earlier litrev subsections.

The amount of potential investigation which can be done into the dynamic weaving of target-altering / “inside” aspects in simulation & modelling applications is vast. While literature on the complete topic is absent, each individual component of this research angle is well-studied on its own. These opportunities might be related to existing literature through the following questions:

- How does PyDySoFu compare to existing aspect orientation frameworks, particularly those with a focus on dynamic weaving? Related frameworks are summarised and compared in section 2.2.
- What is the use of aspect orientation in simulation & modelling? How does the approach taken in PyDySoFu ’s prior work relate to existing approaches? This will be discussed for simulation in section 2.4, and for modelling in section 2.3.
- Variability is important to capture in any socio-technical model or simulation. How is variability treated in existing literature, and how does this relate to PyDySoFu ’s approach? This will be explored in section 2.5.

2.2 Dynamism in AOP

Aspect orientation frameworks have supported “dynamic behaviour” in different ways for a long time. This is largely through a technique referred to as

dynamic- or runtime-weaving.

2.2.1 Dynamic and static weaving

Dynamic weaving integrates advice into a target program during its execution, as opposed to during compilation or a pre-processing step. The advantage of this is flexibility: dynamic aspect-oriented approaches have been proposed for deploying hotfixes in safety-critical scenarios where software systems cannot be taken offline to apply patchesCITECITECITE, and in adaptive mobile scenarios where software may need to alter its properties in response to its environment[7], or when debugging code to apply potential patches without reloading an entire software system[8].

To meet these needs, software systems need to check for available aspects to weave at any join point, as it is always possible that the set of applied advice has changed since the program last encountered this point. The technique therefore presents a tradeoff compared to traditional (static) aspect weaving, as illustrated in [9]. Chitchyan and Sommerville generalise this tradeoff by describing different mechanisms used to implement aspect orientation into three main categories⁵, each with their own strengths:

“Total hook weaving” alters all join points where advice may be applied before runtime, so that during execution each join point “watches” for applied advice. The benefit of this approach is that aspects can be applied at any point at runtime, but this flexibility is bought at the cost of maximum overhead: at all points where weaving *may* be possible, checks for applied advice must be made.

“Actual hook weaving” weaves hooks only to join points that are expected to be in use. This limits overhead from watching for applied advice, at the cost of flexibility: during program execution, advice may be applied or retracted *only at specific points within the system*.

“Collected weaving” weaves aspects directly into code at compilation / pre-processing*surely this isn’t dynamic, Ian...?!*, so as to collect advice and target codebase into a single unit. This provides exactly the necessary amount of overhead, and in many cases may result in requiring no “watching” for applied advice at all, but this limits a developer’s ability to amend advice supplied at runtime.

⁵Drawing from [8], [10] where “PROSE”, a particularly influential dynamic aspect orientation library, is detailed.

There is an almost direct tradeoff between the number of potential join points actively checking for applied advice at runtime, and the overhead of dynamism in any aspect oriented framework, with “total hook weaving” providing complete adaptability at the expense of checking at all possible points whether advice is applied.

Another tradeoff could be seen to be the clarity of dynamically woven aspect oriented code. Aspect orientation is already criticised for the lack of clarity as to what woven code will *do* when run, and where weaving can change during program execution, static tools are less useful in making these predictions. Some tools have been produced which do provide tooling for achieving understanding as to what dynamically woven code will do when executed (also called an “Aspect Monitor”, as discussed in [8]), but they are often limited or missing from a dynamically weaving framework’s implementation (such as [11]). **Find more citations for both dynamic weavers with aspect monitors and without. Nanning aspects? Nu?**

2.2.2 PROSE

One implementation of dynamic weaving is PROSE[8], [10], a library which achieves dynamic weaving by use of a Just-In-Time compiler for Java. The authors saw aspect orientation as a solution to software’s increasing need for adaptivity: mobile devices, for example, could enable a required feature by applying an aspect as a kind of “hotfix”, thereby adapting over time to a user’s needs. Other uses of dynamic aspect orientation they identify are in the process of software development: as aspects are applied to a compiled, live product, the join points being used can be inspected by a developer to see whether the pointcut used is correct. If not, a developer could use dynamic weaving to remove a mis-applied aspect, rewrite the pointcut, and weave again without recompiling and relaunching their project.

Indeed, the conclusion Popovici, Alonso, and Gross provide in [10] indicates that the performance issues generalised by Chitchyan and Sommerville in [9] may prevent dynamic aspect orientation from being useful in production software, but that it presented opportunities in a prototyping or debugging context.

PROSE explores dynamic weaving as it could apply in a development context, but the authors do not appear to have investigated dynamic weaving as it could apply to simulation contexts, or others where software making use of aspects does not constitute a *product*.

2.2.3 Handi-Wrap

Handi-Wrap[11] is a Java library allowing for dynamic weaving via a third-party language designed for metaprogramming, called Maya. Do I want a citation for this? Probably not, but worth revisiting.. At the time of development Handi-Wrap’s dynamic aspect weaving feature was novel: the aspect orientation library of note, AspectJ, wove only statically⁶, and Handi-Wrap’s purpose was to show that DSLs for metaprogramming could pave a way to dynamic weaving.

Baker and Hsieh implemented an aspect orientation framework which is reasonably performant, weaves dynamically, and allows for aspect orientation features to be implemented natively for greater control as compared to Handi-Wrap’s then competitor, AspectJ. As a tool, Handi-Wrap demonstrated a promising approach to dynamic weaving, but the project appears to have enjoyed less attention than similar work (such as PROSE, described in ??).

The technique used to implement Handi-Wrap (implementation via a metaprogramming-specific DSL, Maya) is familiar, in that it shares a perspective on dynamic weaving with early PyDySoFu work. The fuzzers used in [2] applied transformations to abstract syntax trees, not unlike a LISP-style macro. To quote [12] by way of contrast: “Maya generalizes macro systems by treating grammar productions as generic functions...”[Revisit this inline quote format](#) The two approaches have clear differences. Most notably, PyDySoFu’s entire implementation *and use* is performed in Python directly, and Maya’s intended purpose is metaprogramming in a more general sense. It is possible that, while Maya provided a useful foundation to explore the dynamic weaving of aspects, its lack of adoption as a language limited handi-wrap’s reach; nevertheless, it is encouraging to see another use of metaprogramming for weaving aspects at runtime.

2.2.4 Nu

[Some extra things here about Nu, such as \[13\].](#)

Nu is an aspect orientation framework written in Java which achieves dynamic weaving by way of the Nu virtual machine [13]. This introduced new primitives in Java for the application and removal of aspects: BIND and REMOVE. [Write more NU writeup — requires more citations etc.](#)

⁶AspectJ now supports what it calls “load-time weaving” — that is, weaving aspects as classes are loaded into the JVM — but not weaving to things that are *already* loaded, meaning AspectJ still allows for only a particular flavour of dynamic behaviour.

2.2.5 Binary Component Adaptation

Binary Component Adaptation[14] (BCA) is a technique for performing adaptations on software components after compilation. Though it works on already-compiled code it does provide dynamic behaviour: the technique can adapt software components via rewriting before or during the loading of its target. Like some aspect orientation techniques^{which?!}, BCA adapts a Java class loader to make its adaptations, but unlike aspect oriented approaches it does not require access to the original source of the software. For scientific simulation purposes, it could therefore be appealing in situations where adaptations are made to another researcher's simulations — assuming the original source code is not published — or in security settings investigating trust in compilers and runtimes[15]. In the present context of developing socio-technical simulations however, this does not appear to be an advantage, particularly at a time when the source code of software components of research projects are increasingly published.

An important distinction to be made is that BCA provides an example of runtime adaptation, but does not enable an aspect oriented approach and is not developed with separation of concerns in mind. It is presented here as a useful contrast to PyDySoFu: it demonstrates an alternative technique for achieving dynamic runtime source manipulation, even if the lack of separation of concerns means it would not be well applied for this thesis.

2.3 Aspect Orientation & Modelling

Having discussed aspect orientation as it is used in a simulation context, it is natural to investigate its use in modelling research, too.

Simulation and modelling are similar topics and are often combined into a single study. However, their goals differ. Simulation typically involves the study of processes or behaviour: there is an expectation that simulations are *executed* or *run*. This often produces data. The intent of modelling is more structural in nature: models are typically observed or analysed to gain insights. Quoting Maria's introduction [16]:

“Modeling [*sic*] is the process of producing a model; a model is a representation of the construction and working of some system of interest. A model is similar to but simpler than the system it represents.

[...]

A simulation of a system is the operation of a model of the

system. The model can be reconfigured and experimented with; usually, this is impossible, too expensive or impractical to do in the system it represents. The operation of the model can be studied, and hence, properties concerning the behavior of the actual system or its subsystem can be inferred.”

Maria’s definition implies that to simulate is to operate a model. Whether this model is constructed for the purpose of simulation or for study in its own right, a simplified representation of the system being studied is implicitly required for any simulation. However, modelling does not imply simulation. Models can be studied for their own merits, and many modelling frameworks exist which are made explicitly for their own study, without regard to their use in simulation⁷. Aspect orientation has seen some study in modelling, particularly for socio-technical modelling, and while aspect-oriented socio-technical modelling is not generally researched with subsequent simulation in mind, an important body of work is still present, and therefore important to discuss.

2.3.1 Aspect Orientation in Business Process Modelling

Aspect orientation for socio-technical systems is particularly well studied in the business process modelling community[19], [20] [find more citations for AOBPM](#)

2.3.2 MAML & SWARM

[Is MAML/SWARM really modelling, or simulation? Simulation, right?](#)

2.3.3

2.4 Aspect Orientation & Simulation

[The simulation section *badly* needs revisiting.](#)

Surprisingly, little literature exists pertaining specifically to the use of aspect-orientation in a simulation context. Aspect orientation is often applied to modelling as discussed in section 2.3, used to compose a perspective of the world from individual parts, but in a way which isn’t necessarily executable or able to produce data.

⁷Consider UML, a well-studied modelling framework which is generally not used for any kind of simulation — depending on its use, it often cannot be — and for which many alternatives now exist specifically to address this limitation [17], [18].[surely one more here to round this out.](#)

Early in the history of aspect orientation as an emerging paradigm, there was some interest in its use for scientific simulation. [21] discuss that computer simulations require code for both observation of a simulation and the simulation itself, and that misuse of this could cause what is in effect a kind of Hawthorne Effect **does hawthorne effect need a citation?**, where the inclusion of observation code intertwined with simulation code might influence the outcome of an experiment. They suggest that improving simulation technologies could combat this approach. Aspect Orientation, being developed specifically with obliviousness in mind, is an ideal candidate which Gulyás and Kozsik identify.

Much of the literature concerning aspect-oriented programming and simulation focus on tooling support for aspect-oriented simulation, rather than investigations into its efficacy. For example, attempts have been made to integrate aspect orientation into new tools [22] **throw more in here from excel**, or into existing ones [23] **throw more in here from excel. Rethink the argument around AOP and sim in the lit. The angle is, “much of the lit strays from any kind of real-world testing of sim tech — why?”**

Some experiments specifically using aspect orientation in the implementation of process-based simulations also exist [24] **include more!**. For example, Ionescu, Piater, Scheuermann, *et al.* apply aspect orientation in a nuclear disaster prevention simulation. Their motivation is that code can become complex to maintain over time and changes to the scientific zeitgeist or to regulatory requirements become costly as technical debt mounts. Aspect orientation therefore allows developers to separate functionality into distinct modules more easily, without disturbing the underlying codebase.

2.4.1 Aspect-oriented L-Systems

Aspect-orientation is also applied in other simulation paradigms. Cieslak, Seleznyova, Prusinkiewicz, *et al.* investigated the use of aspect orientation in L-system based simulations [25]. An L-system [26] is defined by a set of symbols, an initial string composed of these symbols, and a set of rules for rewriting substrings. While being a powerful tool for representing fractal structures, they were originally conceived of for plant modelling (and still see the most use in this field).

Cieslak, Seleznyova, Prusinkiewicz, *et al.* note that some details of plant modelling are actually cross-cutting concerns against many plants or families of plants. To represent these, they introduce a new language to describe plant models which makes use of aspect orientation to represent these cross-cutting

concerns. They test the approach by representing carbon dynamics, apical dominance and biomechanics as cross-cutting concerns that are integrated into a previously published model of kiwifruit shoot development. Cieslak, Seleznyova, Prusinkiewicz, *et al.* hope that these cross-cutting concerns might work in other models too, but this is untested. The use of an aspect in a new model, when developed for another, seems untested in the community's literature writ large and is a noted omission in the conclusion of this particular work.

2.5 Variable Behaviour in Simulation and Modelling

In [2], PyDySoFu was used to model behaviour that changed as the simulation progressed. Behaviour undergoing variance appears in literature from many fields, but some themes stand out. Researchers are often interested in:

- Removing small variations from datasets in order to mine the original process (that real-world actors might be deviating from), referred to as sanitisation,
- Inserting variations so as to produce datasets with

2.5.1 Business Process Modelling & variation in behaviour

Is this going from data to a model? Models to data? Potentially multiple models either way depending on their use? Consider this and possibly restructure.

In real-world business processes, natural variation is difficult to avoid. This is because business processes are inherently socio-technical, and so can be expected to exhibit at least slight variations due to the mistakes of human actors executing those processes. Variations can effectively take two forms:

- ① Some variations are expected, where predictable shifts in behaviour emerge over time. Examples would be habits forming which deviate from prescribed processes, skipped steps, or paths of a fork in a process becoming effectively ignored as others become the default (essentially producing redundancy in the model).
- ② Unexpected variations can occur if an actor behaves erratically, information is improperly recorded in a log (and so *seems* to exhibit variance), or if some accident occurs. This appears as random noise in collected data,

and is difficult to statically embed in any model, as change might take myriad forms and occur at an arbitrary number of points in the process.

As these two forms of variance must be modelled differently, they are treated differently in a business process model exhibiting variance. Typically the second is treated as noise: undesired and a distraction from a model built to reflect a prescribed process. They are therefore removed via sanitisation, and are discussed in ???. The first, variation which might be reflective of a model as it can be expected to be *executed* — even when this was not an intended or prescribed version — might be interesting to modellers. These situations might arise, for example, where sociotechnical variance within the context of the broader system is the specific subject of investigation. Degraded modes in these systems are a good example of this [27].

2.5.2 Variations in Process Models

Discussing

2.5.3 Process mining & variation in data

Write a short subsection on the trend of sanitising data / coping with noisy datasets. Sometimes behavioural variance isn't desirable. Consider making noisy data a subsubsection of ???.

So, there's some stuff to cite on sanitisation and mining in the presence of noise — see On process model synthesis based on event logs with noise, [28].

Add more stuff to be cited here, at the very least...

Process Mining is a field which necessarily deals with erroneous data. As processes are identified within event logs sourced from real systems, inconsistencies in data collected or execution of a prescribed process results in data fed to a mining algorithm which is at best not indicative of the desired result, and at worst indicative of a different one altogether. As a result, variation in process logs is a subject of active research in the community.

There are two main research efforts involving mining on log data with variance:

- ① Some researchers look to minimise the impact of log variance on the outcome of process mining. This can be done through the development of mining algorithms which are able to cope somewhat with variance. Many algorithms attempting to solve this problem have been developed, but

their effectiveness depends on the kind of variance present and the degree to which those different variances are expressed in the data[28].

- ② Other researchers look to identify noise in event logs before they are mined, processing them to eliminate any variance before mining begins. This requires classification of noise and the removal of suspect traces[29].

Another perspective on the problem is that noise cannot be successfully eliminated, but that training on empirical noise limits a researcher’s control over an experiment. The argument here is usually along the lines that empirical noise is effectively impossible to predict, exert control over, or classify entirely, so any testing of tools using that data is flawed. Without knowing an algorithm’s response to specific kinds of variance, a researcher can’t compare one approach properly or reproducibly against another. It is therefore important to *produce logs with controlled kinds of variance*, so as to create a kind of synthetic workspace where algorithms are tested against synthetic data with known kinds and degrees of variance. Once they are reproducibly tested against known good datasets, they can undergo empirical verification by using data captured “in the wild”.

Naturally, similar approaches exist outside of process mining, as the requirement for synthetic data is a common one. In potentially sensitive data collected on the public — census or health data for example — there may be a need to publish data which is at least partially synthetic [30]–[32]. Drechsler and Reiter presents an array of simple statistical methods for producing this [31]. Koenecke and Varian note that, depending on the nature of the data *needed* for a given application, different methods are appropriate, meaning a variety of techniques are required [33], and give an overview of methods suitable in economics. MetaSim[34] produces data using probabilistic grammars for training neural nets in a manner naturally resilient to variance by including an appropriate amount via the trained grammar, thereby injecting a guaranteed correct degree of noise. Admittedly neural nets are a common source of synthetic data in modern literature and a research subject with a growing need for training data, perhaps best exemplified in the community’s production of another kind of neural net specifically for this purpose: Generative Adversarial Nets, or GANs [35].

Approaches specific to the generation of synthetic event logs are also abundant[18], [36]–[39], and as PyDySoFu’s original use was in socio-technical modelling, this is our primary interest. **Shugurov paper from 2014 should be checked for citing here too. There’s a useful summary in On Business Process Variants Generation (Pourmasoumi 2015).**

In [36], [40] Stocker and Accorsi describe a method for injecting variance

into synthetic event logs (“traces”). In [36], a method is described whereby security-specific alterations to traces can be made which represent the behaviour of an attacker in some socio-technical system. Variance can be injected by statically manipulating a process before simulating it to generate traces or making modifications to traces after simulation. A supporting tool, “Secsy”, is provided in [40].

A similar approach approach is provided in [37], where alterations are made to a process before simulation occurs. In terms of alterations made to the model directly (and not produced traces), Secsy only supports a limited number of operations on a model: transformation of AND and OR gateways to the alternative kind, and swapping the ordering of modelled activities. The method proposed in [37] is able to make use of a much broader gamut of alterations, by limiting themselves to mutating only block-structured processes which they represent as “structure trees”. The authors claim this limitation does not impact the broad applicability of their approach, as Li claims in their own PhD thesis notes that around 95% of BPMN models can be represented this way [41]. However, that claim should be held with some scepticism. The citations for this claim are [42] and a paper by Polyvyanny, García-Bañuelos, and Dumas which is most likely [43]⁸. The first work checks 214 process models against a set of patterns which are specifically not formalised, and the second presents some formal work on the translation of process models following the block structure relied on in [37] and [41]. However, the two works never cite each other, the application of the formal translations to the patterns presented is non-trivial, and no further explanation as to the application required appears to be presented in the thesis. One could suppose that the translation of the patterns to block-structuring could be automated by an implementation of the theory presented in [43]. In any case, Li notes in [41] that the requirement of block-structuring on a process model is a limiting factor in the application of their own work in their conclusions, and so the broad applicability of [37] should be taken with healthsome caution. After some exhaustive citation reading we can conclude that neither approach supports effective production of synthetic event logs exhibiting a wide gamut of variances by statically manipulating a BPMN model prior to simulation, although the existence of both methods suggests it would be a valuable research outcome.

In [18], we can find good things I should write about.

In [38], we can find good things I should write about.

In [39], we can find good things I should write about.

⁸The citation indicates a paper presented a year earlier than [43], and the author has given talks with the same title and published other works with similar titles — although it is possible a paper with the same name was published a year earlier, and this could change Li’s claim, some confidence can be had that this is a simple referencing error or typo.

In [44], Eno and Thompson develop a workflow for generating synthetic data by converting from PMML, a markup language for predictive modelling, into their own DSL for storing synthetic data specifically. PMML is designed specifically with mining in mind, but SDDL is a flexible language useful for storing arbitrary synthetic data, and they note that the technique might be applicable to anybody generating synthetic datgasets for privacy or security concerns, or for those looking to generate much larger datasets than can be feasibly collected empirically. Their goal is to synthesise datasets indistinguishable in specified characteristics to real-world datasets.

- ① [39] — unread, v interesting
- ② [36] for secsy, and [40], the associated tooling paper
- ③ [37] generates synthetic logs with variance, just like secsy, but instead of making edits to the process before simulation using a “structure tree” representation and identifying points suitable for mutation.
- ④ [38] — unread, v interesting
- ⑤ [18] — Aalst generating logs from models. No variance but they make the case that synthetic data is needed by the community *and* it’s a big name taking a swing, too. Could combine well with [37] to get variance without actually producing new techniques, assuming a limitation of the sim approach to block-structured models (which I think they already impose anyway...)

2.6 Research Opportunities in the Literature

One notable omission from the set of research themes outlined in section 2.5 is that variations on processes aren’t well studied in their own right. That is to say, behavioural variation is typically treated as a nuisance to the researcher or practitioner interested in a model or dataset, and the variations and their impact on simulations are not studied on their own. Some work exists **cite work that focuses on variations specifically**, but the majority of this is done in the context of tooling, i.e. the representation of variation for their use in another research context, where they are not the subject.

As an aspect orientation framework capable of runtime adaptation of a target system which can manipulate a join point from inside **what are our research opportunities given the above...?**

Bibliography

- [1] W. Wallis, W. Kavanagh, A. Miller, and T. Storer, "Designing a mobile game to generate player data-lessons learned," 2020.
- [2] T. Wallis and T. Storer, "Modelling realistic user behaviour in information systems simulations as fuzzing aspects," in *International Conference on Advanced Information Systems Engineering*, Springer, 2018, pp. 254–268.
- [3] —, *Pydysofu*, <https://github.com/twsswt/pydysofu>, 2018.
- [4] —, "Process fuzzing as an approach to genetic programming," in *Proceedings of the SICSA Workshop on Reasoning, Learning and Explainability*, K. Martin, N. Wiratunga, and L. S. Smith, Eds., ser. CEUR Workshop Proceedings, vol. 2151, Aberdeen, UK: CEUR-WS.org, Jun. 2018. [Online]. Available: http://ceur-ws.org/Vol-2151/Paper_S3.pdf.
- [5] G. Kiczales, J. Lamping, A. Mendhekar, C. Maeda, C. Lopes, J.-M. Loingtier, and J. Irwin, "Aspect-oriented programming," in *European conference on object-oriented programming*, Springer, 1997, pp. 220–242.
- [6] G. Kiczales, E. Hilsdale, J. Hugunin, M. Kersten, J. Palm, and W. Griswold, "An overview of aspectj," in *ECOOP*, 2001.
- [7] J. C. Hveding, "An aspect-oriented approach to adaptive systems," M.S. thesis, Institutt for datateknikk og informasjonsvitenskap, 2005.
- [8] A. Popovici, T. Gross, and G. Alonso, "Dynamic weaving for aspect-oriented programming," in *Proceedings of the 1st International Conference on Aspect-Oriented Software Development*, ser. AOSD '02, Enschede, The Netherlands: Association for Computing Machinery, 2002, pp. 141–147, ISBN: 158113469X. DOI: [10.1145/508386.508404](https://doi.org/10.1145/508386.508404). [Online]. Available: <https://doi.org/10.1145/508386.508404>.
- [9] R. Chitchyan and I. Sommerville, "Comparing dynamic ao systems," 2004.

- [10] A. Popovici, G. Alonso, and T. Gross, "Just-in-time aspects: Efficient dynamic weaving for java," in *Proceedings of the 2nd International Conference on Aspect-Oriented Software Development*, Association for Computing Machinery, 2003, pp. 100–109.
- [11] J. Baker and W. Hsieh, "Runtime aspect weaving through metaprogramming," in *Proceedings of the 1st international conference on Aspect-oriented software development - AOSD '02*, ACM Press, 2002. DOI: [10.1145/508386.508396](https://doi.org/10.1145/508386.508396). [Online]. Available: <https://doi.org/10.1145/508386.508396>.
- [12] J. Baker and W. C. Hsieh, "Maya: Multiple-dispatch syntax extension in java," in *Proceedings of the ACM SIGPLAN 2002 Conference on Programming Language Design and Implementation*, ser. PLDI '02, Berlin, Germany: Association for Computing Machinery, 2002, pp. 270–281, ISBN: 1581134630. DOI: [10.1145/512529.512562](https://doi.org/10.1145/512529.512562). [Online]. Available: <https://doi.org/10.1145/512529.512562>.
- [13] R. Dyer and H. Rajan, "Supporting dynamic aspect-oriented features," *ACM Transactions on Software Engineering and Methodology*, vol. 20, no. 2, pp. 1–34, Aug. 2010. DOI: [10.1145/1824760.1824764](https://doi.org/10.1145/1824760.1824764). [Online]. Available: <https://doi.org/10.1145/1824760.1824764>.
- [14] R. Keller and U. Hölzle, "Binary component adaptation," in *European Conference on Object-Oriented Programming*, Springer, 1998, pp. 307–329.
- [15] K. Thompson, "Reflections on trusting trust," in *ACM Turing award lectures*, 2007, p. 1983.
- [16] A. Maria, "Introduction to modeling and simulation," in *Proceedings of the 29th Conference on Winter Simulation*, ser. WSC '97, Atlanta, Georgia, USA: IEEE Computer Society, 1997, pp. 7–13, ISBN: 078034278X. DOI: [10.1145/268437.268440](https://doi.org/10.1145/268437.268440). [Online]. Available: <https://doi.org/10.1145/268437.268440>.
- [17] D. DORI, "Object-process Analysis: Maintaining the Balance Between System Structure and Behaviour," *Journal of Logic and Computation*, vol. 5, no. 2, pp. 227–249, Apr. 1995, ISSN: 0955-792X. DOI: [10.1093/logcom/5.2.227](https://doi.org/10.1093/logcom/5.2.227). eprint: <https://academic.oup.com/logcom/article-pdf/5/2/227/6244720/5-2-227.pdf>. [Online]. Available: <https://doi.org/10.1093/logcom/5.2.227>.
- [18] A. A. Mitsyuk, I. S. Shugurov, A. A. Kalenkova, and W. M. van der Aalst, "Generating event logs for high-level process models," *Simulation Modelling Practice and Theory*, vol. 74, pp. 1–16, May 2017. DOI: [10.1016/j.simpat.2017.01.003](https://doi.org/10.1016/j.simpat.2017.01.003). [Online]. Available: <https://doi.org/10.1016/j.simpat.2017.01.003>.

- [19] I. Machado, R. Bonifácio, V. Alves, L. Turnes, and G. Machado, "Managing variability in business processes," in *Proceedings of the 2011 international workshop on Early aspects - EA '11*, ACM Press, 2011. DOI: [10.1145/1960502.1960508](https://doi.org/10.1145/1960502.1960508). [Online]. Available: <https://doi.org/10.1145/1960502.1960508>.
- [20] C. Cappelli, J. C. Leite, T. Batista, and L. Silva, "An aspect-oriented approach to business process modeling," in *Proceedings of the 15th workshop on Early aspects - EA '09*, ACM Press, 2009. DOI: [10.1145/1509825.1509828](https://doi.org/10.1145/1509825.1509828). [Online]. Available: <https://doi.org/10.1145/1509825.1509828>.
- [21] L. Gulyás and T. Kozsik, "The use of aspect-oriented programming in scientific simulations," in *Proceedings of Sixth Fenno-Ugric Symposium on Software Technology, Estonia*, 1999.
- [22] A. Aksu, F. Belet, and B. Zdemir, "Developing aspects for a discrete event simulation system," in *Proceedings of the 3rd Turkish Aspect-Oriented Software Development Workshop*, Bilkent University, 2008, pp. 84–93.
- [23] M. Chibani, B. Belattar, and A. Bourouis, "Using aop in discrete event simulation: A case study with japosim," *International Journal of Applied Mathematics, Computational Science and Systems Engineering*, vol. 1, 2019.
- [24] T. B. Ionescu, A. Piater, W. Scheuermann, E. Laurien, and A. Iosup, "An aspect-oriented approach for disaster prevention simulation workflows on supercomputers, clusters, and grids," in *2009 13th IEEE/ACM International Symposium on Distributed Simulation and Real Time Applications*, IEEE, 2009. DOI: [10.1109/ds-rt.2009.35](https://doi.org/10.1109/ds-rt.2009.35). [Online]. Available: <https://doi.org/10.1109/ds-rt.2009.35>.
- [25] M. Cieslak, A. N. Seleznyova, P. Prusinkiewicz, and J. Hanan, "Towards aspect-oriented functional-structural plant modelling," *Annals of Botany*, vol. 108, no. 6, pp. 1025–1041, Jul. 2011. DOI: [10.1093/aob/mcr121](https://doi.org/10.1093/aob/mcr121). [Online]. Available: <https://doi.org/10.1093/aob/mcr121>.
- [26] A. Lindenmayer, "Mathematical models for cellular interactions in development i. filaments with one-sided inputs," *Journal of theoretical biology*, vol. 18, no. 3, pp. 280–299, 1968.
- [27] C. W. Johnson and C. Shea, "A comparison of the role of degraded modes of operation in the causes of accidents in rail and air traffic management," in *2007 2nd Institution of Engineering and Technology International Conference on System Safety*, IET, 2007, pp. 89–94.

- [28] A. A. Mitsyuk and I. S. Shugurov, "On process model synthesis based on event logs with noise," *Automatic Control and Computer Sciences*, vol. 50, no. 7, pp. 460–470, Dec. 2016. DOI: [10.3103/s0146411616070154](https://doi.org/10.3103/s0146411616070154). [Online]. Available: <https://doi.org/10.3103/s0146411616070154>.
- [29] H.-J. Cheng and A. Kumar, "Process mining on noisy logs — can log sanitization help to improve performance?" *Decision Support Systems*, vol. 79, pp. 138–149, Nov. 2015. DOI: [10.1016/j.dss.2015.08.003](https://doi.org/10.1016/j.dss.2015.08.003). [Online]. Available: <https://doi.org/10.1016/j.dss.2015.08.003>.
- [30] R. J. Little, "Statistical analysis of masked data," *Journal of Official statistics*, vol. 9, no. 2, p. 407, 1993.
- [31] J. Drechsler and J. P. Reiter, "An empirical evaluation of easily implemented, nonparametric methods for generating synthetic datasets," *Computational Statistics & Data Analysis*, vol. 55, no. 12, pp. 3232–3243, Dec. 2011. DOI: [10.1016/j.csda.2011.06.006](https://doi.org/10.1016/j.csda.2011.06.006). [Online]. Available: <https://doi.org/10.1016/j.csda.2011.06.006>.
- [32] D. B. Rubin, "Discussion statistical disclosure limitation," *Journal of official Statistics*, vol. 9, no. 2, p. 461, 1993.
- [33] A. Koenecke and H. Varian, *Synthetic data generation for economists*, 2020. arXiv: [2011.01374](https://arxiv.org/abs/2011.01374) [econ.GN].
- [34] A. Kar, A. Prakash, M.-Y. Liu, E. Cameracci, J. Yuan, M. Rusiniak, D. Acuna, A. Torralba, and S. Fidler, "Meta-sim: Learning to generate synthetic datasets," in *2019 IEEE/CVF International Conference on Computer Vision (ICCV)*, IEEE, Oct. 2019. DOI: [10.1109/iccv.2019.00465](https://doi.org/10.1109/iccv.2019.00465). [Online]. Available: <https://doi.org/10.1109/iccv.2019.00465>.
- [35] I. J. Goodfellow, J. Pouget-Abadie, M. Mirza, B. Xu, D. Warde-Farley, S. Ozair, A. Courville, and Y. Bengio, *Generative adversarial networks*, 2014. arXiv: [1406.2661](https://arxiv.org/abs/1406.2661) [stat.ML].
- [36] T. Stocker and R. Accorsi, "Secsy: Security-aware synthesis of process event logs," in *Proceedings of the 5th International Workshop on Enterprise Modelling and Information Systems Architectures, St. Gallen, Switzerland*, Citeseer, 2013.
- [37] A. Pourmasoumi, M. Kahani, E. Bagheri, and M. Asadi, "On business process variants generation.," in *CAiSE Forum*, 2015, pp. 179–188.
- [38] D. Loreti, F. Chesani, A. Ciampolini, and P. Mello, "Generating synthetic positive and negative business process traces through abduction," *Knowledge and Information Systems*, vol. 62, no. 2, pp. 813–839, Jun. 2019. DOI: [10.1007/s10115-019-01372-z](https://doi.org/10.1007/s10115-019-01372-z). [Online]. Available: <https://doi.org/10.1007/s10115-019-01372-z>.

- [39] A. Yousfi, R. Saidi, and A. K. Dey, "Variability patterns for business processes in BPMN," *Information Systems and e-Business Management*, vol. 14, no. 3, pp. 443–467, Aug. 2015. DOI: [10.1007/s10257-015-0290-7](https://doi.org/10.1007/s10257-015-0290-7). [Online]. Available: <https://doi.org/10.1007/s10257-015-0290-7>.
- [40] T. Stocker and R. Accorsi, "Secsy: A security-oriented tool for synthesizing process event logs," in *BPM (Demos)*, 2014, p. 71.
- [41] C. Li, "Mining process model variants: Challenges, techniques, examples," Ph.D. dissertation, University of Twente, The Netherlands, 2010.
- [42] L. Thom, M. Reichert, and C. Iochpe, "Activity patterns in process-aware information systems: Basic concepts and empirical evidence," *Int. J. Bus. Process. Integr. Manag.*, vol. 4, pp. 93–110, 2009.
- [43] A. Polyvyanyy, L. García-Bañuelos, and M. Dumas, "Structuring acyclic process models," in *Business Process Management*, R. Hull, J. Mendling, and S. Tai, Eds., Berlin, Heidelberg: Springer Berlin Heidelberg, 2010, pp. 276–293, ISBN: 978-3-642-15618-2.
- [44] J. Eno and C. W. Thompson, "Generating synthetic data to match data mining patterns," *IEEE Internet Computing*, vol. 12, no. 3, pp. 78–82, 2008.